

Recall, not verification, is the bottleneck when frontier LLMs elucidate molecular structures from real spectra

Ilkham Yabbarov¹ (*corresponding: ilkhamfy@gmail.com*), Rudra Sondhi¹, Rodrigo A. Vargas-Hernández¹ ¹ Department of Chemistry and Chemical Biology, McMaster University, Hamilton, Ontario, Canada

Abstract

Given the molecular formula together with the infrared band list and $^1\text{H}/^{13}\text{C}$ shift lists exactly as reported in an open-access paper, how often does a frontier large language model (here, Claude) recover the correct molecular *constitution*? We find: **28% of the time** (top-1, $n=194$; 95% CI 22–35) — far below the near-100% implied by curated demonstrations. The bottleneck is not the model’s judgment but its *proposal*: on a 60-compound forward-verification subset the model proposes the true structure for only **19 of 60** compounds (31%, 95% CI 21–44), and where it does, forward-verification selects it in **16 of 19** (84%, 95% CI 62–94, against a 66% derangement floor; on the 13 of those 19 where more than one candidate existed and a choice was actually made, $10/13 = 77\%$, §5.2). **Recall, not verification, is the wall.** We establish this with three contributions. First, **IRexp** — the largest openly redistributable collection of *experimental* infrared **band lists** (121,233 records, a third structure-linked), mined from open-access literature and released with its mining pipeline (§2). Second, **IRSpectra-Bench** — an open, blind, mechanically scored benchmark of 194 compounds on which accuracy splits sharply by structural complexity, and a within-compound control showing that reported numbers are sensitive to how a problem is posed, not to raw capability alone (§3–§4). Third, a training-free **forward-verification** method that acts on the diagnosis: generating regiochemistry-aware candidates and re-ranking them by predicted-versus-observed ^{13}C raises both candidate recall and top-1, though at this sample size the top-1 gain is directional rather than resolved (§5.3). The gain is real but bounded, because within the training-free regime accuracy is capped by the product of generation recall and verification precision. That ceiling is a property of training-free elicitation rather than an established limit of the task: a small generator fine-tuned on IRexp lifts recall substantially and gives the highest full-benchmark accuracy we report (§5.6). But the wall moves rather than falls — the true structure remains outside the candidate pool for roughly half of compounds, and verification precision is untouched. The recall-bound pattern recurs in a battery-electrolyte domain (§4.5), where the true structure enters the candidate set for 13 of 46 compounds and is ranked first in 12 of those 13. A formula-only control confirms the spectra are doing the work: masking them drops top-1 from 23% to 5% on the same compounds, with every formula-only success also solved from spectra and none the reverse (§4.6). Every result is single-vendor (Claude); a cross-vendor test is the key open question. The core protocol uses no model training and no paid API — two clearly-fenced probes (§5.4, §5.6) are the only exceptions — and all data, predictions, and code are released.

1. Introduction

Determining a molecule’s structure from its spectra is a central, time-consuming task in synthetic and analytical chemistry. The dominant machine-learning approach trains specialised encoder–decoder models to map spectra to structures; the recent *Spectro* model, for example,

learns $^1\text{H}/^{13}\text{C}/\text{IR} \rightarrow \text{SELFIES}$ from a corpus of 6,833 molecules.¹ In parallel, general-purpose LLMs have been reported to perform the same task off-the-shelf: a 2026 industrial white paper (a non-peer-reviewed company report) found that Claude Opus matched or beat commercial NMR-prediction software in the forward direction (structure $\rightarrow$ spectrum, ± 0.08 ppm ^1H) and “recovered all eight simpler structures on every attempt” in the inverse direction (spectrum $\rightarrow$ structure).² We treat its numbers as a motivating claim to be tested against peer-reviewed benchmarks, not as an established baseline.

These results are striking, but the evaluation that supports them is narrow: 15 inverse problems on curated single-ring or two-fragment molecules, NMR only, with the seven harder targets additionally given the *starting-material structure* as a hint, and “recovery” scored leniently over three runs and three ranked candidates. The question a practising chemist actually faces — *take an arbitrary experimental spectrum from a paper and recover the structure* — is left open. Answering it requires (i) a large, diverse, real benchmark; (ii) a blind, reproducible scoring protocol; and (iii) honest accounting for the methodological choices that inflate or deflate the apparent score.

We provide all three. We build IRexp (§2), a literature-mined dataset that supplies the IR modality absent from prior LLM evaluations alongside paired NMR and resolved structures. We define a blind benchmark on it (§3) and measure inverse-elucidation performance under matched and mismatched methodologies (§4), isolating a large solver-methodology effect with a within-compound control. We then introduce forward-verification elucidation (§5), a training-free method that turns the model’s *strong* direction (forward prediction) against its *weak* one (inverse regiochemistry), and use it to localise the bottleneck. The central finding is a sharp asymmetry: given a candidate set that contains it, the model *verifies* the correct structure 84% of the time, yet *proposes* it for only 31% of compounds — recall, not verification, is the wall (Fig. 1). The end-to-end study design is summarised in Fig. S1. Throughout, the solver and verifier are LLM agents run under a consumer subscription — no fine-tuning, no API spend. That makes the protocol cheap to *re-run*, but cheap is not the same as reproducible: the subscription harness pins no model snapshot, exposes no temperature or seed, and carries an undisclosed product system prompt, so an outside group reproduces it distributionally rather than exactly (§8). What *is* exactly reproducible is the scoring — all predictions, ground truth and scorers are released, so every number in the training-free core (§3–§5.3, §4.6) regenerates from the released artifacts. The two trained probes are the exception: §5.4’s HOSE lookup and GNN are built from an nmrshiftdb2 dump that we cannot redistribute, so those rows regenerate only once a reader supplies the same dump (see *Data and code availability*). Two trained probes (§5.6, §5.7) are reported separately as fenced complements. By construction this is an open-resource contribution in the remit of *Digital Discovery*: an openly licensed, structure-linked spectral dataset; a pre-registered, mechanically scored benchmark with released ground truth and scorer; and a training-free, zero-paid-API core pipeline (with two trained probes fenced in the SI) in which every main-text figure regenerates from the released code and predictions. We prioritise an honest, reusable measurement of where current models stand over a leaderboard-topping number.

Contributions. 1. **IRexp** — the largest openly *redistributable* experimental-IR dataset (121,233 records; 43,060 structure-linked; 33,201 IR+ $^1\text{H}+^{13}\text{C}$ +structure): larger by raw record count than any prior collection and uniquely released under permissive licences for bulk model development (view-only resources such as SDBS are larger in structure-linked spectra but are not redistributable; §2.1), with a reproducible mining pipeline. 2. An **open multimodal benchmark** and the first blind, mechanically scored, complexity-stratified evaluation of frontier-LLM structure elucidation on real data — measured in depth on one model family (Claude) — reconciling the gap to op-

timistic prior reports. Crucially, the ground truth is experimental structures from the published literature, resolved deterministically (OPSIN/RDKit) and expert-audited: no LLM curates the labels or scores the predictions. The model is the system under test, not the source of its own answers. 3. A **diagnostic decomposition** of LLM elucidation into recall and verification, showing that **recall — not verification — bounds current performance**, obtained via a training-free forward-verification probe that yields a bounded improvement (top-1 23%→30%) but, by our own measurements, cannot exceed a recall/precision ceiling without sharper verification or 2D-NMR data.

1.1 Related work

Trained spectra→structure models. The dominant line trains sequence or graph decoders to emit structures from spectra: *Spectro* ($^1\text{H}/^{13}\text{C}/\text{IR}\rightarrow\text{SELFIES}$)¹, the multitask CNN+transformer of routine 1D-NMR³, and set/graph transformers such as NMRTrans⁴. These reach high accuracy *in-distribution* but require a labelled spectra→structure corpus, which is exactly the scarce resource our IRexp pipeline targets, and they are retrained per modality. Closest in spirit to our multimodal setting is **NMIRacle**⁵, a generative model conditioned jointly on IR + ^1H + ^{13}C ; it is a strong trained baseline, whereas our contribution is the *open experimental-IR data it (and others) can train on*, plus a *training-free* protocol and a blind benchmark to measure it.

LLMs as elucidators, and how this work differs. General-purpose LLMs have been applied off-the-shelf — Anthropic’s forward/inverse white paper², SpectraLLM and MolSpectLLM (multimodal LLMs over multi-spectral input)^{6,7}, and knowledge-enhanced tree-search reasoning⁸ — and dedicated multimodal *benchmarks* now exist, most prominently **MolPuzzle**⁹ (IR+MS+ ^1H + ^{13}C elucidation puzzles with the molecular formula supplied). These establish that the task is worth benchmarking and that LLMs can be scored on multi-spectral input; we claim priority on neither. What is new here is three things that, to our knowledge, no prior benchmark combines. **(i) Real, literature-mined experimental spectra at scale:** prior multimodal benchmarks and trained baselines^{1,5,10} rely on simulated, software-predicted, or hand-curated puzzle spectra, whereas IRSpectra-Bench is drawn from IRexp’s experimental IR + ^1H + ^{13}C of out-of-distribution compounds, so it measures the simulated-to-real gap rather than in-distribution skill. The closest trained precedent, Alberts et al.¹⁰, learns an IR→structure transformer from ~635k *simulated* spectra with a 3,453-spectrum experimental fine-tune (top-1 44% on 6–13 heavy atoms); IRexp instead releases experimental IR at scale as an open, redistributable resource and pairs it with a blind LLM benchmark. **(ii) Blind, fully specified, mechanical scoring:** reported accuracies on the same task vary enormously with inference method and scoring harness — GPT-4o on MolPuzzle ranges from 1.4%⁹ to 27.8%⁸, the latter using knowledge-enhanced tree-search reasoning⁸ — so numbers across papers are not directly comparable; we therefore fix and release a single, pre-registered, RDKit-InChIKey protocol with bootstrap CIs. **(iii) A recall/verification decomposition:** existing benchmarks report a single aggregate score, whereas we factor performance into generation recall and verification precision and show *where* the task is lost — the result that transfers across models, domains, and verifier choices.

Computational NMR for structure validation. Our forward-verification method is the LLM analog of a workflow chemists already trust: assign a structure by computing the spectrum each candidate *would* give and matching it to experiment. In solution, this is the DP4 / DP4+ probabilistic framework over GIAO-DFT shifts^{11,12}; in the solid state it is **NMR crystallography**, where GIPAW-computed shifts adjudicate between candidate structures^{13,14}. We replace the quantum-chemical predictor with a forward LLM, trading accuracy for zero setup cost, and inherit the same

core principle — *verification by forward prediction is easier than inverse generation.*

2. The IRExp dataset

2.1 Motivation

What an IRExp record is. Each record is a **band list** — the peak positions in cm^{-1} that an author transcribed into the text of a paper (e.g. “IR (KBr): 3373, 3045, 2914, 1664 cm^{-1} ”) — together with the $^1\text{H}/^{13}\text{C}$ shift lists reported alongside it and, where resolvable, the compound’s 2D structure. **IRExp contains no absorbance traces.** This is the form in which experimental IR is overwhelmingly *published*, and it is the form a language model consumes, but it is a different object from a digitised spectrum and the two should not be counted against one another.

Open experimental IR is scarce *in a redistributable, ML-ready form*, and the existing resources divide into two kinds that are not directly comparable.

Digitised spectra (absorbance traces). The largest freely downloadable collections are the NIST WebBook¹⁵ (~16k IR spectra) and the Chemotion electronic-lab-notebook deposit¹⁶ (~2k). The AIST SDBS¹⁷ is larger (~54k FT-IR, all structure-linked) but is *view-only*: it caps downloads at 50 spectra/day, forbids commercial use, and ships no permissive licence or bulk export, so it cannot be used to train or redistribute open models at scale. Commercial libraries (e.g. Wiley KnowItAll, ~ 10^5 spectra) are larger still but closed. **On this axis IRExp is not the largest resource, and does not compete: SDBS alone holds more structure-linked spectra than IRExp holds structure-linked band lists.**

Band lists. Against text-derived peak-list resources, no large, permissively-licensed, structure-linked, bulk-reusable IR collection existed. IRExp fills that gap, and to our knowledge is the largest *openly redistributable* collection of experimental IR **band lists** by record count. The claim is scoped to that object type deliberately; the contribution is redistributable scale in the published-data modality, not a spectral library.

2.2 Construction

A per-compound IR band list, together with co-reported $^1\text{H}/^{13}\text{C}$ NMR, follows a remarkably stable textual convention in the experimental sections of organic chemistry papers. We exploit this with a browser-free harvesting agent:

- **Discovery.** Open-access primary literature is enumerated through the NCBI E-utilities and harvested in bulk from the PMC Open-Access Subset¹⁸ on AWS S3 (plain HTTPS, no anti-bot, fully redistributable CC-BY content), supplemented by the Chemotion FT-IR deposit (RADAR4Chem, CC-BY-SA-4.0).
- **Extraction.** A deterministic parser segments experimental text into per-compound records and extracts IR wavenumbers and $^1\text{H}/^{13}\text{C}$ shift lists, with quality gates that reject instrument scan-range artefacts and prose false-positives (band-list density, ≥ 4 bands, plausible 400–4000 cm^{-1} window).
- **Structure resolution.** In-text IUPAC names are converted to SMILES with OPSIN¹⁹, canonicalised with RDKit²⁰ (InChIKey, SELFIES²¹), with a PubChem²² fallback for trivial/natural-product names. A key engineering finding was that the dominant open-access

main-text convention labels compounds with *letter-prefixed* series labels (e.g. "...carbothioamide (**B1**):") rather than the digit-first "(3a)" of supporting-information sections; capturing these and cleaning narrative/PDF artefacts before resolution raised structure coverage from 24% to **35%**.

The pipeline is browser-free by design. This is not an anti-bot bypass — the bulk corpus is fully open — but a throughput choice: plain HTTP from a bulk corpus parses at hundreds of papers per second, which is what makes a 10^5 -scale harvest feasible where a per-page browser driver (as used to build prior 10^3 -scale sets) is not.

2.3 Contents and licensing

Table 1. IRExp dataset contents and provenance.

field	value
experimental IR band-list records	121,233
...co-reporting ^1H and/or ^{13}C NMR	87,075 (72%)
...with a resolved 2D structure	43,060 (35.5%)
...of these, with ^1H and/or ^{13}C NMR	40,702
...of these, with both ^1H and ^{13}C (full quadruples)	33,201
<i>by provenance</i> : author-transcribed from PMC-OA text (CC-BY)	119,345
<i>by provenance</i> : peak-picked from Chemotion ELN deposits (CC-BY-SA)	1,888

Both pools are stored in the same band-list form, but they are not the same object and we keep them separable. The PMC pool is transcribed from the experimental section of a paper and carries a median of 9 bands per record; the Chemotion pool derives from deposited spectra and, being peak-picked rather than author-selected, carries a median of 39. Users training on band density should treat the pools separately (the `ir_source` and `license` fields distinguish them).

Extraction fidelity is measured, not assumed. Because every record cites its source accession and PMC is open, transcription can be checked directly. On a seed-fixed random sample of 60 PMC-sourced records (`scripts/audit_extraction.py --n 60 --seed 0`), we re-fetched each source article and asked whether every recorded wavenumber appears in that article’s text: **560/560 bands and 60/60 records were confirmed** (Wilson 95% CI 99.3–100% and 94–100% respectively). This bounds *transcription* error — hallucinated, mis-parsed or unit-mangled values — at under 1% of bands. It does **not** measure whether the parser found every IR string in every paper, which is a recall question that requires human reading; that manual audit is prepared but not yet run (§7). The audit record is released at `data/audit/extraction_audit.json`.

A structure-complete split, **irexp_resolved** (43,060 records, 100% structure-linked), is the training-/benchmark-ready subset and is $\sim 6\times$ the 6,833-molecule set used to train Spectro¹ (Fig. S2). Provenance is 119,345 PMC-OA records (CC-BY) plus 1,888 Chemotion records (CC-BY-SA); the two licences are kept as separable pools. Each record is DOI-/accession-traceable. Re-resolution is additive and content-keyed, so the dataset can be re-enriched without re-mining.

3. Benchmark design (IRSpectra-Bench)

From `irexp_resolved` we draw **IRSpectra-Bench**, 194 blind elucidation problems. Each problem presents the **molecular formula** (as from HRMS), the **IR band list**, and the **^1H and ^{13}C shift lists** (with multiplicities and J-couplings where reported), and asks for the structure. No name, SMILES, or hint is given. Every main-round ground-truth structure is **spectrally validated** by an automated RDKit consistency check (^{13}C peak count vs symmetry-unique carbons, molecular-formula match, SELFIES round-trip), excluding records with merged or incomplete spectra (6/140 in the main round, leaving 134). The exclusions are itemised rather than asserted: five report more ^{13}C peaks than the structure has carbons (R03 28>24, R14 11>8, R65 26>18, R67 23>16, R131 34>17 — merged or contaminated spectra) and one is too sparse to constrain (R82, 5 peaks for 22 symmetry-unique carbons). The filter is deterministic and runs over all three rounds in `scripts/validate_benchmark.py`, which regenerates every `clean_qids.json` from the released questions and answers, so the cohort behind the headline number is reproducible rather than a shipped artifact. The 60 compounds of the controlled rounds are retained as fixed, pre-registered sets for the difficulty and within-compound controls, with the same self-consistency audit reported separately rather than used to exclude (57/60 pass; §7). Problems are stratified by RDKit ring analysis: a compound is **simple** iff it has at most two rings (single ring or two separate ring fragments), no fused/spiro/bridgehead system, and ≤ 22 heavy atoms; every other compound — any fused/spiro/bridged or ≥ 3 -ring system, or more than 22 heavy atoms — is **complex** (98 simple / 96 complex). This binary difficulty axis is distinct from the continuous size gradient in §4.1, which bins by heavy-atom count (≤ 15 / 16–25 / > 25). The criterion is therefore exhaustive (the 23–24-atom band, 13 compounds, is classed complex by size). InChIKey de-duplication is applied across all rounds to prevent leakage.

Scoring is mechanical. A prediction is *correct* if its RDKit InChIKey connectivity layer (first 14 characters) matches the reference — i.e. we score *constitution* (atom-and-bond connectivity); stereochemistry is reported separately. This is a deliberate floor: a candidate with the correct constitution but wrong stereochemistry counts as correct. We report the strict alternative rather than assert it is immaterial. Scoring the *full* InChIKey (stereochemistry-sensitive) gives **21.1% top-1 and 25.8% recovered** (41/194 and 50/194), against 28.4% / 33.5% at the connectivity layer — 7.3 points lower, so the choice of layer is not immaterial and the constitution figure should be read as an upper bound on full-stereochemistry accuracy. We nevertheless take constitution as the headline metric because 1D $^1\text{H}/^{13}\text{C}/\text{IR}$ rarely fixes absolute configuration, so a stereochemistry-aware score is not well-posed from the given data: only 10.3% (20/194) of benchmark answers carry a *defined* (assigned R/S) stereocentre, and a model is penalised there for information the prompt never contained. Both numbers are produced by `scripts/score_main.py --stereo`.

Metrics. We report five quantities. **Top-1 (exact constitution)** is the fraction of compounds whose single best-ranked candidate matches the reference at the InChIKey connectivity layer. **Recovered (top-3)** is the fraction for which the reference appears among the up-to-three ranked candidates returned (matching the lenient “recovery” protocol of ref.²). Where this paper writes *recovery* it is always reporting another study’s metric under that study’s own name; our two quantities are always **recovered (top-3)** and **generation recall**, which are different denominators and are never used interchangeably. **Generation recall** is the fraction for which the reference is present in the candidate pool *before* re-ranking — the ceiling any verifier can reach. **Conditional-on-recall precision** is the verifier’s hit rate over recall-positive compounds only, isolating verification quality

from generation. The forward verifier ranks candidates by a symmetric **chamfer distance** between predicted and observed ^{13}C peak sets (for each predicted peak, the distance to its nearest observed peak, summed, and symmetrically for each observed peak; lower is better, with no equal-count requirement). We also report Morgan(2, 2048)²³ Tanimoto as a graded "right scaffold/family" signal. CIs are bootstrap 95% over compounds; model-vs-model differences use McNemar’s exact test with Holm correction.

Solvers are LLM agents run under a consumer subscription. A frontier LLM (Claude Opus) is invoked as an independent sub-agent per batch of problems; agents are closed-book — configured with no web/tool access beyond an RDKit formula check and no access to ground truth, verified by grep-auditing their task transcripts at run time (audit logs available on request; the committed artefacts are the parsed per-compound predictions). Adherence to the supplied formula is high but imperfect: 91.3% (115/126) of forward-verification candidates match the given molecular formula exactly (100% of true structures, 89.6% of decoys), and 76.6% across the full top-3 pool — a residual generator error orthogonal to regiochemistry. This makes the benchmark free to run and reproducible without API credits.

4. How well do LLMs elucidate real structures?

4.1 Headline performance

Solver agents work blind from formula + IR + ^1H + ^{13}C and return up to three ranked candidates. Each agent handles a small batch of problems in a bounded context that is reset between batches (2–12 compounds per context in the released run; see `data/benchmark_main/raw/`), rather than one long context over the whole benchmark — the arm §4.3 shows matters. Over the full **194-compound benchmark** (134 spectrally-validated compounds + the 60 from the controlled rounds), with bootstrap 95% confidence intervals:

Table 2. Headline elucidation performance on IRSpectra-Bench (n=194), overall and by difficulty stratum, with bootstrap 95% CIs.

metric	overall (n=194)	simple (n=98)	complex (n=96)
top-1 exact constitution	28.4% [22–35]	48.0% [39–57]	8.3% [3–15]
recovered (within top-3)	33.5% [27–40]	54.1% [44–63]	12.5% [6–20]
scaffold-level (best Tanimoto ≥ 0.45)	56%	73%	39%
mean best Tanimoto	0.59	0.73	0.45

As a cohort-robustness check, restricting the controlled rounds to their spectrally-clean subsets as well (134 main-clean + 57 controlled-clean = 191 strictly-validated compounds) leaves the headline unchanged: top-1 **28.8%** (95% CI 23–36), recall 34.0%, with simple 49.0% / complex 8.4% — within ~1 point of the n=194 figures on every metric shown, with fully overlapping intervals, so the asymmetric inclusion of the pre-registered controlled sets does not drive the result.

The gradient is sharp and the intervals are tight. The 48%→8% simple→complex separation (Fig. 2) confirms the benchmark is discriminating across a realistic difficulty range, and accuracy falls

monotonically with molecular size in step with it — top-1 **60.5%** for ≤ 15 heavy atoms, 28.3% for 16–25, and **7.0%** above 25 (Fig. S3). The model recovers molecular formula and functional groups reliably and the scaffold often (best Tanimoto ≥ 0.45 for 56% of compounds), but the exact constitution far less often, failing predominantly on **regiochemistry** — *which* position a substituent occupies. This is consistent with an information limit of 1D data: many regioisomers have similar $^1\text{H}/^{13}\text{C}$ shifts, which is precisely why 2D experiments (HMBC, NOESY) exist.

4.2 Reconciling with prior reports

Our 28% top-1 sits far below the $\sim 100\%$ on “simple” molecules reported, in a non-peer-reviewed company white paper, for the same model class.² Because that figure has not been independently scored, we reconcile it against the *peer-reviewed* record (the MolPuzzle benchmark⁹ and its re-scorings⁸, and the trained baselines^{1,5}) rather than treating it as settled. The gap is largely attributable to methodology and scoring choices, not raw capability:

- **Difficulty.** “Single ring” by ring-count is not “easy”: our simple stratum includes, e.g., a hexasubstituted benzene whose regiochemistry has many realisations. Ring count is a poor proxy for elucidation difficulty.
- **Scoring.** Prior work counts a recovery if the reference appears among three ranked candidates over three independent runs; we report single-run top-1 and top-3.
- **Hints.** Prior hard targets received the starting-material SMILES, which fixes most of the scaffold; we give none.
- **Curation.** Prior compounds were hand-selected; ours are scraped and unfiltered for solvability.

On a like-for-like easy/hinted/lenient setup the numbers rise; on the realistic, hint-free, scraped regime, $\sim 28\%$ is the honest figure.

Versus trained models — a bound, not a leaderboard. A like-for-like comparison against specialised trained models is not available: no system has been scored on an identical test set, and published numbers differ in the three respects that most move the score — spectrum realism (simulated/curated vs. real), hints, and how “exact match” is defined. The strongest trained baselines report their accuracy *in-distribution on simulated spectra*. Spectro ($^1\text{H}/^{13}\text{C}/\text{IR} \rightarrow \text{SELFIES}$, 6,833 training molecules) reaches $\sim 90\%$ top-1 exact recovery — but on a 1,366-molecule held-out split whose IR is plotted from reference data and whose NMR is software-*predicted*, not experimental¹; and NMIRacle, which like us conditions jointly on $\text{IR}+^1\text{H}+^{13}\text{C}$ with *no* hints, reports 48% top-1 / 66% top-15 exact-SMILES recovery — again on held-out molecules from a *simulated* corpus drawn from the training distribution⁵. We measure 28.4% top-1 (33.5% top-3) on the full $n=194$ benchmark — rising to 30% with forward verification on the $n=60$ forward-verify set (§5) — on **blind, real, literature-mined experimental** spectra of out-of-distribution compounds. These are not comparable as a leaderboard; read as a *bound on the simulated-to-real gap*, the contrast suggests that high in-distribution accuracies substantially overstate real-world performance — the same gap we document for the LLM above. The instability of the metric itself reinforces the caution: the same $\sim 20\times$ MolPuzzle swing documented in §1.1 (GPT-4o: 1.4%⁹–27.8%⁸, method- and harness-dependent) bounds how much weight any single unaudited near-100% claim² can bear.

4.3 Methodology dominates: a within-compound control

The same 20 molecules were solved two ways: (a) by a single LLM context handling all of them sequentially with no tools, and (b) by four independent agents of five compounds each, with RDKit

formula-checking — i.e. bounded, reset contexts. On the *identical* compounds, **recovered (top-3)** rose from **5% to 15%** ($1/20 \rightarrow 3/20$) and top-1 from 0% to 15% ($0/20 \rightarrow 3/20$), with zero sample confound. Because §3 pins McNemar’s exact test for paired comparisons we apply it here too, and it does not reach significance: the top-1 arm is necessarily nested (arm (a) solved none) at $b=0$, $c=3$, giving $p=0.25$, and the recovered arm reaches at best $p=0.5$. The 15% point estimate carries a Wilson 95% CI of roughly 5–36%. This is therefore a directional within-compound demonstration, not an established multiplier, and we do not claim the $3\times$ as a measured effect size. Small rounds also swing widely (15–40% across $n=20$ –40 draws), which is why the headline is the full **194-compound** figure (28.4% top-1, 95% CI 22–35) rather than any single round. The practical lesson — bounded, frequently-reset contexts with tool access raise apparent performance substantially — also explains part of the gap to optimistic prior reports, whose per-problem API calls implicitly used method (b).

4.4 Model comparison: the benchmark ranks capability (and separates the extremes)

A benchmark is only useful if it separates models. On a fixed 24-compound subset solved blind by four Claude models — spanning a wide capability range, including the newest (Table 5) — under the identical protocol (Fig. 3):

One asymmetry must be disclosed, because §4.3 shows the variable it concerns has a large effect. The three comparison models each solved the 24 compounds as four six-compound contexts, whereas the Opus column reuses the headline run, in which those same 24 items were packed as one six-compound and two twelve-compound contexts. Context packing is exactly the factor §4.3 measures at 5%→15%, so the Opus point estimate is not strictly protocol-matched to the other three and, if anything, is the arm handicapped by longer contexts. This does not affect the nesting or the Fable-vs-Haiku contrast, but a clean re-run of the Opus arm under four six-compound contexts is the right fix and is listed as outstanding in `docs/MODELS.md`.

Table 3. Four-model comparison on a fixed 24-compound subset under the identical blind protocol.

model	top-1	recovered	top-1 95% CI
Claude Fable 5	45%	54%	[25, 67]
Claude Opus	25%	29%	[8, 42]
Claude Sonnet	20%	25%	[8, 38]
Claude Haiku	0%	4%	[0, 14]

CI’s are bootstrap except Haiku’s, which is the Clopper–Pearson exact interval for $0/24$ (the percentile bootstrap is degenerate at a boundary count of zero).

Three signals matter here. First, the four models rank in **monotonic capability order**, and the outcomes are **strictly nested** ($\text{Haiku} \subset \text{Sonnet} \subset \text{Opus} \subset \text{Fable}$ — each stronger model solves a superset of the compounds the weaker one does), with the smallest **floored at 0% exact** (4% recovered) on the same problems — so the benchmark is **capability-sensitive** and not a coin-flip. The nesting makes the *ranking* robust, but at $n=24$ the subset is **underpowered to separate adjacent models**: by McNemar’s exact test only the Fable-vs-Haiku gap survives multiple-comparison correction (Holm-adjusted $p=0.006$); the Fable-vs-Opus gap is not significant (uncorrected $p=0.063$, Holm-adjusted $p=0.19$), and Opus vs Sonnet are statistically indistinguishable — the bootstrap

CIs above overlap accordingly. Second, two mid-tier frontier models agree closely (Opus 25%, Sonnet 20%), so the recall-bound regime of §4.1 (top-1 in the high-20s%) is **not an artefact of a single model**. Third, the newest model nearly **doubles** the next-best top-1 (45% vs 25% on identical compounds; a large but, at this n, not-yet-significant gap) yet still misses the majority — IRSpectra-Bench is **hard and far from saturated even for the strongest model available**, with clear headroom. We ran four Claude-family models because they are callable for free under one subscription; a true cross-vendor sweep (GPT-, Gemini-class, open models) needs API access we deliberately did without, but the monotonic, large capability spread makes it unlikely the recall-bound pattern is specific to one model lineage. This is the behaviour a benchmark needs to remain informative as models improve.

4.5 Domain case study: battery-electrolyte chemistry

Structure elucidation from spectra is the daily inverse problem of the electrolyte-and-interphase community: the soluble decomposition products of carbonate and sulfone electrolytes, the additives that build the solid-electrolyte interphase, and the fluorophosphate/fluorosulfonyl species that NMR is routinely used to assign are exactly the molecules whose constitution must be read back out of IR and multinuclear NMR. To test whether the recall-bound regime above holds for *this* chemistry — rather than for organic molecules at large — we curated **IRSpectra-Bench-Electrolyte**, a 48-compound subset (eight per class as curated; two later yielded no parseable candidate, leaving **46 scored**) of structure-resolved IRexp records selected by substructure for the six functional families that dominate lithium- and sodium-battery electrolytes and their breakdown products: **carbonate** (linear/cyclic, the EC/DMC backbone), **sulfonyl/sulfonate** (sulfones, sulfonamides, the $-\text{SO}_2\text{CF}_3$ motif of imide salts), **nitrile** (acetonitrile- and adiponitrile-type high-voltage additives), **sp³-C-F** (fluorinated solvents and additives), **phosphoryl** (phosphates/phosphonates, the LiPF_6 -derived OPF chemistry), and **glyme/oligoether** (the ether solvents and PEG linkers). These are literature compounds bearing the *functional chemistry* of electrolytes, drawn from the open corpus; they are **not** operando or in-cell degradation spectra, and we make no claim that they are. They are excluded from every other split in this paper. Each compound was solved blind under the identical decoupled-agent protocol (Opus, closed-book, up to three ranked candidates, RDKit only for formula/parse checks).

Performance lands in the **same broad regime as the headline benchmark** — overall **top-1 26%, recovered 28%** (12/46 and 13/46; two compounds received no parseable candidate; at n=46 the 95% CI is wide and overlaps the headline) — consistent with the bottleneck being a property of the elucidation task rather than of any one chemical neighbourhood. The recall-bound *signature* is present in the same form: the true structure enters the candidate set for only 13 of 46 compounds, and of those 13 the solver already ranks it first in 12, so ranking is not what limits this subset either. We did **not** run forward-verification here, so this reproduces the recall-bound pattern under the solver’s own ranking; it does not independently reproduce the §5 verification-precision measurement. The per-class breakdown (Fig. S4) is itself informative:

Table 4. Per-class performance on IRSpectra-Bench-Electrolyte (n=46).

electrolyte class	n	top-1	95% CI	recovered (top-3)	95% CI
sp ³ -C-F	8	50%	[22, 78]	50%	[22, 78]
carbonate	7	29%	[8, 64]	43%	[16, 75]
phosphoryl	8	25%	[7, 59]	25%	[7, 59]

electrolyte class	n	top-1	95% CI	recovered (top-3)	95% CI
glyme / oligoether	7	29%	[8, 64]	29%	[8, 64]
sulfonyl / sulfonate	8	12%	[2, 47]	12%	[2, 47]
nitrile	8	12%	[2, 47]	12%	[2, 47]

Intervals are Wilson score intervals for a binomial proportion, appropriate at these counts where the percentile bootstrap is degenerate. They overlap almost completely.

At $n=7-8$ per class these per-class differences are within sampling noise (six-class χ^2 $p \approx 0.56$; best-vs-worst Fisher exact $p \approx 0.28$), so the ordering is a hypothesis-generating pattern rather than an established ranking. Read that way, it is chemically legible. **sp³-C-F** centres are the easiest: a C-F coupling fingerprint (the large $^1J/2J(C,F)$ splittings) localises the fluorine and pins regiochemistry, removing exactly the degeneracy that defeats elucidation elsewhere. **Sulfonyl** and **nitrile** are the hardest, for two distinct reasons that recur in real electrolyte analysis: sulfur/phosphorus **oxidation-state ambiguity** — sulfide vs. sulfoxide vs. sulfone, $-SCF_3$ vs. $-SO_2CF_3$, the very distinctions an interphase study must make — is poorly constrained by $^1H/^{13}C$ shifts alone (it lives in the IR and in heteronuclei the benchmark does not score); and nitrile-bearing targets in the corpus sit on heavily substituted heteroaromatic cores whose regiochemistry the 1D spectra underdetermine. This is the recall-bound failure of §4.1 reappearing in a domain-specific guise, and it is precisely where the **forward-verification** recipe of §5 — compute the spectrum each candidate *would* give and match it to the one observed — plays a role *analogous* to (not a replacement for) computational-NMR / NMR-crystallography structure validation, sharing its forward-prediction logic while trading DFT’s calibrated accuracy for zero setup cost. The subset, its per-class answers, and the scorer are released with the benchmark.

4.6 Is the model reading the spectra? A formula-only control

Because every benchmark compound is mined from open-access literature, a frontier model may have encountered it during pretraining. The cheapest decisive test of whether that explains the headline number is to take the spectra away. We re-ran the identical blind protocol with every spectral channel masked, so the solver receives the molecular formula and nothing else (`scripts/modality_ablation.py`, condition `formulaonly`; solvers were barred from reading any repository file or searching the web, and RDKit was permitted only to check that a proposed SMILES parses and matches the formula). A molecular formula does not determine constitution, so accuracy materially above the floor would indicate recall rather than reasoning.

Table 5. Formula-only control, paired on the same 60 compounds as §5.

condition	top-1	recovered (top-3)
formula only	3/60 (5%)	3/60 (5%)
formula + IR + 1H + ^{13}C	14/60 (23%)	19/60 (32%)

The outcomes are perfectly nested: **eleven** compounds are solved with the spectra and not without, and **none** the other way round (McNemar exact $p=0.001$). Removing the spectra removes the result.

A second, independent control: publication recency. The formula-only arm shows the spectra carry the signal, but not whether memorisation contributes to the part the spectra explain. A model can only have memorised a compound whose source paper was in its training corpus, and older papers have had longer in more corpora — so if recall drove the headline number, accuracy should fall with publication recency. We resolved the publication year of the source paper for **all 194** benchmark compounds from their accessions (`scripts/contamination_recency.py`); they span 2008–2026. Accuracy is flat: **28.6%** for the older half (≤ 2020 , $n=112$) against **28.0%** for the newer ($n=82$), a point-biserial correlation between publication year and correctness of $r = -0.007$, and no monotone trend across year buckets (Fig. 5b). The most recent bucket (≥ 2024 , $n=25$) is in fact the highest at 40% [23, 59].

The raw split is if anything biased *against* the newer half, because newer papers skew to larger molecules (median 22 heavy atoms against 20) and size is the dominant driver of lower accuracy (§4.1). Stratifying by heavy-atom band removes that confound. Newer compounds lead in the two bands that carry most of the accuracy (≤ 15 : 64% vs 58%; 16–25: 34% vs 25%) and trail slightly in the largest band, where both are near the floor (> 25 : 6% vs 8%). The size-adjusted older-minus-newer difference is **−5.1 points, 95% CI [−17.2, +7.0]** (Cochran–Mantel–Haenszel $\chi^2=0.42$, $p=0.51$, continuity-corrected). Its point estimate has the opposite sign from what recall out of pretraining would produce, but the interval comfortably includes zero — so this **bounds** any recency effect rather than demonstrating a reversed one, and by the same standard §5.3 applies to its own adjacent conditions we do not read it as directional.

The 5% is not zero, and it is worth saying what those three compounds are rather than rounding them away: a 2-(trimethylsilyl)aryl sulfonate, whose Si/S/Cl/F₂ composition is a near-unique benzyne-precursor signature; N-tosyl leucine; and a vanillyl alkanone. In each case an unusual element combination or a very common derivative class makes the formula close to determining, which is chemical inference from composition rather than evidence of having memorised this benchmark — though this experiment cannot separate the two. What it does establish is a bound: formula-level recall accounts for at most about a fifth of the headline accuracy on this set. Together with the recency result, the two controls are independent and agree. Neither is a randomised experiment — publication year is observational and the formula-only arm cannot distinguish memorisation from inference on a near-determining formula — so we claim a strong bound rather than exclusion (§7). The control record is released at `data/modality/formulaonly_control.json`.

5. Forward-verification elucidation

5.1 Method

The inverse direction is the model’s hard, isomer-blind direction; the **forward** direction (structure→spectrum) is its easy, accurate one.² Regioisomers, crucially, have *different* forward-predicted ¹³C shifts. We therefore close a generator–verifier loop:

generate candidate structures (inverse) → **forward-predict** each candidate’s ¹³C spectrum, blind to the observed spectrum → **re-rank** candidates by the distance between predicted and observed ¹³C → return the best match.

This mirrors how a chemist confirms a structure (“if it were that isomer, C-3 would be at ~120 ppm; we observe 135, so it is the other”), exploits the standard principle that verification is easier than generation, and requires no training. **Fig. 4** shows the mechanism on a real benchmark

pair — the picolinamide / nicotinamide regioisomers, which the inverse task cannot separate but whose forward-predicted ^{13}C spectra match the observed one at 0.42 vs 1.30 ppm. We implemented the verifier as independent LLM agents that predict ^{13}C shift lists from SMILES alone, using pure reasoning with no tools; candidates were shuffled and anonymised so that isomers of one target never co-occur. Predicted and observed ^{13}C peak sets were then compared with a symmetric chamfer distance.

For balance — and because conditional precision is 72–84%, not 100% — the verifier also produces clear **false positives** on near-degenerate pairs. On the 2-(nitrophenyl)-2,3-dihydroquinazolin-4(1*H*)-one targets ($\text{C}_{14}\text{H}_{11}\text{N}_3\text{O}_3$), the forward predictor cannot separate the *ortho*- and *meta*-nitrophenyl regioisomers: their predicted ^{13}C lie within its ~ 2 ppm error, and it ranks the wrong (2-nitrophenyl) isomer first by a chamfer of 1.35 ppm versus 1.36 ppm for the true (3-nitrophenyl) one — a 0.01 ppm margin, i.e. effectively a coin flip. This is the precision-loss mechanism of §5.3 made concrete: when candidates are near-degenerate the cheap predictor cannot resolve them, and it is exactly these cases — not generation failures — that a sharper (DFT-level) or 2D-NMR-grounded verifier would need to fix.

Conceptually this is **analogous to NMR-crystallography logic, with an LLM in place of the quantum chemistry** — an analogy, not an equivalence: where DP4/DP4+ rank candidates by GIAO-DFT shifts^{11,12} and NMR crystallography adjudicates polymorphs and connectivity by GIPAW-computed shifts^{13,14}, we rank by the shifts a forward LLM predicts. The shared principle is *verification by forward prediction*; what we do **not** inherit is DFT’s calibrated error model. The trade is deliberate — we forgo the $\approx 1\text{--}2$ ppm accuracy and the rigorous, probabilistic error model of DFT (the very thing that lets DP4 emit a posterior probability) for a predictor that runs in seconds at zero setup, and we quantify below exactly how far that cheaper, uncalibrated verifier carries the inverse problem.

5.2 Result

On the 60 benchmark compounds (126 candidate structures from the solver agents):

Table 6. Forward-verification decomposition on the 60 forward-verify compounds.

	value
generation recall (true structure among candidates)	31%
top-1, solver self-ranking	23%
top-1, forward-verified re-ranking	26%
conditional on recall — self-ranking	14/19 (73%)
conditional on recall — forward-verification	16/19 (84%)
...of which had a single candidate (no choice to make)	6/19
conditional on recall, multi-candidate only — self-ranking	8/13 (62%)
conditional on recall, multi-candidate only — forward-verification	10/13 (77%)

The decomposition is the finding. **When the true structure is among the candidates, for-**

ward verification selects it in 16 of 19 cases (84%) — high in absolute terms and ~18 points above the 66% derangement chance floor of §5.5 (given how few and near-degenerate the recall-positive candidate sets are).

One composition detail must be stated, because it inflates every ranker equally and §5.5 excludes exactly these cases from its own analysis. Of the 19 recall-positive compounds, **6 had a single candidate**, so the verifier faced no choice and any ranker — ours, the solver’s, a coin — scores them by construction. On the **13 compounds where a choice actually existed**, forward-verification selects the true structure in **10/13 (77%)** and the solver’s own self-ranking in **8/13 (62%)**. The asymmetry the paper rests on is unaffected: 31% generation recall against 77% verification precision is the same wall, and the margin over self-ranking is if anything wider on this subset (+15 points against +11). We report both denominators throughout and treat the 13-compound figure as the one that measures verification.

Either way that margin is a two- to three-compound difference at $n \leq 19$ and is **not** statistically significant (McNemar exact $p \approx 0.63$); the load-bearing claim is that verification precision is high in absolute terms while recall binds — not that forward-verification beats self-ranking. The overall top-1 moves only 23%→26% because the binding constraint is **generation recall**: the true structure was never proposed for 41 of 60 compounds, which no re-ranking can repair.

LLM elucidation therefore factorises into two near-independent levers: the **verifier is already strong (84%)**; the **generator (31% recall) is the wall** (Fig. 1).

5.3 Generate-wide: testing the recipe

The decomposition implies a recipe — *generate wide, verify by forward prediction* — a chemistry analog of self-consistency sampling, where many independent generations are pooled and the answer is chosen by agreement rather than from a single pass.²⁴ We tested it directly: ten independent solver agents proposed up to **six regiochemistry-aware candidates per compound**, pooled with the originals, and the 65 new candidates were forward-predicted and re-ranked as before. The measured result:

Table 7. Generate-wide vs original on the same 60 compounds.

	original	generate-wide
recall (true structure among candidates)	31%	41%
forward-verified top-1	26%	30%
verification precision (conditional on recall)	84%	72%

Wide generation lifts recall +10 points (31%→41%, i.e. 19/60→25/60) and exact top-1 +7 points over the self-ranking baseline (23%→30%, 14/60→18/60) — equivalently +4 over the original forward-verified top-1 (26%→30%, 16/60→18/60, Table 7) — on the same 60 compounds (Fig. 6), with no training. One asymmetry in the arm should be stated: of the 217 distinct new candidates the wide pass proposed, only **65 (29%) were forward-predicted** — predicting all of them was beyond the run’s budget — and an unpredicted candidate is assigned an infinite match distance, so it can never be selected. Recall counts every candidate (a true structure need only be *present*), but top-1 can only benefit from the predicted subset. The reported 30% is therefore a **lower bound** on what this recipe achieves with full prediction coverage; it cannot be inflated by the gap. Table 7 regenerates from the released artifacts via `scripts/score_generate_wide.py`. **These**

top-1 differences are directional, not statistically resolved at n=60: the 14/60→18/60 improvement is a four-compound difference, but the stages are **not** nested: going from self-ranking to generate-wide gains seven compounds and loses three, so the paired test is McNemar exact **p=0.34** (b=3, c=7). The intermediate steps are weaker still — self-rank→forward-verify p=0.63 (b=1, c=3) and forward-verify→generate-wide p=0.69 (b=2, c=4). Discordant counts are computed from the released per-compound outcomes by `scripts/ladder_significance.py`. The recall gain is the better-supported effect. We therefore read the ladder as a demonstration that the *mechanism* works and that recall is the movable factor, not as a precise measurement of the size of the top-1 gain. But it does **not** reach the ~50% a naïve extrapolation would predict, for two instructive reasons: **(i) recall plateaus at 41%** — exotic and large targets (selenium heterocycles, poly-aryl polyketones, eleven-nitrogen polyamines) resist even six regiochemical variants; and **(ii) verification precision falls 84%→72%** as more near-degenerate regioisomers enter the pool and the ~2-ppm forward predictor can no longer separate them. This recall/precision tension is the honest ceiling of the training-free approach: the recall-bound diagnosis stands, and closing the gap further requires either sharper verification or 2D-NMR constraints, not merely more candidates.

5.4 Non-LLM verifiers: a deterministic lookup and a learned model

§5.3 suggests an obvious fix: swap the LLM forward-predictor for a non-LLM ^{13}C predictor with a rigorous error model. We test two, both trained on the **same nmrshiftdb2 dump**²⁵ and both dropped into the verifier slot on the **same §5.2 candidate set**, so that only the predictor changes.

The first is the canonical choice: a **HOSE-code**²⁶-**style lookup** (RDKit radial-environment bins, spheres r=4→1 with a hybridisation prior fallback; 31,000 molecules, 332,595 assigned carbons; held-out MAE 3.23 ppm, median 1.73). The second is a small **message-passing GNN** (4 layers, per-carbon ^{13}C regression) trained on the identical dump, which reaches held-out MAE 1.70 ppm (median 1.02) — roughly twice as sharp, and an independently competent predictor.

Table 8. Verifier comparison, conditional on recall (n=19), on the identical §5.2 set.

verifier	top-1 recall	held-out ^{13}C MAE
solver self-ranking	14/19 (73%)	—
deterministic HOSE <i>lookup</i>	14/19 (73%)	3.23 ppm
learned GNN (same data)	16/19 (84%)	1.70 ppm
LLM forward-verifier (§5.2)	16/19 (84%)	—

The lookup does not help: it matches the solver’s own ranking and sits below the LLM verifier. Its diagnosis is coverage — of the 2,244 candidate carbons, only **2%** match a training environment at the most specific sphere (r=4) and **67%** resolve only at coarse spheres (r≤2), because the benchmark’s exotic chemistry (selenium heterocycles, polyaryl ketones, large polyamines) is under-represented in nmrshiftdb2. A coarse environment cannot separate the regioisomers the verifier exists to separate.

The learned model, on the same data, reaches the LLM verifier’s 84% (Fig. S6). **That is a two-compound difference at n=19 and is not statistically resolved** (three compounds gained, one lost against the lookup: McNemar exact p=0.63) — the same standard by which §5.2 declines to separate its own adjacent conditions. Taken together the two rows are *suggestive and directional*: they are consistent with the lookup’s failure being substantially **method** — an

inability to generalise across novel environments — rather than coverage alone, and with a cheap learned predictor being sufficient to match the LLM without compound-specific DFT accuracy or 2D-NMR. They do not establish it, and the coverage reading is not excluded; separating the two would need a substantially larger conditional-on-recall set. What none of these reach is the near-degenerate-regioisomer precision ceiling (§5.3/§5.5), where DFT-level accuracy or orthogonal 2D-NMR constraints remain the genuine fix.

The GNN and the LLM reach 84% on *partly different* compounds (two independent verifiers concurring, not one), and the result is **generalisation, not memorisation**: candidate-training overlap is **0/126** by InChIKey-14 (zero overlap against the entire nmrshiftdb2 database; every true structure and all three GNN-wins absent), Morgan-Tanimoto nearest-neighbour to training is a modest 0.44 median with the wins sitting *at* — not above — the median, and a Y-randomisation control (1,000 derangements) places the real 84% above the 97.5th percentile of the chance distribution (mean 58.6%, 95% range 42.1–73.7%; one-sided $p < 0.05$). The margins are thin (one decisive win at 0.79 ppm, two tight, one 0.08-ppm coin-flip loss), so the +2 is confirmatory of the LLM number rather than a precise quantitative gain. Like §5.6, the learned verifier is a **trained complement**, reported outside the training-free protocol. Artifacts and both leakage checks are released (`scripts/gnn_predict.py`, `data/fverify/gnn_results.txt`, `data/nmrshiftdb/gnn_c13.pt`; `docs/VERIFIER_PROBE.md`).

5.5 Negative control

Permutation negative control (Y-randomisation analog). A re-ranker earns the interpretation that it exploits genuine predicted-vs-observed ^{13}C agreement only if that agreement is destroyed when the pairing is broken. Borrowing the Y-randomisation discipline of QSAR validation, we re-paired — as a *derangement*, so no compound keeps its own spectrum — which observed ^{13}C spectrum each candidate set is scored against, and re-ran the verifier 1,000 times. Conditional-on-recall precision falls from the true **84.2% (16/19)** to a permuted mean of **66.4%** (95% range 52.6–78.9%; one-sided empirical $p = 0.019$, two-sided $p = 0.038$) — the verifier acts on real spectral agreement, not a candidate-list artefact. The honest caveat is the height of the chance floor: because recall-positive compounds carry few and near-identical (regioisomeric) candidates, even a random pairing lands on the correct structure ~66% of the time, so the verifier’s genuine margin over chance is **~18 points**, not the full 84% — real and significant, but to be read against this high baseline.

Confidence calibration (a negative result). We also asked whether the chamfer match-distance doubles as a confidence signal that would support selective prediction. It does not. Ranking the 48 multi-candidate compounds by their chamfer margin (the gap between the best and second-best predicted-vs-observed distance) and answering only the most-confident fraction leaves top-1 essentially flat and non-monotonic with coverage ($\approx 21\%$ overall, neither rising nor falling reliably as the threshold tightens). Compounds with a single proposed candidate have no margin at all and must be excluded; an earlier analysis that retained them produced a spurious “improvement” that was entirely an artefact of those trivial cases. We therefore do **not** claim the verifier distance as a calibrated abstention gauge — a sharper, DFT/2D-NMR-grounded confidence estimate (§5.4) remains the route to reliable selective prediction.

5.6 Is the recall wall task-intrinsic? A trained-generator probe

The ceiling above is a property of *training-free LLM elicitation*; it leaves open whether the ~41% recall plateau is intrinsic to 1D-data elucidation or specific to how LLMs are elicited. We test this with a complementary probe — a small (~16M-parameter) $^1\text{H}/^{13}\text{C}\rightarrow\text{SMILES}$ transformer (simulated-spectra pretraining, then fine-tuned on IReXP; ensemble of four), held **deliberately separate from the training-free protocol**. Its candidates are filtered to the given molecular formula, pooled with Claude’s, and re-ranked by the same verifiers.

The generator supplies candidates enumeration cannot. On the 194-compound benchmark the true structure enters the candidate pool for **54.1%** of compounds (versus 41.8% for scaffold enumeration and 33.5% for Claude alone). Crucially, where enumeration’s near-degenerate regioisomers *collapse* the deterministic HOSE verifier (top-1 28.4% $\rightarrow$ 16.0%, the §5.3 precision-loss mechanism), the generator’s formula-correct, ^{13}C -separable candidates **convert**: HOSE top-1 rises **28.4% $\rightarrow$ 35.1%** (McNemar exact $p=0.015$; +6.7 points, 95% CI [+2.1, +11.9]; Fig. S5). With the stronger LLM forward-prediction verifier on the 60 forward-verify compounds, recall rises 41% $\rightarrow$ **56%**; the corresponding forward-verified top-1 (measured at **41%**, precision 72% $\rightarrow$ 73%) was obtained with blind in-house prediction agents and is reported as **provisional** pending deposit of those forward-prediction outputs and a re-run under the pinned model snapshot (the released `forward_verify_gen.py` reproduces the recall arm and the deterministic HOSE re-rank; the LLM forward-prediction JSONs are not yet in the bundle).

Two controls guard against memorisation, both reproducible from the released bundle (`contrib/generator_probe`): (i) the fine-tuning split removes every benchmark InChIKey-14 ($\text{train}\cap\text{benchmark}=0$, $\text{val}\cap\text{benchmark}=0$), and **none of the 40 newly-recovered compounds appear in either training stage** (0/40 in simulated pretraining, 0/40 in fine-tuning); (ii) the simulated-pretrained model alone recovers **0/248** benchmark structures zero-shot, rising to 25% only after IReXP fine-tuning. The IReXP data — not the architecture — is the active ingredient, and the recall ceiling is therefore **elicitation-specific, not task-intrinsic**. We report this as a probe, not part of the headline training-free protocol; it answers whether the wall is breakable, not whether to abandon training-free methods. (The released public split `irexp_release/train` does *not* hold out the benchmark — it overlaps it by 117/200 InChIKey-14 — so downstream users must de-leak as in `build_exp_manifest.py`; see Data and code availability.)

6. Discussion

The frontier LLMs we tested (the Claude family) are neither “solving structure elucidation” on realistic 1D data nor failing at it. They are reliable **scaffold-level** elucidators (best-candidate Tanimoto ≥ 0.45 for 56% of compounds, 73% of simple ones; mean best Tanimoto 0.59) and good **verifiers** (84% conditional on recall). Exact top-1 is throttled by candidate recall and by the regiochemical underdetermination intrinsic to 1D NMR. The contribution of this paper is therefore best read as a **diagnosis with a bounded, training-free improvement attached**, not as a method that solves the task. The diagnostic result is the durable take-away: across models, domains, and verifiers, the wall is **generation recall**, not verification. The accompanying improvement is real but deliberately modest — forward verification with wide generation lifts top-1 from 23% to 30%, and we show by direct measurement (§5.3) that this stays *below its own ceiling* (recall plateaus at 41%; verification precision falls to 72% as near-degenerate regioisomers accumulate), while the match distance fails as a confidence gauge for selective prediction (§5.5, a reported null result).

This reframes the engineering problem. Rather than training a bespoke spectra→structure model — a target the frontier already meets at the scaffold level and that ages out each model generation — the durable levers are (i) **open, hard, honestly scored benchmarks** that expose specific failure modes (here, regiochemistry and recall), and (ii) **inference-time scaffolding** (decoupled solving, generate-and-verify) that needs no training and improves with each model. IReXp’s IR modality and 2D-NMR-ready provenance position it for the obvious next step: supplying the HMBC/COSY/NOESY constraints that would attack regiochemistry at the source. Our trained-generator probe (§5.6) sharpens the point: the recall wall is *elicitation-specific, not task-intrinsic* — a small generator fine-tuned on IReXp lifts it (recall 41→56%) — and because that gain comes entirely from the released data (0→25% with fine-tuning, 0/248 without), it is the **open dataset**, not a bespoke architecture, that does the work when training does help.

Protocol is a lever of the same order as capability. How a problem is posed and verified — bounded, reset contexts with tool access; generate-and-verify — moves accuracy enough that it must be reported alongside any capability claim. We do *not* claim it dominates capability: on the fixed 24-compound subset the model axis spans 0% to 45% top-1 (Table 3), a wider range than the 5%→15% protocol effect in §4.3, and the two are measured on different sets. The point is that a number quoted without its protocol is uninterpretable, which echoes a recurring lesson across cheminformatics: careful pipeline and protocol design competes with the fashionable component. The same pattern appears in molecular property prediction, where a recent systematic benchmark reports that learned, pretrained molecular embeddings rarely outperform classical ECFP fingerprints once evaluation is controlled²⁷ — the modelling fashion underperforms the well-engineered baseline. We read our within-compound control (§4.3) and forward-verification recipe (§5) in the same light: the durable gains in LLM elucidation come from honest benchmarking and inference-time scaffolding rather than from waiting on the next, larger model. We draw this parallel narrowly and claim no formal equivalence between the two settings.

That the bottleneck reproduces inside a single application domain (§4.5: 26% top-1 on battery-electrolyte chemistry, n=46, against 28% overall) is consistent with it being structural rather than incidental — and it makes the forward-verification recipe directly relevant to the magnetic-resonance workflows (electrolyte-decomposition assignment, NMR crystallography) where computing a candidate’s spectrum to confirm it is already standard practice.

For practitioners, two operational findings transfer immediately: solve each problem in bounded, frequently-reset contexts with tool access (5%→15%, 1/20→3/20, directional at n=20, McNemar $p \geq 0.25$, §4.3), and use forward-predicted-vs-observed ¹³C agreement to *re-rank* candidates, not to decide whether to trust the winner. §5.5 is explicit on this: match distance is a strong relative re-ranker but failed as an absolute confidence gauge in our selective-prediction test, so it does not support abstention thresholds.

7. Limitations

Several limitations of earlier drafts are now resolved with controlled experiments; we state what remains plainly.

Resolved. **Extraction noise:** an RDKit self-consistency audit (¹³C peak count vs symmetry-unique carbons, formula, SELFIES round-trip) finds **57/60 ground truths spectrally clean**, and the metrics are unchanged on that clean subset (self-ranking top-1 24% vs 23%, recall 33% vs 31% on the

n=60 forward-verify set) — the conclusions are not driven by scraping artefacts. **Verifier precision / abstention:** the generate-wide experiment (§5.3) quantifies the recall/precision tension directly — verification precision is 72–84% conditional on recall and degrades as near-degenerate regioisomers accumulate, so forward-match distance is a strong *re-ranker* but a soft *confidence* gauge. We further tested non-LLM verifiers (§5.4): a nmrshiftdb2-trained HOSE-code *lookup* does **not** beat the LLM verifier (73% vs 84% conditional), while a *learned* GNN on the same data reaches 84% — a two-compound difference at n=19 that is directional only (p=0.63), so it is suggestive that the deterministic failure was method rather than coverage alone, not a demonstration of it. What remains beyond all of them is the near-degenerate-regiochemistry precision ceiling, where DFT-level accuracy or 2D-NMR constraints are still the genuine fix. **Projection:** §5.2’s extrapolation is replaced by the **measured** §5.3 result (top-1 30%, recall 41%) — and is honestly below the optimistic estimate.

Independence checks. Scoring throughout is mechanical RDKit (not LLM-judged), and all solver/verifier runs were transcript-audited *at generation time* for zero web and zero ground-truth access; the committed artefacts are the parsed per-compound predictions, and the underlying transcripts are available on request. To confirm that the forward-verifier’s measured advantage reflects real predicted-vs-observed spectral agreement rather than leakage or a candidate-list artefact, we ran a permutation negative control (Y-randomisation analog, §5.5): re-pairing — as a derangement — which observed ^{13}C spectrum each candidate set is scored against, over 1,000 permutations, collapses conditional-on-recall precision from 84.2% to a chance mean of 66.4% (two-sided p=0.038) — the signal weakens to the chance floor under label-shuffle, as a correctly isolated blind evaluation requires. A second Claude model (Sonnet) re-solved a 12-compound subset under the identical generate-wide protocol and was comparably recall-bound (recall 33% vs Opus 41% on those compounds), consistent with the recall bound being a property of the task rather than of one model. As this is within the Claude family, it speaks to model-instance robustness, not cross-vendor generality.

Remaining. (i) **Pretraining contamination is bounded, not excluded.** Every benchmark compound was mined from open-access literature, so a frontier model may have encountered it in training. The formula-only control (§4.6) bounds how much of the headline number pure recall can explain — masking the spectra drops top-1 from 23% to 5%, perfectly nested, McNemar exact p=0.001 — and the three formula-only successes are all compounds whose composition is close to determining. The second control (§4.6) resolves the source publication year for all 194 compounds and finds accuracy flat in it (r=−0.007), with the size-adjusted older-versus-newer difference bounded at −5.1 points, 95% CI [−17.2, +7.0].

Two things these controls do not do. Neither is randomised: publication year is observational, and the formula-only arm cannot separate memorisation from inference on a near-determining formula. And we do not anchor the recency test to a *known* training cutoff — the subscription harness does not disclose one (§8), so we test recency as a continuous proxy rather than partitioning at a cutoff date. A replication restricted to compounds published after a disclosed cutoff would settle what these bound, and remains open.

(ii) **Human audit — prepared but not yet run.** This covers two things we have not validated by hand: the elucidation and forward-prediction outputs, and the *recall* side of dataset extraction (whether the parser found every IR string in every source paper). Transcription fidelity of the records we do hold is measured and high (§2.3), but record-level recall is not, and only reading papers can settle it. Solver and verifier are both LLMs, so the one validation we cannot perform ourselves is an expert-chemist review of a sample of elucidations and forward

predictions. We have therefore **built and frozen** a blinded, pre-registered audit package — a difficulty-stratified 30-compound sample (9 recall-positive), rendered candidate structures, a separate withheld answer key, and a mechanical scoring sheet — released at `data/audit/` with its protocol at `docs/EXPERT_AUDIT_PROTOCOL.md` (regenerable via `scripts/make_audit_sample.py`), designed to test the two load-bearing claims (misses are predominantly consistent regioisomers, §4; forward verification is a trustworthy re-ranker, §5). **We have not yet run the panel**; until expert results are in, those claims should be read as machine-validated (RDKit InChIKey) but not yet human-validated.

(iii) Single vendor and an underpowered cross-model comparison. Every number comes from one vendor’s models. The headline (28.4% top-1) is a single frontier model (Claude Opus), and the cross-model evidence (§4.4) is four **Claude-family** models on a shared **24-compound** subset. As §4.4 sets out, that comparison is underpowered: the ranking is robust (strictly nested, weakest floored at 0%) but no adjacent gap is established at $n=24$. Quantitative claims should therefore be read as holding **for the Claude family**, and a true cross-**vendor** sweep (GPT-, Gemini-class, open-weight models) — the test of whether the pattern is a property of the task rather than one lineage — was forgone because it needs paid API access incompatible with our zero-cost protocol; we flag it as the most important external-validity experiment left open.

(iv) Domain-subset scope: the battery-electrolyte case study (§4.5) comprises *literature compounds bearing electrolyte-relevant functional chemistry* drawn from the open corpus — not operando or in-cell decomposition spectra — so it demonstrates functional-class transfer of the elucidation bottleneck, not direct assignment of authentic interphase/degradation products, which would require a dedicated operando set.

(v) Constitution-only scoring: correctness is judged on InChIKey connectivity, so a correct-constitution / wrong-stereochemistry prediction is counted correct. Scoring the full InChIKey instead gives 21.1% top-1 (§3), 7.3 points lower, so the headline should be read as an upper bound on full-stereochemistry accuracy. We keep constitution as the headline because 1D NMR/IR rarely determines absolute configuration — the 10.3% (20/194) of targets with a defined stereocentre would be penalised for information the prompt never carried — but a stereochemistry-sensitive benchmark would need 2D/chiroptical data.

(vi) Single-sample scoring: each headline compound is scored from one solver prediction set (one decoupled run), so reported top-1/recall carry no run-to-run (LLM-sampling) variance estimate — the bootstrap CIs reflect compound sampling only. The generate-wide experiment (§5.3) pools ten independent generation passes (recall 31%→41%), indicating generator stochasticity that single-pass scoring understates.

(vii) Chemical-space coverage: the benchmark is drawn from open-access organic methodology and total-synthesis literature, covering small-to-medium drug-like and synthetic-organic space; it is **not** representative of organometallic/coordination compounds, large biomolecules (peptides, oligonucleotides, oligosaccharides), or stereochemistry-heavy targets, and our results should not be extrapolated to them.

8. Methods

Mining and resolution. PMC-OA full text was fetched from `s3://pmc-oa-opendata`; records were extracted with a deterministic parser and resolved with OPSIN, RDKit, and SELFIES, with

an optional cached PubChem fallback. Durable content-keyed caches make resolution additive and restart-safe.

Benchmark and agents. Problems were sampled from `irexp_resolved`, stratified by RDKit ring analysis, with cross-round de-duplication by InChIKey. The battery-electrolyte subset (§4.5) was drawn from the same corpus by SMARTS substructure filters for six electrolyte functional classes (carbonate, sulfonyl/sulfonate, nitrile, $\text{sp}^3\text{-C-F}$, phosphoryl, glyme/oligoether), balanced to eight compounds per class (48 curated; 46 scored after two yielded no parseable candidate), excluding every compound used elsewhere, and J-enriched and spectrally validated identically to the main rounds. Solver and forward-prediction agents were independent Claude-Opus sub-agents invoked under a consumer subscription; agents were instructed closed-book and audited via automated transcript search for zero web/answer access. Scoring used RDKit InChIKey-connectivity matching and Morgan(2, 2048) Tanimoto; forward-verification used a symmetric chamfer distance over ^{13}C peak sets. The core protocol trains no model and uses no paid API; the two trained probes — the §5.6 generator and the §5.4 learned ^{13}C verifier — are the only separately-trained exceptions, each reported as a complement and fenced from the headline results.

Models and versions. All experiments used Anthropic Claude models via the consumer subscription (claude.ai). The §4.4 comparison spanned, in capability order, Claude Haiku, Claude Sonnet, Claude Opus, and Claude Fable 5; the headline benchmark (§4.1) and forward-verification (§5) used Claude Opus. `docs/MODELS.md` records, for each experiment, the model used, its data directory, the collection date, the harness and tool access, and the scoring code path; all LLM predictions were collected between 2026-06-09 and 2026-06-11 (UTC). Decoding parameters are not exposed by the subscription harness and were neither set nor recorded, so re-running reproduces the protocol distributionally rather than exactly; the dated model snapshot identifiers are listed there as the outstanding items to be pinned on submission.

Reproducibility. Every round is frozen: questions, ground-truth answers, per-agent raw outputs, predictions, and scorer outputs are released, and the sampler, scorer, and forward-verification harness are scripted end-to-end.

Figures (main text)

- **Fig. 1** (`docs/figures/fig_wall.png`) — the diagnosis in one glance, as a part-to-whole bar of the 60 forward-verify compounds: 16 verified top-1, 3 recalled but mis-ranked, 41 never proposed — *the wall*, 68% of the bar. Generation recall is 19/60 (31%); conditional verification precision is 16/19 (84%); their product is the 26% end-to-end top-1. The two rates have different denominators and are not differenced.
- **Fig. 2** (`docs/figures/fig1_difficulty.png`) — top-1 and recovered accuracy on IRSpectra-Bench by difficulty (all / simple / complex, $n=194$) with bootstrap 95% CIs.
- **Fig. 3** (`docs/figures/fig5_models.png`) — four-model comparison on a 24-compound subset: Fable 5 45% > Opus 25% > Sonnet 20% > Haiku 0% top-1 (strictly nested); capability-sensitive but underpowered to separate adjacent models at $n=24$.
- **Fig. 4** (`docs/figures/fig_mechanism.png`) — forward-verification on a real benchmark regioisomer pair (picolinamide vs nicotinamide): forward-predicted ^{13}C matches the true isomer (chamfer 0.42 vs 1.30 ppm), an analog of NMR-crystallography.
- **Fig. 5** (`docs/figures/fig_contamination.png`) — two contamination controls.
 - (a) Removing the spectra: formula-only reaches 3/60 against 14/60 with IR + ^1H + ^{13}C on

the same compounds, nested (11 solved only with the spectra, none only without). (b) Accuracy against source publication year (n=194, Wilson 95% CIs): flat, point-biserial $r = -0.007$.

- **Fig. 6** (docs/figures/fig3_method.png) — forward-verification inference ladder on the same 60 compounds: solver self-ranking $\rightarrow$ + forward-verify $\rightarrow$ + generate-wide (23%/26%/30% top-1).

Supporting Information figures

These are supplied as a separate Electronic Supplementary Information document (docs/paper_eSI.pdf, built by scripts/build_pdf.py), as RSC requires; they are listed here for reference.

- **Fig. S1** (docs/figures/fig0_overview.png) — study design: open multimodal data (IRexp) $\rightarrow$ blind, complexity-stratified benchmark $\rightarrow$ decoupled blind solving $\rightarrow$ forward-verification re-ranking; training-free core pipeline.
- **Fig. S2** (docs/figures/fig4_dataset.png) — IRexp composition: IR records $\rightarrow$ NMR-paired $\rightarrow$ structure-linked $\rightarrow$ full IR+ ^1H + ^{13}C +structure quadruples.
- **Fig. S3** (docs/figures/fig2_size.png) — accuracy vs molecular size (heavy-atom bucket); the monotonic 60% $\rightarrow$ 7% top-1 gradient.
- **Fig. S4** (docs/figures/fig6_electrolyte.png) — top-1 and recovered accuracy on IRSpectra-Bench-Electrolyte by battery-electrolyte chemical class (n=46): $\text{sp}^3\text{-C-F}$ easiest (50%), sulfonyl and nitrile hardest (12%); overall 26%/28%.
- **Fig. S5** (docs/figures/fig_generator_probe.png) — trained-generator probe (§5.6; a complement, not part of the training-free protocol): candidate recall and deterministic-HOSE top-1 on the 194-compound benchmark for Claude / + scaffold enumeration / + trained generator. Enumeration’s near-degenerate isomers collapse the verifier (28.4 $\rightarrow$ 16.0%) while the generator’s formula-correct candidates convert (28.4 $\rightarrow$ 35.1%).
- **Fig. S6** (docs/figures/fig_verifier.png) — learned-verifier probe (§5.4; a complement, not part of the training-free protocol). (A) Conditional-on-recall top-1 (n=19) for the four verifiers: a GNN trained on the same nmrshiftdb2 data as the HOSE lookup recovers the LLM verifier’s 84% that the lookup (73%) could not. (B) Why: held-out ^{13}C MAE — the learned model is $\sim 2\times$ sharper (1.70 vs 3.23 ppm).

Data and code availability

All data and code are released in the project repository: IRexp and the irexp_resolved split (data/irexp/, data/irexp_resolved/); the benchmark rounds and within-compound control (data/benchmark*/, scored by scripts/benchmark_v2.py); the battery-electrolyte case-study subset (data/benchmark_electrolyte/, built by scripts/build_electrolyte_bench.py, scored by scripts/score_electrolyte.py); the ground-truth integrity audit (scripts/validate_benchmark.py, data/benchmark*/clean_qids.json); the forward-verification and generate-wide experiments (data/fverify/, data/gw/, data/fverify2/, scripts/forward_verify.py); the deterministic HOSE-code verifier ablation (scripts/hose_predict.py, data/fverify/hose_results.txt); and the dataset-mining pipeline (spectro_scraper/). Companion technical notes: docs/BENCHMARK.md and docs/FORWARD_VERIFY.md. The verifier negative-control and selective-prediction analyses (scripts/verifier_diagnostics.py), the recall-headroom and scaffold-enumeration study (scripts/analyze_recall_headroom.py, scripts/enumerate_isomers.py, scripts/closing_the_gap.py),

the modality-ablation harness (`scripts/modality_ablation.py`, `docs/MODALITY_ABLATION.md`), and the blinded expert-audit package (`data/audit/`, `docs/EXPERT_AUDIT_PROTOCOL.md`) are released likewise. The §5.6 trained-generator probe ships as a self-contained bundle (`contrib/generator_probe/`): generator candidates, the de-leaked split with its InChIKey-14 manifest, and the verification scripts (`scripts/closing_the_gap_gen.py`, `scripts/forward_verify_gen.py`, `scripts/verify_leakage_exact40.py`); model checkpoints are deposited on Zenodo. Downstream note: the public `irexp_release/train` split does **not** hold out IRSpectra-Bench (117/200 InChIKey-14 overlap); de-leak with `contrib/generator_probe/build_exp_manifest.py` before training models that will be evaluated on the benchmark. The §5.4 learned-verifier probe ships its full reproducer — `scripts/gnn_predict.py` (extract/train/score/control), the trained model `data/nmrshiftdb/gnn_c13.pt`, per-compound results and both leakage checks (`data/fverify/gnn_results.txt`), and the write-up `docs/VERIFIER_PROBE.md`; the GNN trains on the same `nmrshiftdb2` dump as the §5.4 HOSE lookup. **Archival deposit:** a complete frozen snapshot (dataset, benchmark, answer keys, predictions, scripts, figure-regeneration, and the expert-audit package) will be archived on Zenodo under DOI [TODO: 10.5281/zenodo.XXXXXXX — mint on submission]; the GitHub repository is the development mirror and the Zenodo record the citable version of record. (The held-out answer keys `data/audit/key.jsonl` and `data/modality/key.json` are deposited but flagged “withhold from blinded reviewers”.)

Licensing and attribution

IRexp is derived from two open-access source pools and redistributed under terms compatible with each (see `data/NOTICE`). **(a) Redistribution:** we release only *extracted numeric data* — IR band lists, $^1\text{H}/^{13}\text{C}$ shift lists, and resolved structures (SMILES/SELFIES/InChIKey) — plus each record’s source DOI/accession, not source full text, figures, or PDFs. **(b) Two separable pools:** 119,345 records derive from the PMC Open-Access Subset (**CC-BY-4.0**) and 1,888 from the Chemotion RADAR4Chem FT-IR deposit (**CC-BY-SA-4.0**); the pools are kept separate and individually labelled, so users may take the CC-BY pool alone, while any combined or Chemotion-derived release carries CC-BY-SA-4.0 to honour the ShareAlike term. Code is released under the MIT License. **(c) Attribution:** re-users must cite this dataset (Zenodo DOI above) and attribute the original publications via each record’s `source_doi`.

Author Contributions

I.Y.: conceptualization, methodology, software, formal analysis, investigation, data curation, visualization, writing — original draft. **R.S.:** methodology, software, formal analysis, investigation, validation (trained-generator and learned-verifier probes, §5.4 and §5.6), writing — review and editing. **R.A.V.-H.:** conceptualization, methodology, supervision, writing — review and editing.

Conflicts of Interest

The authors declare no competing interests.

Use of AI tools

This work studies a large language model, and LLMs were also used as instruments and as writing aids; we state both roles explicitly. **As an object of study:** all reported elucidation, forward-prediction and verification results were produced by Claude models invoked under the protocol of §3 and §8 — these are the measurements the paper reports, not assistance in producing it. **As**

a writing aid: the authors used an LLM-based coding assistant for figure generation, analysis scripting, manuscript copy-editing, and for the internal review pass that produced several of the corrections recorded in the repository history. **No text, number, figure or citation in this manuscript was accepted without author verification against the released data and code;** every quantitative claim regenerates from the scripts in `scripts/`. The authors take full responsibility for the content.

Acknowledgements


(To be completed before submission: funding sources, compute/infrastructure, and any individual acknowledgements. — AUTHORS)

References

- 1 E. Chacko, R. Sondhi, A. Praveen, K. L. Luska and R. A. Vargas-Hernández, *ChemRxiv*, 2024, preprint, DOI: [10.26434/chemrxiv-2024-37v2j](https://doi.org/10.26434/chemrxiv-2024-37v2j).
- 2 D. Kamber, *Making Claude a chemist: How Claude performs on NMR prediction and structure elucidation*, Anthropic, 2026.
- 3 F. Hu, M. S. Chen, G. M. Rotskoff, M. W. Kanan and T. E. Markland, *arXiv*, 2024, preprint, DOI: [10.48550/arXiv.2408.08284](https://doi.org/10.48550/arXiv.2408.08284).
- 4 L. Yang, Z. Yang, J. Xie, Y. Wang, B. Gao, T. Fu, X. Wei, J. Sun, J. Wu, C. He, Y. Li and Q. Gu, *arXiv*, 2026, preprint, DOI: [10.48550/arXiv.2602.10158](https://doi.org/10.48550/arXiv.2602.10158).
- 5 F. Ottomano, Y. Li and A. M. Ganose, *arXiv*, 2025, preprint, DOI: [10.48550/arXiv.2512.19733](https://doi.org/10.48550/arXiv.2512.19733).
- 6 Y. Su, J. Chen, Z. Jiang, Z. Zhong, L. Wang, Q. Liu and Z. Zhang, *arXiv*, 2025, preprint, DOI: [10.48550/arXiv.2508.08441](https://doi.org/10.48550/arXiv.2508.08441).
- 7 S. Shen, J. Xie, Z. Yang, A. Zhang, S. Sun, B. Gao, T. Fu, B. Qi and Y. Li, *arXiv*, 2025, preprint, DOI: [10.48550/arXiv.2509.21861](https://doi.org/10.48550/arXiv.2509.21861).
- 8 X. Zhuang, B. Wu, J. Cui, K. Feng, X. Li, H. Xing, K. Ding, Q. Zhang and H. Chen, *arXiv*, 2025, preprint, DOI: [10.48550/arXiv.2506.23056](https://doi.org/10.48550/arXiv.2506.23056).
- 9 K. Guo, B. Nan, Y. Zhou, T. Guo, Z. Guo, M. Surve, Z. Liang, N. V. Chawla, O. Wiest and X. Zhang, in *Advances in neural information processing systems 37 (NeurIPS 2024), datasets and benchmarks track*, 2024.
- 10 M. Alberts, T. Laino and A. C. Vaucher, *Commun. Chem.*, 2024, **7**, 268.
- 11 S. G. Smith and J. M. Goodman, *J. Am. Chem. Soc.*, 2010, **132**, 12946–12959.
- 12 N. Grimblat, M. M. Zanardi and A. M. Sarotti, *J. Org. Chem.*, 2015, **80**, 12526–12534.
- 13 C. J. Pickard and F. Mauri, *Phys. Rev. B*, 2001, **63**, 245101.
- 14 S. E. Ashbrook and D. McKay, *Chem. Commun.*, 2016, **52**, 7186–7204.
- 15 National Institute of Standards and Technology, NIST chemistry WebBook, SRD 69, <https://webbook.nist.gov>, (accessed 1 January 2026).

- 16 Chemotion Repository, *RADAR4Chem: FT-IR spectroscopy data*, Karlsruhe Institute of Technology, 2024.
- 17 National Institute of Advanced Industrial Science and Technology (AIST), Spectral database for organic compounds (SDBS), <https://sdb.db.aist.go.jp>, (accessed 1 January 2026).
- 18 National Library of Medicine, NCBI PMC open access subset, <https://www.ncbi.nlm.nih.gov/pmc/tools/openftlist/>, (accessed 1 January 2026).
- 19 D. M. Lowe, P. T. Corbett, P. Murray-Rust and R. C. Glen, *J. Chem. Inf. Model.*, 2011, **51**, 739–753.
- 20 G. Landrum and others, RDKit <https://www.rdkit.org>.
- 21 M. Krenn, F. Häse, A. Nigam, P. Friederich and A. Aspuru-Guzik, *Mach. Learn.: Sci. Technol.*, 2020, **1**, 045024.
- 22 S. Kim, J. Chen, T. Cheng and others, *Nucleic Acids Res.*, 2023, **51**, D1373–D1380.
- 23 D. Rogers and M. Hahn, *J. Chem. Inf. Model.*, 2010, **50**, 742–754.
- 24 X. Wang, J. Wei, D. Schuurmans, Q. Le, E. Chi, S. Narang, A. Chowdhery and D. Zhou, in *International conference on learning representations (ICLR)*, 2023.
- 25 S. Kuhn and N. E. Schlörer, *Magn. Reson. Chem.*, 2015, **53**, 582–589.
- 26 W. Bremser, *Anal. Chim. Acta*, 1978, **103**, 355–365.
- 27 M. Praski, J. Adamczyk and W. Czech, *arXiv*, 2025, preprint, DOI: 10.48550/arXiv.2508.06199.

Table of contents entry



28% top-1 on blind, real spectra — recall (31%), not verification (84%), is the wall

A frontier LLM recovers the correct molecular constitution from real, blind IR + ^1H + ^{13}C literature spectra for 28% of 194 compounds. The bottleneck is candidate *recall*, not verification: forward-predicting ^{13}C and re-ranking selects the true structure 84% of the time it is proposed — but it is proposed only 31% of the time.

Figure plates




Fig. 1: The diagnosis as a single part-to-whole bar of the 60 forward-verify compounds. The true structure is verified top-1 for 16 (green), recalled but mis-ranked for 3 (vermilion), and never proposed for 41 (grey) — *the wall*, 68% of the bar. Forward-verification recovers $16/60 = 26\%$ exact top-1 end-to-end: the model proposes the true structure for only $19/60 = 31\%$ of compounds and, of those, verifies 84% (77% on the 13 of 19 where more than one candidate existed; §5.2). Recall, not verification, is the wall.




Fig. 2: Top-1 and recovered accuracy on IRSpectra-Bench by difficulty (all / simple / complex, n=194) with bootstrap 95% CIs.




Fig. 3: Four-model comparison on a 24-compound subset: Fable 5 45% > Opus 25% > Sonnet 20% > Haiku 0% top-1 (strictly nested; underpowered to separate adjacent models at $n=24$).




Fig. 4: Forward-verification on a real benchmark regioisomer pair (picolinamide vs nicotinamide): forward-predicted ^{13}C matches the true isomer (chamfer 0.42 vs 1.30 ppm) — an analog of NMR-crystallography.




Fig. 5: Two contamination controls. (a) Removing the spectra: with only the molecular formula the solver reaches 3/60, against 14/60 with IR + ^1H + ^{13}C on the same compounds; the outcomes are nested, with 11 compounds solved only with the spectra and none only without. (b) Accuracy against the publication year of the source paper ($n=194$, Wilson 95% CIs): flat, with a point-biserial correlation of -0.007 . Recall from pretraining would predict a decline with recency; none is observed.




Fig. 6: Forward-verification inference ladder on the same 60 compounds: solver self-ranking $\rightarrow$ + forward-verify $\rightarrow$ + generate-wide (23%/26%/30% top-1).