

# Recall, not verification, is the bottleneck when frontier LLMs elucidate molecular structures from real spectra

Ilkham Yabbarov<sup>a\*</sup>, Rudra Sondhi<sup>a</sup>, Rodrigo A. Vargas-Hernández<sup>a</sup>

<sup>a</sup> Department of Chemistry and Chemical Biology, McMaster University, Hamilton, Ontario L8S 4L8, Canada.

\*E-mail: yabbaroi@mcmaster.ca

Given the molecular formula together with the infrared band list and  $^1\text{H}/^{13}\text{C}$  shift lists exactly as reported in an open-access paper, how often does a frontier large language model (here, Claude) recover the correct molecular *constitution*? We find 28% (top-1,  $n=194$ ; 95% CI 22–35) — far below the near-100% implied by curated demonstrations, and 15% once the benchmark’s deliberate 50/50 difficulty balance is reweighted to the corpus it came from (§3).

The bottleneck lies in the model’s *proposal* rather than its judgment. The true structure is proposed for only 34% of compounds; where it is, forward-verification — predicting each candidate’s  $^{13}\text{C}$  spectrum and re-ranking by agreement with the observed one — selects it 89% of the time (58/65), 81% where a ranker had a choice, and 71% under the same corpus weights against a reweighted recall of 20%. **Recall, not verification, is the wall.**

That split needs no new data to apply to published work: a system reporting top-1 and top- $k$  has already bounded its own recall and its own ranking loss. Read that way, three groups have run this control without reporting it — hold the system, the scoring criterion and the candidate budget fixed, change only the data, and recall carries 68–83% of the resulting collapse (§6). Ranking degrades too; it is simply the smaller term where spectra are real and heterogeneous, and the larger one on simulated spectra and single curated instrument libraries.

We release *IRexp*, the largest openly redistributable collection of *experimental* infrared band lists (121,233 records, a third structure-linked); *IRSpectra-Bench*, a blind, mechanically scored benchmark of 194 compounds; and a training-free *forward-verification* recipe, which lifts top-1 from 28% to 30% where generating wider moves recall from 32% to 42% — both directional rather than statistically resolved (§5.2, §5.3). Masking the spectra drops top-1 from 23% to 5% on the same compounds (§4.6). All data, predictions and code are released.

## 1. Introduction


Determining a molecule’s structure from its spectra is a central, time-consuming task in synthetic and analytical chemistry. Machine learning attacks it with encoder–decoder models trained on paired corpora; *Spectro*, for one, learns  $^1\text{H}/^{13}\text{C}/\text{IR} \rightarrow \text{SELFIES}$  from 6,833 molecules<sup>1</sup>. General-purpose LLMs are now reported to do it off-the-shelf: a non-peer-reviewed 2026 industrial white paper found that Claude Opus matched or beat commercial NMR-prediction software in the forward direction (structure  $\rightarrow$  spectrum,  $\pm 0.08$  ppm  $^1\text{H}$ ) and “recovered all eight simpler structures on every attempt” in the inverse<sup>2</sup>. We treat those numbers as a motivating claim to test against peer-reviewed benchmarks, not as an established baseline.

That evaluation is narrow: 15 inverse problems on curated single-ring or two-fragment molecules, NMR only, seven of them given the *starting-material structure* as a hint, and “recovery” scored leniently over three runs and three ranked candidates. The question a chemist actually faces —

take an arbitrary experimental spectrum from a paper and recover the structure — is left open. It needs a large, diverse, real benchmark; blind, reproducible scoring; and honest accounting for the methodological choices that inflate or deflate the apparent score.

We provide all three. *IRexp* (§2) is, to our knowledge, the largest openly *redistributable* collection of experimental IR *band lists* (121,233 records; 43,060 structure-linked; 33,201 IR +  $^1\text{H}$  +  $^{13}\text{C}$  + structure); the superlative is scoped to that object type, since view-only libraries such as SDBS and commercial libraries hold more structure-linked *spectra* but cannot be redistributed (§2.1). On it, *IRSpectra-Bench* (§3) is a blind, mechanically scored, complexity-stratified evaluation of frontier-LLM elucidation on real spectra. It is not the first blind, mechanically scored one — MolQuest scores 530 post-2025 compounds by exact canonical SMILES<sup>3</sup> — and what is ours is the stratification, the multimodal IR +  $^1\text{H}$  +  $^{13}\text{C}$  input, and the decomposition it supports: measured in depth on Claude, replicated across three other model families (§4.7), and isolating a large solver-methodology effect with a within-compound control (§4). Its ground truth is literature structures resolved deterministically (OPSIN/RDKit) and checked mechanically against the source articles (560/560 bands on a seed-fixed sample, §2.3; the *expert-chemist* review is prepared but not yet run, §8); no LLM curates labels or scores predictions.

Forward-verification elucidation (§5) turns the model’s *strong* direction (forward prediction) against its *weak* one (inverse regiochemistry); the loop is prior art (§1.1), the decomposition it enables is ours, run training-free over every benchmark compound. The finding is a sharp asymmetry: given a candidate set containing it, the model *verifies* the correct structure 89% of the time, yet *proposes* it for only 34% of compounds, so generation is what binds the result (Fig. 1). The gain is bounded (top-1 28%  $\rightarrow$  30% on  $n=194$ ; 23%  $\rightarrow$  30% on the 60-compound arm, the two pre-registered controlled rounds on which wide generation was also tested) and, by our own measurements, cannot exceed a recall/precision ceiling without sharper verification or 2D-NMR data.



**Fig. 1.** All 194 benchmark compounds as a single part-to-whole bar. The true structure is verified top-1 for 58 (green), recalled but mis-ranked for 7 (vermilion), and never proposed for 129 (grey): *the wall*, 66% of the bar. Generation recall (the fraction of compounds whose true structure appears anywhere in the solver’s candidate list) is  $65/194 = 34\%$ . Of those 65, forward-verification then ranks 58 first: verification precision (conditional on recall) is  $58/65 = 89\%$ . Their product is the 30% end-to-end top-1. The two rates have different denominators and are not differenced.

Solver and verifier (Fig. S1) are LLM agents on a consumer subscription, with no fine-tuning and no API spend. They are cheap to *re-run* but not reproducible: the harness pins no model snapshot, exposes no temperature or seed, and carries an undisclosed product system prompt. Our own dispatch instruction for the main rounds was also not captured, only the per-compound data it wrapped, so an outside group reproduces this distributionally rather than exactly (§9). Scoring is exactly reproducible: predictions, ground truth and scorers are released, and every number in

the training-free core (§3–§5.3, §4.6) regenerates from them. The two exceptions are §5.4’s HOSE lookup and GNN, which need an nmrshiftdb2 dump we cannot redistribute (see *Data availability*); that GNN and the §5.6 generator are trained probes, fenced as complements to the training-free core.

## 1.1 Related work

**Trained spectra  $\rightarrow$  structure models.** The dominant line trains sequence or graph decoders on spectra: *Spectro*, from our own group<sup>1</sup>; a multitask CNN+transformer for routine 1D-NMR<sup>4</sup>; set/graph transformers such as NMRTrans<sup>5</sup>; and — closest to our multimodal setting — *NMIRacle*<sup>6</sup>, conditioned jointly on IR +  $^1\text{H}$  +  $^{13}\text{C}$ . They are accurate *in-distribution* but retrained per modality and dependent on a labelled spectra  $\rightarrow$  structure corpus. That resource is scarce rather than absent, and one concurrent effort attacks it the same way we do: NMRTrans builds NMRSpec, a literature-mined corpus of experimental  $^1\text{H}/^{13}\text{C}$  spectra<sup>5</sup>. IReXP is the complement rather than the competitor — infrared band lists, released under permissive licences for redistribution (§2.3) — and the two together are the open, experimental, multimodal corpus neither supplies alone. Our own group’s prior infrared work sits upstream of elucidation rather than at it: *vIR-OLO* recasts functional-group identification as object detection, localising bands as well as labelling them<sup>7</sup>, and *j-IR-vis* learns transferable spectral embeddings from IR<sup>8</sup>. Neither generates candidate structures, which is the step this paper finds binding.

**LLMs as elucidators, and how this work differs.** LLM-orchestrated tools already plan and execute real syntheses (ChemCrow<sup>9</sup>, Coscientist<sup>10</sup>); such systems *call* an elucidation routine, and this paper measures what one returns. General-purpose LLMs have been applied off-the-shelf<sup>2</sup>, as have multimodal models over multi-spectral input (SpectraLLM, SpecMol)<sup>11,12</sup> and knowledge-enhanced tree-search reasoning<sup>13</sup>; *MolPuzzle*<sup>14</sup> supplies IR+MS+ $^1\text{H}$ + $^{13}\text{C}$  puzzles with the formula given, and *IR-Agent*<sup>15</sup> emulates expert IR interpretation on experimental spectra with a multi-agent framework. Priessner *et al.* run the same generate-then-rerank shape in this journal, re-ranking generator candidates with a reasoning LLM<sup>16</sup>. We claim priority on none of them: IR-Agent improves *how the model reads a spectrum*, where we ask what limits the outcome once it has.

Prior benchmarks and trained baselines evaluate on simulated or software-predicted spectra<sup>1,6</sup> or on curated single-instrument libraries — Alberts *et al.* use NIST gas-phase IR<sup>17,18</sup> — where IRSpectra-Bench uses spectra as reported by the authors of the source papers, across thousands of laboratories and instruments; the contrast is with both simulation *and* curated uniformity. Recent benchmarks bound the task elsewhere: NMRGym’s 269,999-molecule scaffold-split experimental  $^1\text{H}/^{13}\text{C}$  suite<sup>19</sup>; MolQuest, agentic and multi-turn, where the best models reach roughly 50% and most fall below 30%<sup>3</sup>; and SpecX, where multimodal LLMs “lack precise spectral grounding”<sup>20</sup> — the limit §5.1 measures. That line is also the state of the art for *trained, formula-conditioned IR* elucidation — an IR  $\rightarrow$  structure transformer<sup>17</sup> now at 63.8% top-1 / 84.0% top-10 on experimental NIST gas-phase spectra given the formula, pretrained on 1,399,806 simulated spectra<sup>18</sup>. That headline is on a 6–13-heavy-atom subset, the range where our own accuracy is highest (60.5% top-1 for  $\leq 15$  heavy atoms, §4.1), so on comparable molecules the training-free LLM and the purpose-trained transformer are closer than the headline numbers suggest. Size is not the whole story, though, and we do not claim it is: the same work reports 59.94% on a 5–35-heavy-atom set ( $n=5,024$ ), which spans most of our range and falls only four points below that 63.8% headline. What differs there is the data rather than the size: a single-instrument gas-phase library against heterogeneous literature-reported band lists. Neither setting bounds the other. Reported accuracies also swing with inference method and scoring harness: GPT-4o is scored at 1.4% on MolPuzzle by the benchmark’s own authors<sup>14</sup>,

at 27.8% under a plain chain-of-thought harness, and at 57.8% with knowledge-enhanced tree-search reasoning<sup>13</sup> — a factor of forty for one model on one benchmark. Cross-paper numbers are therefore not comparable, and we fix and release a single pre-registered RDKit-InChIKey protocol with bootstrap CIs.

**Computational NMR for structure validation.** Forward-verification is the LLM analog of a workflow chemists already trust — compute the spectrum each candidate *would* give and match it to experiment: DP4 / DP4+ over GIAO-DFT shifts in solution<sup>21,22</sup>, *NMR crystallography* with GIPAW shifts in the solid state<sup>23,24</sup>. We swap the quantum-chemical predictor for a forward LLM, trading accuracy for zero setup cost, and inherit the principle that *verification by forward prediction is easier than inverse generation*. *Computer-assisted structure elucidation* (CASE) has run the loop since the 1990s and overturned published assignments<sup>25,26</sup>; its generator is a strict enumerator with exhaustive recall by construction<sup>26</sup> — the axis on which we find the LLM weak, and the axis the LLM benchmark literature almost never reports: across every paper we could read in full it appears once, in a supplementary figure caption (§6). The decomposition is stable under every perturbation we could run — four Claude models, a second chemical domain, four verifiers, all recall-bound (§4.4, §4.5, §5.4) — but is not tested outside the Claude family (§8).

**The generate-and-forward-verify loop is not ours.** *NMR-Solver*<sup>27</sup> builds it without an LLM, retrieving and fragment-recombining candidates and ranking them by <sup>1</sup>H/<sup>13</sup>C shifts forward-predicted with NMRNet (<sup>13</sup>C MAE 1.098 ppm); on  $\approx 450$  experimental literature spectra with the formula supplied it reports 52.89% top-1 with stereochemistry ignored, well above our 28.4% on the same criterion; the twin supplementary table scores the identical predictions with stereochemistry preserved and gives 31.56%, against our 21.1%. We quote both because the gap between them, 21 points on one system’s own predictions, is larger than most differences anyone compares. §5 asks a different question — *how much of the remaining error is generation and how much is verification* — answered by decomposition (§5.2) rather than by a better solver. Its predictor is roughly twice as sharp as the  $\approx 2$  ppm LLM forward-predictor that §5.1 identifies as the binding constraint, and it delivers roughly twice the top-1. That is what our §5.4 ablation predicts and what §5.1’s separability measurement implies. *NMRAgent*<sup>28</sup> couples spectral tools to a knowledge graph and validates on newly isolated natural products; it is complementary — a best-effort system, where we measure an off-the-shelf model.

**Concurrent work, and what we do not claim.** Kliuev *et al.* give six frontier LLMs the <sup>1</sup>H and <sup>13</sup>C peak lists for 105 molecules: the best solves 24%, against 69% for the best of four specialised solvers, and every model “drift[s] toward smaller molecules”<sup>29</sup> — an independent replication of our headline and of its size dependence. Espejo Morales *et al.* instead put a frozen LLM in a curated environment of processing tools, tabulated shifts and hard formula, unsaturation and integration checks: 80.9% top-1 on a chemistry-education set (N=236), 71.1% on the 15 molecules where graduate students reach 66.7%, and 20.6% on 34 industrial samples where a trained transformer reaches 14.7%<sup>30</sup>. Their frame — search rather than a learned map — is orthogonal to ours: they ask which architecture, we ask which stage of it binds. Their inputs are raw instrument files and ours published reports, which they exclude by design (§8). We claim priority on neither that frame, nor on external simulation tools improving agent elucidation, nor on placing frontier LLMs near a quarter on real peak lists. Nor on the idea that candidate supply binds: Priessner *et al.* measure generator recall as a standalone quantity and state the constraint outright — a re-ranker “cannot overcome cases where the correct structure is entirely absent from the candidate pool”<sup>16</sup>. What we add is the measurement at scale and the generalisation: the split on 194 compounds across four model

families and four verifiers, and then recovered from the published top- $k$  figures of every system we could read (§6), where it turns out to hold in some regimes and fail in others.

---

## 2. The IReXP dataset

### 2.1 Motivation

An IReXP record is a *band list* — peak positions in  $\text{cm}^{-1}$  transcribed by an author into a paper’s text — with co-reported  $^1\text{H}/^{13}\text{C}$  shift lists and, where resolvable, the 2D structure. IReXP contains no absorbance traces. That is the published form, and what a language model consumes, but it is a different object from a digitised spectrum and the two should not be counted against one another.

Among *digitised spectra*, free downloads run to the NIST WebBook<sup>31</sup> ( $\approx 16\text{k}$ ) and the Chemotion ELN deposit<sup>32</sup> ( $\approx 2\text{k}$ ). AIST SDBS<sup>33</sup> is larger ( $\approx 54\text{k}$  FT-IR, all structure-linked) but *view-only* (50 spectra/day, no bulk export), so IReXP does not compete there: SDBS alone holds more structure-linked spectra than IReXP holds structure-linked band lists. Among text-derived *band lists* IReXP is the largest *openly redistributable* collection by record count, a claim deliberately scoped to that object type.

### 2.2 Construction

Experimental sections report per-compound IR band lists with co-reported  $^1\text{H}/^{13}\text{C}$  NMR in a stable textual convention. Mining it is mature<sup>34</sup>; we add a pipeline specialised to the *IR band list*, the modality those tools passed over.

- **Discovery.** Open-access literature is harvested in bulk from the PMC Open-Access Subset<sup>35</sup> and the Chemotion FT-IR deposit (RADAR4Chem).
- **Extraction.** A deterministic parser segments experimental text per compound and extracts IR wavenumbers and  $^1\text{H}/^{13}\text{C}$  shifts, gated against scan-range artefacts and prose false-positives.
- **Resolution.** IUPAC names go to SMILES with OPSIN<sup>36</sup> and RDKit<sup>37</sup> canonicalisation (InChIKey, SELFIES<sup>38</sup>), with a PubChem<sup>39</sup> fallback for trivial names; handling open-access label conventions and narrative artefacts raised structure coverage from 24% to 35%.

### 2.3 Contents and licensing

Table 1 gives what the release holds and where each part came from.

Table 1. IRexp dataset contents and provenance.

| field                                                                     | value          |
|---------------------------------------------------------------------------|----------------|
| experimental IR band-list records                                         | <b>121,233</b> |
| ...co-reporting $^1\text{H}$ and/or $^{13}\text{C}$ NMR                   | 87,075 (72%)   |
| ...with a resolved 2D structure                                           | 43,060 (35.5%) |
| ...of these, with $^1\text{H}$ and/or $^{13}\text{C}$ NMR                 | 40,702         |
| ...of these, with both $^1\text{H}$ and $^{13}\text{C}$ (full quadruples) | 33,201         |
| <i>by provenance</i> : author-transcribed from PMC-OA text (CC-BY)        | 119,345        |
| <i>by provenance</i> : peak-picked from Chemotion ELN deposits (CC-BY-SA) | 1,888          |

The pools differ: PMC records are author-transcribed (median 9 bands), Chemotion records peak-picked from deposited spectra (median 39), so users training on band density should treat them apart. `scripts/split_license_pools.py` separates them on `source_doi` and stamps each record’s licence.

**Extraction fidelity is measured.** On a seed-fixed random sample of 60 PMC-sourced records (`scripts/audit_extraction.py --n 60 --seed 0`) we re-fetched each article and checked every recorded wavenumber against its text: 560/560 bands and 60/60 records were confirmed (Wilson 95% CI 99.3–100% and 94–100%). This bounds *transcription* error — hallucinated, mis-parsed or unit-mangled values — below 1% of bands. It does not measure whether the parser found every IR string in every paper, a recall question requiring human reading; that audit is prepared but not yet run (§8).

`irexp_resolved` (43,060 records, 100% structure-linked) is the benchmark-ready split,  $\approx 6\times$  the 6,833-molecule set used to train Spectro<sup>1</sup> (Fig. S2).

### 3. Benchmark design (IRSpectra-Bench)

From `irexp_resolved` we draw *IRSpectra-Bench*, 194 blind elucidation problems, each giving the *molecular formula* (as from HRMS), the *IR band list* and the  $^1\text{H}$  and  $^{13}\text{C}$  *shift lists*; no name, SMILES or scaffold hint. One exception is measured rather than assumed: the shift strings are the source paper’s, and in 10 of the 194 the authors’ peak assignments name a ring system (2CH<sub>2</sub>·pyrrolidine, pyridazinone H5). Those compounds are solved 1/10 against 54/184 (29.3%) for the rest (`scripts/prompt_leakage.py`), so the annotation does not carry the headline — they are harder than average, which is presumably why their authors annotated them. We keep them rather than drop them. Main-round ground truths are spectrally validated by an automated RD-Kit check ( $^{13}\text{C}$  peaks vs symmetry-unique carbons, formula match, SELFIES round-trip) excluding merged or incomplete spectra — 6/140, leaving 134. `scripts/validate_benchmark.py` runs the deterministic filter over all three rounds, regenerating every `clean_qids.json` from the released questions and answers.

The filter does not gate on  $^1\text{H}$ , and we report what that leaves. In 13 of the 194 retained records the reported  $^1\text{H}$  integral exceeds the reference structure’s hydrogen count — residual solvent, water, exchangeable protons, or a rotamer mixture. They are printed as a diagnostic, not excluded: the cohort was fixed in advance, and re-filtering post hoc on a criterion chosen after seeing results

is a degree of freedom a benchmark should not take. We report the sensitivity instead: dropping all 13 moves the headline from 28.4% to 29.3% (53/181), +0.9 points, leaving every conclusion unchanged. The controlled rounds’ 60 compounds are fixed, pre-registered sets, audited the same way but reported rather than filtered (57/60 pass; §8).

**Difficulty is declared in advance, as a property of the compound.** By RDKit ring analysis a compound is *simple* iff it has at most two rings, no fused/spiro/bridgehead system and  $\leq 22$  heavy atoms; all others are *complex* (98 simple / 96 complex), so the split is exhaustive. The threshold is not load-bearing: sweeping it from 18 to 26 leaves the simple-minus-complex top-1 gap inside a 36–40-point band throughout (36.1 at the narrowest, 39.6 as released; `scripts/difficulty_sensitivity.py`). This axis is distinct from the continuous size gradient of §4.1 ( $\leq 15$  / 16–25 /  $> 25$  heavy atoms), and InChIKey de-duplication across rounds prevents leakage.

**The 50/50 balance is by design, and the corpus is not balanced.** The sampler fills each stratum to half the round, so “98 simple / 96 complex” is a property of the draw and not an observation about the literature. Under the same eligibility filters the eligible corpus is 17.5% simple and 82.5% complex (5,059 / 23,929 of 28,988; median 26 heavy atoms against the benchmark’s 20), so simple compounds are enriched 2.9 $\times$  (`scripts/corpus_reweight.py`). Balancing is the right choice for measuring a gradient — an unbalanced draw would put almost no mass in the stratum where the model succeeds — but it means the headline is a weighted average under weights we chose. Reweighted to the corpus, top-1 falls from 28.4% to 15.2% [10.6–20.4] and recall from 33.5% to 19.8% [14.4–25.8], on bootstrap 95% CIs. **Verification precision has to be reweighted with them**, or the flattering half of the decomposition keeps the favourable composition: it is 94.3% (50/53) on simple compounds and 66.7% (8/12) on complex ones, so the pooled 89.2% becomes **71.5% [50.2–92.5]** under corpus weights — a wide interval on twelve complex compounds, and we report it as such. The inequality is unaffected and if anything sharpened: 71.5% against a 19.8% recall. We report the benchmark figure as the headline because it is the one the released set reproduces, and the reweighted figure as the better estimate of what a chemist meets on an arbitrary paper. The reweighting assumes the one sampler filter it cannot apply — recovery of the raw  $^1\text{H}$  string, which needs a network fetch per record — is independent of difficulty.

**Scoring is mechanical.** A prediction is *correct* if its RDKit InChIKey connectivity layer (first 14 characters) matches the reference: we score *constitution*, so correct constitution with wrong stereochemistry counts as correct. We report the strict alternative rather than assert it is immaterial — the *full*, stereochemistry-sensitive InChIKey gives 21.1% top-1 and 25.8% recovered (41/194 and 50/194) against 28.4% / 33.5% — 7.3 points lower on top-1 and 7.7 on recovered — so the constitution figure is an upper bound on full-stereochemistry accuracy. Constitution is the headline because 1D  $^1\text{H}/^{13}\text{C}$ /IR rarely fixes absolute configuration: only 10.3% (20/194) of answers carry a *defined* (assigned R/S) stereocentre, and a model is penalised there for information the prompt never contained (`scripts/score_main.py --stereo`).

**Metrics.** *Top-1 (exact constitution)* is the fraction whose best-ranked candidate matches at the connectivity layer; *recovered (top-3)* the fraction whose reference appears among the up-to-three returned candidates, the lenient “recovery” protocol of ref.<sup>2</sup> — their metric, under their name. *Generation recall* is the fraction whose reference is in the candidate pool *before* re-ranking, the ceiling any verifier can reach; it coincides with recovered (top-3) except in the generate-wide arm of §5.3, where the pool is larger. *Verification precision (conditional on recall)*, measured over the *recall-positive* compounds alone — those whose reference the solver did propose — isolates verification from generation. The forward-verifier ranks by a symmetric *chamfer distance* between predicted and observed  $^{13}\text{C}$  peak sets (lower is better, no equal-count requirement); Morgan(2,

2048) Tanimoto<sup>40</sup> gives a graded scaffold-family signal. CIs are bootstrap 95%; model-vs-model differences use McNemar’s exact test<sup>41</sup> with Holm correction<sup>42</sup>.

Solvers are frontier LLM agents (Claude Opus), one sub-agent per batch, closed-book under a consumer subscription: no tools beyond an RDKit formula check, no ground truth, verified by grep-auditing transcripts.

Formula adherence is high but imperfect. Of the 126 candidates carried through forward-verification on the original arm (§5.2), 91.3% (115/126) match the given formula exactly, and 76.6% do across the full top-3 pool. Adherence is also not uniform across rounds: on top-1 answers it is 95.0% (38/40) in v3 and 90.0% (18/20) in the v2-control — the two pre-registered controlled rounds, 60 compounds between them — against 77.6% (104/134) in the headline main round. That is a structure of the wrong composition, in violation of a constraint the solver was explicitly handed, for about one main-round compound in five. Such answers score as misses, so this inflates nothing, but the main-round arm is the weakest-constrained (`scripts/analyze_misses.py`).

## 4. How well do LLMs elucidate real structures?

### 4.1 Headline performance


Solver agents work blind from formula + IR + <sup>1</sup>H + <sup>13</sup>C and return up to three ranked candidates, in bounded contexts reset every 2–12 compounds rather than one long context — the variable that §4.3 shows matters. The benchmark is all 194 compounds (134 spectrally-validated + 60 controlled-round; `data/benchmark_main/raw/`).

**Table 2. Headline elucidation performance on IRSpectra-Bench (n=194).** Overall and by difficulty stratum, with bootstrap 95% CIs.

| metric                                      | overall (n=194)      | simple (n=98) | complex (n=96) |
|---------------------------------------------|----------------------|---------------|----------------|
| top-1 exact constitution                    | <b>28.4%</b> [22–35] | 48.0% [39–57] | 8.3% [3–15]    |
| recovered (top-3)                           | 33.5% [27–40]        | 54.1% [44–63] | 12.5% [6–20]   |
| scaffold-level (best Tanimoto $\geq 0.45$ ) | 56%                  | 73%           | 39%            |
| mean best Tanimoto                          | 0.59                 | 0.73          | 0.45           |

A strictly-validated subset with the controlled rounds restricted to their clean parts too (134 main-clean + 57 controlled-clean = 191) reproduces this within  $\approx 1$  point on every metric — top-1 28.8%, 95% CI 23–36; recall 34.0%; simple 49.0% / complex 8.4% — so the asymmetric inclusion of the pre-registered controlled sets does not drive it. What does move the overall figure is the deliberate 50/50 difficulty balance: reweighted to the 17.5%-simple composition of the eligible corpus, top-1 is 15.2% [10.6–20.4] and recall 19.8% [14.4–25.8] (§3). The per-stratum figures are unaffected — only the weights change — and the reweighted number is the one to read as performance on an arbitrary paper. Accuracy falls monotonically with size in step with the difficulty gradient (Fig. 2): 60.5% top-1 at  $\leq 15$  heavy atoms, 28.3% at 16–25, 7.0% above 25 (Fig. S3).

Of the 137 analysable top-1 misses (139 in all; two predictions did not parse; `scripts/analyze_misses.py`), 76.6% are constitutional isomers of the true structure against 23.4% with the wrong formula. That is not mostly regiochemistry: only 22.6% share the true Murcko scaffold,



**Fig. 2.** Top-1 and recovered accuracy on IRSpectra-Bench by difficulty (all / simple / complex, n=194), with bootstrap 95% confidence intervals. The benchmark separates a realistic difficulty range: simple targets are solved four to six times as often as complex ones on both metrics.

2.9% reach Tanimoto  $\geq 0.85$ , and the median isomeric-miss Tanimoto is 0.39. Strict regiochemistry accounts for a fifth of failures; most misses are larger rearrangements of the same atoms. Near-degenerate  $^1\text{H}/^{13}\text{C}$  shifts under-determine connectivity, hence HMBC and NOESY; what the misses show is *wrong connectivity at the right composition*.

## 4.2 Reconciling with prior reports

Our 28% top-1 sits far below the  $\approx 100\%$  on “simple” molecules reported for the same model class in a non-peer-reviewed company white paper<sup>2</sup>. Unsourced independently, it is best read against the *peer-reviewed* record: MolPuzzle<sup>14</sup> and its re-scorings<sup>13</sup>, and the trained baselines<sup>1,6</sup>. Four methodology and scoring choices differ, each in the direction that raises a reported number; we name them without apportioning the gap. **Difficulty:** ring count bounds difficulty from below rather than proxying for it — our simple stratum includes a hexasubstituted benzene, while  $\geq 4$  rings mark 38% of recall-negative compounds (those whose reference was never proposed) against 3% of recall-positive ones (§5.3). **Scoring:** prior work counts a recovery if the reference appears among three ranked candidates over three independent runs; we report single-run top-1 and top-3. **Hints:** prior hard targets got the starting-material SMILES and we give none, though the formula we do give is worth 5% top-1 alone with the spectra masked (§4.6). **Curation:** prior compounds were hand-selected; ours scraped, unfiltered for solvability.

**Versus trained models: a bound on the gap.** No system has been scored on our test set, and published numbers differ in the three respects that most move the score: spectrum realism, hints, exact-match definition. That is a reason not to read the comparisons below as a ranking; it is not a reason no baseline was run, and none was. We built and released the benchmark, so running an existing open elucidator on it is possible and is the obvious next comparison — it needs a system whose weights and inference code are public and that accepts IR +  $^1\text{H}$  +  $^{13}\text{C}$  with a supplied formula, which we did not have. The nearest thing here is the trained generator of §5.6, and it is ours, not independent. Most trained baselines report in-distribution accuracy on simulated spectra; Alberts et al. do not — 63.8% top-1 from IR alone, formula supplied, on experimental NIST gas-phase spectra of a curated single-instrument library of 6–13-heavy-atom molecules<sup>18</sup>. Ours span

8–60 heavy atoms (median 20), only 15 of 194 (8%) in that window; against our  $\leq 15$ -heavy-atom stratum (60.5%) their 63.8% is level — and level is not parity in our favour, because their model reads infrared alone where ours reads infrared and  $^1\text{H}$  and  $^{13}\text{C}$ . A three-modality frontier LLM matching a single-modality trained transformer is the comparison, and it does not flatter the LLM. Size alone does not explain the rest: the same work reports 59.94% over 5–35 heavy atoms ( $n=5,024$ ), a range that covers most of ours. The remaining difference is what the spectra are — one gas-phase instrument against thousands of laboratories — and the comparison bounds the gap between those regimes rather than ranking the two systems. Multimodal systems report more still, on data of their own making: Spectro 93% (82% with fixed embeddings) on a split whose IR is plotted from reference data and whose NMR is software-predicted<sup>1</sup>; NMIRacle 48% top-1 / 66% top-15 on molecules held out from a *simulated* corpus inside its own training distribution, as its authors note<sup>6</sup>. Ours is 28.4% top-1 (33.5% recovered, top-3), 29.9% with forward verification (§5), on blind, real, literature-mined experimental spectra of out-of-distribution compounds.

One asymmetry cuts against us: NMIRacle takes no molecular formula where we supply it (§3), so the settings differ on opposite axes — our spectra real and out-of-distribution against their simulated and in-distribution, their input strictly harder — and neither number bounds the other. Read only as a *bound on the simulated-to-real gap*, the contrast suggests high in-distribution accuracies substantially overstate real-world performance. Metric instability reinforces the caution: the  $\approx 40\times$  MolPuzzle swing in §1.1 (GPT-4o: 1.4%<sup>14</sup> to 27.8% to 57.8%<sup>13</sup>, method- and harness-dependent) bounds the weight any unaudited near-100% claim<sup>2</sup> can bear.

### 4.3 Methodology dominates: a within-compound control


The same 20 molecules were solved two ways: (a) one LLM context handling all sequentially with no tools, (b) four independent agents of five compounds each with RDKit formula-checking. Recovered (top-3) rose from 5% to 15% ( $1/20 \rightarrow 3/20$ ) and top-1 from 0% to 15% ( $0/20 \rightarrow 3/20$ ). McNemar’s exact test (§3) does not reach significance, and at this size it could not have: the exact test conditions on the discordant pairs, of which the top-1 arm has three ( $b=0$ ,  $c=3$ ,  $p=0.25$ ) and the recovered arm four ( $b=1$ ,  $c=3$ ,  $p=0.625$ ). Below six discordant pairs the smallest attainable two-sided  $p$  is above 0.05, so “not significant” here reports the design, not the data. The 15% carries a Wilson 95% CI of roughly 5–36%. Bounded, frequently-reset contexts with tool access therefore appear to raise measured performance, consistent in direction but not established in size at  $n=20$  ( $p=0.25$ ): a directional within-compound demonstration, and it does not measure a  $3\times$  effect. Small rounds swing widely (15–40% across  $n=20$ –40 draws), hence the full 194-compound headline. It plausibly explains *part* of the gap to optimistic prior reports, whose per-problem API calls implicitly used method (b) — a hypothesis, not a quantified contribution.

### 4.4 Model comparison: the benchmark orders capability but separates only the extremes

We solved a fixed 24-compound subset blind with four Claude models spanning a wide capability range, including the newest (Table 5), under one prompt, one scorer and one candidate budget (Fig. 3, Table 3).

One protocol asymmetry must be disclosed: the three comparison models solved the 24 compounds as four six-compound contexts each, whereas the Opus column reuses the headline run, where those items sat in one six-compound and two twelve-compound contexts — the variable that §4.3 measures at  $5\% \rightarrow 15\%$ . The Opus estimate is therefore not protocol-matched and, if anything, handicapped

by the longer contexts; this touches neither the nesting nor the Fable-vs-Haiku contrast, and a clean re-run is outstanding in docs/MODELS.md.



**Fig. 3.** Four-model comparison on a fixed 24-compound subset, solved blind under one protocol: Fable 5 46% > Opus 25% > Sonnet 21% > Haiku 0% top-1. The outcomes are strictly nested — each stronger model solves a superset of the weaker one’s compounds — so the benchmark is capability-sensitive, but at n=24 it is underpowered to separate adjacent models (§4.4).

**Table 3. Four-model comparison on a fixed 24-compound subset.** All four models ran the identical blind protocol.

| model          | top-1      | recovered | top-1 95% CI |
|----------------|------------|-----------|--------------|
| Claude Fable 5 | <b>46%</b> | 54%       | [25–67]      |
| Claude Opus    | 25%        | 29%       | [8–42]       |
| Claude Sonnet  | 21%        | 25%       | [8–38]       |
| Claude Haiku   | 0%         | 4%        | [0–14]       |

CI’s are bootstrap except Haiku’s: Clopper–Pearson exact for 0/24, where the percentile bootstrap is degenerate.

The models rank in monotonic capability order with strictly nested outcomes (Haiku  $\subset$  Sonnet  $\subset$  Opus  $\subset$  Fable), so the benchmark is capability-sensitive. The nesting makes the *ranking* robust, but at n=24 the subset is underpowered to separate adjacent models. By McNemar’s exact test only Fable-vs-Haiku survives multiple-comparison correction (Holm-adjusted p=0.006). Fable-vs-Opus does not: uncorrected p=0.063, which is the *floor* for its five discordant pairs, so the comparison had no way to reach 0.05 (Holm-adjusted p=0.19). Opus and Sonnet are indistinguishable. Two mid-tier models agreeing closely puts the recall-bound regime of §4.1 beyond a single model; the newest nearly doubles the next-best top-1 yet still misses the majority, leaving the benchmark far from saturated. We used Claude-family models because they are callable for free under one subscription; a cross-vendor sweep needs API access we did without, though the large monotonic spread makes it unlikely the pattern is lineage-specific.

## 4.5 Domain case study: battery-electrolyte chemistry

To test the recall-bound regime on battery-electrolyte chemistry rather than on organic molecules at large, we curated *IRSpectra-Bench-Electrolyte*: a 48-compound subset of IRExp records across the six functional families that dominate lithium- and sodium-battery electrolytes, eight compounds per class, held out of every other split and solved blind under the identical protocol. Two yielded no parseable candidate, leaving 46 scored. These are literature compounds, not operando or in-cell spectra, and we make no claim that they are.

Performance lands in the same broad regime as the headline benchmark — top-1 26%, recovered 28% (12/46 and 13/46; at  $n=46$  the 95% CI is wide and overlaps the headline), scored over the 46 that yielded a parseable candidate — consistent with a bottleneck in the elucidation task rather than in any one chemical neighbourhood. The main round instead scores an unparseable answer as a miss, and on that rule this arm reads  $12/48 = 25\%$ ; we report both rather than let two denominators sit unexplained. The true structure enters the candidate set for only 13 of 46 compounds and ranks first in 12 of those 13. Forward-verification was not run here, so this is the recall-bound pattern under the solver’s own ranking rather than the §5 measurement (Fig. S4). The per-class breakdown is Table 4; at eight compounds a class the intervals overlap throughout.

**Table 4. Per-class performance on IRSpectra-Bench-Electrolyte ( $n=46$ ).**

| electrolyte class    | n | top-1      | 95% CI  | recovered (top-3) | 95% CI  |
|----------------------|---|------------|---------|-------------------|---------|
| sp <sup>3</sup> -C-F | 8 | <b>50%</b> | [22–78] | 50%               | [22–78] |
| carbonate            | 7 | 29%        | [8–64]  | 43%               | [16–75] |
| phosphoryl           | 8 | 25%        | [7–59]  | 25%               | [7–59]  |
| glyme / oligoether   | 7 | 29%        | [8–64]  | 29%               | [8–64]  |
| sulfonyl / sulfonate | 8 | 12%        | [2–47]  | 12%               | [2–47]  |
| nitrile              | 8 | 12%        | [2–47]  | 12%               | [2–47]  |

The intervals overlap almost completely and the spread is within sampling noise (six-class  $\chi^2$   $p \approx 0.56$ ; best-vs-worst Fisher exact  $p \approx 0.28$ ), so the ordering is hypothesis-generating rather than an established ranking, though chemically legible: C–F couplings pin sp<sup>3</sup>-C–F regiochemistry, while sulfonyl and nitrile targets turn on oxidation-state ambiguity and heteroaromatic substitution that <sup>1</sup>H/<sup>13</sup>C shifts underdetermine. It is the recall-bound failure of §4.1 in a domain-specific guise, and it is where §5’s *forward-verification* recipe plays a role *analogous* to (not a replacement for) computational-NMR validation.

## 4.6 Is the model reading the spectra? A formula-only control

Every benchmark compound is mined from open-access literature, so a frontier model may have met it in pretraining, a well-documented hazard for LLM evaluation<sup>43</sup>. We reran the identical blind protocol with every spectral channel masked, leaving the solver the molecular formula and nothing else. A formula does not determine constitution, so accuracy materially above the floor would indicate recall rather than reasoning.


The arms were not generated together: the *formula-only* arm was fresh (2026-07-28), the *full-modality* arm archived from June. The bias therefore runs *toward* formula-only — fresh agents reason harder per compound than the archived batched run — and it still collapsed to 5%; a confound that works against the finding cannot manufacture it.

**Table 5. Formula-only control.** Paired on the same 60 compounds as §5.

| condition                                     | top-1            | recovered (top-3) |
|-----------------------------------------------|------------------|-------------------|
| formula only                                  | <b>3/60 (5%)</b> | 3/60 (5%)         |
| formula + IR + $^1\text{H}$ + $^{13}\text{C}$ | 14/60 (23%)      | 19/60 (32%)       |

The outcomes are perfectly nested: eleven compounds are solved with the spectra and not without, and none the other way round (McNemar exact  $p=0.001$ ). Removing the spectra removes the result.

**A second, independent control: publication recency.** The formula-only arm shows the spectra carry the signal, not whether memorisation contributes to the part they explain; if recall drove the headline number, accuracy should fall with recency. Across all 194 compounds, source papers spanning 2008–2026, accuracy is flat: 28.6% for the older half ( $\leq 2020$ ,  $n=112$ ) against 28.0% for the newer ( $n=82$ ), a point-biserial year–correctness correlation of  $r=-0.007$ , and no monotone trend across buckets (Fig. 4b), the most recent ( $\geq 2024$ ,  $n=25$ ) being highest at 40% [23–59].



**Fig. 4.** Two contamination controls. (a) The solver reaches 3/60 on formula alone against 14/60 with IR +  $^1\text{H}$  +  $^{13}\text{C}$ , nested outcomes: 11 solved only with the spectra, none only without (McNemar exact  $p=0.001$ ). (b) Accuracy against source-paper year ( $n=194$ , Wilson 95% CIs) is flat, point-biserial  $r=-0.007$ , where pretraining recall predicts a decline.

Newer papers skew larger (median 22 heavy atoms against 20) and size dominates accuracy, biasing the raw split *against* them (§4.1); size-adjusted, the older-minus-newer difference is  $-5.1$  points, 95% CI  $[-17.2$  to  $+7.0]$  (Cochran–Mantel–Haenszel  $\chi^2=0.42$ ,  $p=0.51$ , continuity-corrected). Its sign is opposite to what pretraining recall would produce, but the interval comfortably includes zero, so this bounds any recency effect rather than demonstrating a reversed one, and by the standard that §5.3 applies to its own adjacent conditions we do not read it as directional.

The 5% is not zero, and its three compounds do not all read the same way: one is absent from PubChem and near-determined by its benzyne-precursor composition, while a catalogued sulfonamide and a named natural product are as consistent with recognition as with inference. Either way, formula-level recall accounts for at most about a fifth of the headline accuracy on this set. The two controls are independent and agree, but neither is randomised — publication year is observational, and the formula-only arm cannot distinguish memorisation from inference on a near-determining

formula — so we claim a strong bound rather than exclusion (§8). Nor can either control address retrieval on the spectral string itself, which the prompt carries verbatim from the source article; §8 states that objection and names the control that would settle it.

#### 4.7 Does the diagnosis hold outside one vendor? A four-vendor replication

Every number so far comes from one lineage, the external-validity gap §8 (iii) named as most important. Three terms carry the table below, and §5 treats them in full. *Forward-verification* re-ranks a solver’s candidate structures by predicting each one’s  $^{13}\text{C}$  spectrum and matching it against the observed spectrum. *Generation recall* is the fraction of compounds whose true structure the solver proposed at all, and *verification precision (conditional on recall)* is the fraction of those recalled compounds that the re-ranking then puts first. The *60-compound arm* is the two pre-registered controlled rounds, v3 and v2-control, on which the original decomposition was run (Table 7’s first column). One artefact of precision is worth naming in advance: a compound for which the solver returned a single candidate is scored correct by any re-ranker, so a candidate set rich in singletons inflates precision by construction, which is why the last column repeats it over the compounds where more than one candidate existed.

Six non-Claude models solved the same 60 compounds as Table 7’s first column under the identical blind protocol; the three whose output was well-formed enough were carried through forward-verification (Table 6).

**Table 6. Cross-vendor decomposition on the 60-compound arm.** Recall and precision have different denominators, so the criterion is the inequality rather than a difference.

| model                    | generation recall   | verification precision   recall | multi-candidate only |
|--------------------------|---------------------|---------------------------------|----------------------|
| Claude Opus<br>(Table 7) | 19/60 = 32% [21–44] | 16/19 = 84% [62–94]             | 10/13 = 77%          |
| Grok 4.6                 | 32/60 = 53% [41–65] | 20/32 = 62% [45–77]             | 20/32 = 62%          |
| Gemini 3.7 Flash         | 30/60 = 50% [38–62] | 22/30 = 73% [56–86]             | 22/30 = 73%          |
| GPT-5.6 Sol              | 25/60 = 42% [30–54] | 17/25 = 68% [48–83]             | 16/24 = 67%          |


The inequality holds in every arm (Fig. 5): verification precision exceeds generation recall for four independent model families. Recall and precision are measured on the same compounds, so what carries the claim is the paired difference rather than whether two marginal intervals overlap. The quantity bootstrapped is that difference *within* each vendor — its verification precision conditional on recall minus its own generation recall — not a difference against Claude. Resampling compounds (`scripts/cross_vendor_gap.py`) separates it from zero in three of the four arms: Claude +52.5 points [+31.2 to +71.7], GPT-5.6 Sol +26.3 [+4.5 to +48.6] and Gemini +23.3 [+3.2 to +44.3]. Grok does not — +9.2 [−11.7 to +30.0] — and we report its gap as directional. Claude’s is the widest of the four and also the least informative: the singleton composition described next inflates it.

Two findings sit underneath. Six of Claude’s nineteen recall-positive compounds carried one candidate, so nothing was ranked and any verifier scores them by construction; §5.2 makes that point about the full benchmark, and here it inflates Claude’s own precision alone. The newer models almost always return three candidates (0–1 singletons), and where a choice existed the four precisions fall in a band that dissolves the apparent Claude advantage. Verification also looks vendor-

independent where generation does not, precision on multi-candidate sets spanning 62–77% while recall spans 32–53%: the limiting factor is the generator. Three models beat ours on generation recall, which is the diagnosis working: recall is the movable factor.

**The candidate budget is not matched, and the correction runs both ways.** Recall above asks whether the true structure is anywhere in the candidate list, and the lists differ in length: ours holds 2.20 candidates per compound on this arm against exactly 3.00 for every comparison model (Composer 2.82). A longer list can only raise recall, so the comparison is tilted toward the other vendors — the mirror of the singleton effect we correct on the precision side. At one candidate, the only budget every model met, recall is  $14/60 = 23\%$  for ours against  $23/60 = 38\%$  for Grok,  $23/60 = 38\%$  for Gemini and  $21/60 = 35\%$  for GPT-5.6 Sol (`scripts/cross_vendor_budget.py`). The ordering survives and three models still lead ours, but Grok’s margin is 15 points. The 21-point figure is not one a budget-matched comparison supports.

Three weaker models are excluded from the decomposition: Composer 2.5 (20% recall) and GPT-5.6 Luna (15%) match the given formula only 67% and 76% of the time, and `nvidia/nemotron-3.5-lightning` managed 2%. §4.1’s scorer now prints formula adherence beside recall. A seventh arm, DeepSeek V4 Pro, returned answers for only 18 of 60 compounds before the run ended; the true structure appears for 8 of those, so its recall is a lower bound and it is excluded from the comparison rather than from the release — both its solve and verify files are deposited.



**Fig. 5.** Every model measured on the 60-compound arm. (a) Generation recall, each model at the candidate budget it used: 3.00 candidates per compound for all but ours (2.20) and Composer (2.82). That axis therefore favours the models that always returned three. Grey marks models below the formula-adherence floor, whose recall is not a chemistry result, and hatching marks an incomplete arm — DeepSeek V4 Pro answered 18 of 60 — whose recall is a lower bound. (b) The claim itself: verification precision against generation recall, with the diagonal. Every model sits above the line, so verification is better than generation for all four families; the hatched hollow marker is again the incomplete arm. The four Claude models of §4.4 are deliberately absent: they ran a different 24-compound subset, and putting them here would imply a comparison that was never made.

**Contamination was controlled.** Blindness rested on an instruction: these runs used cloud agents with repository access to tracked answer keys. Grok re-solved all ten batches from a clone with the answer files removed: recall 28/60 against 32/60, paired McNemar  $p=0.39$ , formula adherence 97%

against 95% of candidates (`scripts/cross_vendor_sweep.py`). The asymmetry is decisive — a key-reader would solve a *superset*, yet four compounds were solved only in the arm that had no key, with 24 of 60 solved by both, which is sampling noise rather than copying. And on constitutionally correct structures whose reference carries assigned stereocentres, which 1D spectra cannot supply, Grok reproduced 0/3 correct descriptors, Gemini 0/2 and GPT-5.6 Sol 0/2.

Two limits. The models were reached through a coding-assistant harness exposing reasoning-effort tiers, a decoding control the Claude arm never had (§9), so these measure a model *as served*, not a bare endpoint — and which tier served each arm was not recorded. The harness allowlist carries both `gpt-5.6-sol-xhigh` and `gpt-5.6-sol-high`, and the run did not report its selection, so the arms are not demonstrably matched to one another either. That bounds what this section can claim, and the bound falls unevenly. Recall and precision within an arm come from the same run at the same tier, whichever it was, so the inequality — the result — is internally valid in every arm. A recall *ranking between named models* is not: it orders models at unknown effort, and the 15-point margin above should be read as such. The other limit is that the clean-clone control covers only Grok; Gemini and GPT-5.6 Sol rest on the stereochemistry argument and the shared protocol.

---

## 5. Forward-verification elucidation

### 5.1 Method

The inverse direction is the model’s hard, isomer-blind direction; the forward direction (structure  $\rightarrow$  spectrum) is its easy, accurate one<sup>2</sup>. We close a training-free generator–verifier loop:


*generate* candidate structures (inverse)  $\rightarrow$  *forward-predict* each candidate’s  $^{13}\text{C}$  spectrum, blind to the observed spectrum  $\rightarrow$  *re-rank* candidates by the distance between predicted and observed  $^{13}\text{C}$   $\rightarrow$  return the best match.

Candidates are the inverse solver’s proposals: 373 deduplicated structures across the 194 targets, forward-predicted in shuffled, anonymised batches, blind to the observed spectrum and the target’s identity. Predicted and observed  $^{13}\text{C}$  peak sets are then compared by symmetric chamfer distance (Fig. 6).

Regioisomers have *different* forward-predicted shifts, but the margin is thin: isomeric candidate pairs for the same target separate by a median 1.21 ppm (quartiles 0.84 and 1.78), and 82% are predicted closer together than the predictor’s own  $\approx 2$  ppm error — usually within the noise. The ranking nevertheless recovers the right one 89% of the time when it is present (§5.2), a real effect on a thin margin, confirmed against a derangement null, in which every candidate set is re-scored against another compound’s observed spectrum so that none keeps its own, at  $p=0.001$  (§5.5).

Verification precision runs 72–89% rather than 100%: near-degenerate candidates the predictor cannot resolve are its false positives. The same spacing sets the limits that §5 reports separately: the 54.0% derangement chance floor on multi-candidate compounds (§5.5) and the ceiling neither the HOSE lookup nor the GNN escapes (§5.4). Sharper predictors would move them, and better search would not; they exist:  $\approx 1$  ppm message-passing ensembles<sup>44</sup>, a DFT-coupled graph network at 0.94 ppm<sup>45</sup>, a community benchmark<sup>46</sup> and curated shift sets<sup>47</sup> for calibration.

The scheme is an analogy to DP4<sup>21,22</sup> and NMR crystallography<sup>23,24</sup>, not an equivalence: we forgo DFT’s  $\approx 1$ –2 ppm accuracy and its calibrated, probabilistic error model for a predictor that runs in seconds.



**Fig. 6.** Forward-verification on a benchmark regioisomer pair: picolinamide and nicotinamide are indistinguishable to the inverse task, but the true isomer's forward-predicted <sup>13</sup>C spectrum matches the observed one at a chamfer of 0.42 ppm against 1.30 ppm for the alternative (§5.1).

## 5.2 Result

We ran the verifier over every compound in the benchmark. The first column below is the 60-compound arm (v3 + v2-control) on which §5.3–§5.5 build; the second is the full benchmark. All 373 candidates were forward-predicted, so nothing here is a lower bound; the self-ranking row re-derives Table 2 from §4’s output.

**Table 7. Forward-verification decomposition.** Original arm and full benchmark.

|                                                     | 60-compound arm | full benchmark<br>(n=194) |
|-----------------------------------------------------|-----------------|---------------------------|
| generation recall                                   | 19/60 (32%)     | 65/194 (34%)              |
| top-1, solver self-ranking                          | 14/60 (23%)     | 55/194 (28%)              |
| top-1, forward-verified re-ranking                  | 16/60 (27%)     | 58/194 (30%)              |
| <i>conditional on recall</i> — self-ranking         | 14/19 (74%)     | 55/65 (85%)               |
| <i>conditional on recall</i> — forward-verification | 16/19 (84%)     | 58/65 ( <b>89%</b> )      |
| ...single-candidate compounds within that set       | 6/19            | 28/65                     |
| <i>multi-candidate only</i> — self-ranking          | 8/13 (62%)      | 27/37 (73%)               |
| <i>multi-candidate only</i> — forward-verification  | 10/13 (77%)     | 30/37 (81%)               |

When the true structure is among the candidates, forward-verification selects it in 58 of 65 cases (89%). Of the 65, 28 had a single candidate, which any ranker scores by construction, so that pooled figure is not the verifier’s margin over chance. On the 37 where a choice existed, forward-verification gets 30/37 (81%) against a 54.0% derangement floor (§5.5) — +27.1 points; self-ranking 27/37 (73%); we treat the 37 as the denominator that measures verification.

The margin over self-ranking remains small and unresolved: seven compounds gained, four lost, McNemar exact  $p=0.55$ . We do not claim it. The load-bearing claim is that precision is high in absolute terms while recall binds: top-1 moves only 28% → 30% because the true structure was never proposed for 129 of 194 compounds, which no re-ranking can repair. The verifier converts recall almost completely on *simple* targets (50/53, 94%) and is exactly tied with self-ranking on *complex* ones (8/12, 67%). Elucidation factorises into two near-independent levers: the verifier is already strong (89%); the generator (34% recall) is the wall (Fig. 1).


## 5.3 Generate-wide: testing the recipe

The decomposition implies a recipe: *generate wide, verify by forward prediction* — a chemistry analog of self-consistency sampling<sup>48</sup>. Ten solver agents proposed up to six regiochemistry-aware candidates per compound, pooled with the originals and re-ranked as before.

**Table 8. Generate-wide vs original.** On the 60-compound arm only, since wide generation was run there; the “original” column is Table 7’s first.

|                                                     | original<br>(60-compound arm) | generate-wide |
|-----------------------------------------------------|-------------------------------|---------------|
| generation recall (true structure among candidates) | 32%                           | <b>42%</b>    |
| forward-verified top-1                              | 27%                           | 30%           |
| verification precision (conditional on recall)      | 84%                           | 72%           |

Wide generation lifts recall +10 points (32%  $\rightarrow$  42%, 19/60  $\rightarrow$  25/60) and top-1 +7 over the self-ranking baseline (23%  $\rightarrow$  30%, 14/60  $\rightarrow$  18/60), i.e. +3 over the forward-verified top-1 (27%  $\rightarrow$  30%, 16/60  $\rightarrow$  18/60, Table 8), on the same 60 compounds (Fig. 7) with no training.



**Fig. 7.** The forward-verification inference ladder on the 60-compound arm: solver self-ranking  $\rightarrow$  + forward-verification  $\rightarrow$  + generate-wide, at 23% / 27% / 30% top-1. Each rung is a training-free change to inference alone, directional rather than individually resolved at this sample size (§5.3).

**These top-1 differences are directional; at n=60 they are not statistically resolved.**

The stages are not nested: self-ranking to generate-wide gains seven compounds and loses three, McNemar exact  $p=0.34$  ( $b=3$ ,  $c=7$ ). The recall gain is the better-supported effect: the ladder shows the *mechanism* works and recall is the movable factor; it does not fix how large the top-1 gain is.

The ladder does not extend. (i) Recall plateaus at 42%, on large polycyclic targets rather than exotic ones —  $\geq 4$  rings in 3.1% of recalled compounds against 38.0% of missed ones. (ii) Verification precision falls 84%  $\rightarrow$  72% as more near-degenerate regioisomers enter the pool. This recall/precision tension is the ceiling of the training-free approach: closing the gap requires sharper verification or 2D-NMR constraints, not merely more candidates.

**How much of the wall is regiochemistry?** Enough to matter and not enough to remove it. Of the 129 compounds the solver misses, 97 (75.2%) have a candidate of the *right molecular formula* and 41 (31.8%) share both formula and generic scaffold with one — the class a deterministic isomer enumerator can reach (`scripts/analyze_recall_headroom.py`). Adding those to the 65 already recalled gives a loose upper bound of  $106/194 = 54.6\%$ , and the enumerator’s realised reach is

41.8% (§5.3), so roughly a third of the remaining wall is substituent placement around a scaffold the model already has, and two thirds is a scaffold it never proposed. The bound is loose because generic-scaffold equality overcounts what an enumerator can actually build; we report the realised figure as the achievable one.

#### 5.4 Non-LLM verifiers: a deterministic lookup and a learned model

§5.3 suggests an obvious fix: a non-LLM  $^{13}\text{C}$  predictor in the verifier slot. We tested a HOSE-code<sup>49</sup>-style lookup and a small message-passing GNN, both trained on the same nmrshiftdb2 dump<sup>50</sup> and applied to the same §5.2 candidate sets, so that only the predictor changes. The GNN is deliberately modest; sharper purpose-built  $^{13}\text{C}$  models exist<sup>44–46</sup>.

**Table 9. Verifier comparison, conditional on recall.**

| verifier                         | 60-compound arm<br>(n=19) | full benchmark<br>(n=65) | held-out $^{13}\text{C}$<br>MAE |
|----------------------------------|---------------------------|--------------------------|---------------------------------|
| solver self-ranking              | 14/19 (74%)               | 55/65 (85%)              | —                               |
| deterministic HOSE <i>lookup</i> | 14/19 (74%)               | 55/65 (85%)              | 3.23 ppm                        |
| learned GNN (same data)          | 16/19 (84%)               | <b>59/65 (91%)</b>       | 1.70 ppm                        |
| LLM forward-verifier (§5.2)      | 16/19 (84%)               | 58/65 (89%)              | —                               |

The lookup ties the solver’s own score at both sample sizes (Table 9). Coverage is the natural explanation, and the next paragraph is the test of it: of the 6,360 candidate carbons, only 2% match a training environment at  $r=4$  and 71% resolve at  $r\leq 2$  or the hybridisation prior, too coarse to separate regioisomers.

The GNN matches and slightly exceeds the LLM verifier, 59/65 against 58/65 (Fig. S6), though not significantly: *suggestive and directional*, consistent with the lookup’s failure being substantially method rather than coverage alone, but not established. Neither predictor reaches the near-degenerate-regioisomer precision ceiling (§5.3/§5.5), where DFT-level accuracy or orthogonal 2D-NMR constraints remain the genuine fix.

Leakage is slight: of the 373 candidate structures the two predictors re-rank (364 distinct InChIKey-14), two appear in nmrshiftdb2, and both are distractors on recall-negative compounds; none of the 65 candidates that *are* benchmark answers appears there (`scripts/verifier_leakage.py`). A Y-randomisation control (1,000 derangements) places the real result above the 97.5th percentile of chance (n=19: real 84% vs mean 58.6%, one-sided  $p<0.05$ ). Like §5.6, the learned verifier is a trained complement, reported outside the training-free protocol.

Concurrent industry work adds a fourth point on the same axis, outside our protocol. Wagen gave a frontier LLM up to ten calls to MagNET, an equivariant graph transformer whose predictions on the true structures of that benchmark reach 1.20 ppm MAE and 1.60 ppm RMSD against experiment<sup>51</sup>, on eight  $^{13}\text{C}$  problems drawn from structures that were misassigned and later corrected computationally<sup>52</sup>, at 56 runs per configuration<sup>53</sup>. Scored as we score, on connectivity, the external predictor moves 46.4% (26/56) to 55.4% (31/56) — the same order as our own ladder (§5.3); the larger 21.4% to 51.8% headline is on exact SMILES including stereochemistry, which our protocol does not test (§8). The author reports the study as underpowered and preliminary at  $n=8$  and flags possible training-set leakage, so we read it as evidence of direction rather than of magnitude.

## 5.5 Negative control

**Permutation negative control (Y-randomisation analog).** A re-ranker exploiting genuine predicted-vs-observed  $^{13}\text{C}$  agreement must lose it when the pairing is broken. We re-paired which observed  $^{13}\text{C}$  spectrum each candidate set is scored against — a *derangement*, so no compound keeps its own spectrum — and re-ran the verifier 1,000 times. Verification precision falls from the true 89.2% (58/65) to a permuted mean of 73.8% (95% range 66.2–81.5%; one-sided empirical  $p=0.001$ , two-sided  $p=0.002$ ). The verifier acts on real spectral agreement rather than a candidate-list artefact. That floor is high because 28 of the 65 recall-positive compounds carry a single scorable candidate, and a compound with nothing to rank is scored correct under *every* pairing. Those compounds enter the permuted score as guaranteed hits while carrying no verification signal, so the control was partly measuring the composition of the recall-positive set. On the 37 of those 65 that carried more than one candidate — the set that §5.2 calls the one that measures verification — the floor falls to 54.0% and the real value is 81.1% (30/37), a margin of +27.1 points rather than +15.4, at the same one-sided  $p=0.001$ . We report the restricted figure as the verifier’s margin over chance and the pooled one for continuity with §5.2’s denominator.

**Confidence calibration (a negative result).** Ranking the 138 multi-candidate compounds of the whole benchmark — every compound that received more than one candidate, of the 192 that received any, and so a wider set than the 37 recall-positive ones above — by chamfer margin and answering only the most-confident fraction leaves top-1 flat and non-monotonic with coverage (22% at full coverage, 24% at 75%, 28% at 50%, 24% at 25%). Single-candidate compounds, which have no margin, must be excluded; retaining them in an earlier analysis produced a spurious “improvement”. We therefore do not claim the verifier distance as a calibrated abstention gauge (§5.4). This does not contradict the chamfer distance being informative in aggregate (§5.2): the *absolute* distance separates right from wrong answers across compounds, while the *within-compound margin* between best and second-best does not rank one compound’s confidence against another’s. Selecting which questions to answer needs the second, and only the first is established here.

## 5.6 Is the recall wall task-intrinsic? A trained-generator probe

The ceiling above is a property of *training-free LLM elicitation*; whether the  $\approx 42\%$  recall plateau is intrinsic to 1D-data elucidation is a separate question. We probed it by substituting a small ( $\approx 16\text{M}$ -parameter)  $^1\text{H}/^{13}\text{C} \rightarrow \text{SMILES}$  transformer for the enumerator; its candidates were pooled with Claude’s and re-ranked by the same verifiers.

On the 194-compound benchmark the true structure enters the candidate pool for 54.1% of compounds, versus 41.8% for scaffold enumeration and 33.5% for Claude alone. Neither of the first two is an independent elucidator, and the comparison should not be read as one: enumeration relocates substituents *on a model candidate* (`scripts/enumerate_isomers.py`), so it is Claude-seeded, and the generator’s candidates are pooled with Claude’s. All three rows measure what is added to the LLM’s pool rather than what a system reaches on its own. Where enumeration’s near-degenerate regioisomers *collapse* the HOSE verifier (top-1 28.4%  $\rightarrow$  16.0%, the §5.3 precision-loss mechanism), the generator’s formula-correct,  $^{13}\text{C}$ -separable candidates convert: top-1 rises 28.4%  $\rightarrow$  35.1% (McNemar exact  $p=0.015$ ; +6.7 points, 95% CI [+2.1 to +11.9]; Fig. S5). The family this is corrected over is the probe’s own three arms — Claude alone, plus enumeration, plus generator — where it survives Holm (0.029) and Benjamini–Hochberg (0.022). It would not survive a correction taken over every  $p$ -value in the paper, and we name the family rather than leave the scope to be guessed. With the LLM forward-verifier on the 60 forward-verify compounds, recall rises 42%  $\rightarrow$

57% (34/60) and top-1 reaches 47% (28/60), at 82% verification precision (28/34). Those two figures come from a re-run rather than a reproduction: the earlier pass reported 41% top-1 at 73% precision but deposited no outputs, so we do not stand behind it, and the re-run lands above it.

Two controls guard against memorisation: the fine-tuning split removes every benchmark InChIKey-14 ( $\text{train} \cap \text{benchmark} = 0$ ,  $\text{val} \cap \text{benchmark} = 0$ ), with none of the 40 newly-recovered compounds in either training stage (0/40 in simulated pretraining, 0/40 in fine-tuning); and the simulated-pretrained model alone recovers 0/248 zero-shot on the probe’s own held-out split — all four cohorts pooled and de-duplicated at InChIKey-14 (140 main, 48 electrolyte, 40 v3, 20 v2-control), a superset of the 194-compound benchmark — rising to 25% on that same 248-structure split, and only after IRexp fine-tuning (`contrib/generator_probe/`). The IRexp data, rather than the architecture, is the active ingredient, so the recall ceiling is elicitation-specific rather than task-intrinsic. We report this as a probe, not part of the headline training-free protocol.

Purpose-trained systems make the point more strongly: NMR-Solver reaches 52.9% top-1 on experimental literature  $^1\text{H}/^{13}\text{C}$  with the formula supplied<sup>27</sup>, and a purpose-trained IR transformer 63.8% top-1 on experimental NIST spectra<sup>18</sup>. Whatever bounds a training-free LLM at 28–30%, it is plainly not the task. The public split `irexp_release/train` does *not* hold out the benchmark (117/200 InChIKey-14 overlap), so downstream users must de-leak; see *Data availability*.

---

## 6. The decomposition across the published literature

The split measured above needs no new data to apply to work already published. A system that returns a ranked list and reports  $\text{top-1} = a$  and  $\text{top-}k = b$  has already said this much: the true structure was in its candidate set at least as often as it reached the top  $k$ , so **recall**  $\geq b$ ; conditional precision is  $\leq a/b$ , since any recall above  $b$  divides the same  $a$  by a larger number; and at least  $b - a$  compounds were proposed and then mis-ranked. Where a paper states that its emitted list is capped at the reported  $k$  the floor is exact — Alberts et al. generate ten ranked SMILES per sample — and where the internal pool is larger it is a genuine bound, as with NMR-Solver’s documented 1,000-molecule pool behind a top-10 column. Table 10 applies this to every system we could read in full (`scripts/literature_decomposition.py`).

One caution governs the whole table. The  $a/b$  ceiling bounds the ranking a pipeline *already does*, not what any verifier could do: our own solver’s ranking is  $55/65 = 84.6\%$ , which  $28.4/33.5 = 84.8\%$  recovers from our published figures alone, and forward-verification exceeds it at  $58/65 = 89.2\%$  precisely by re-ordering. The distance from  $a/b$  to 100% is the room a re-ranker has.

**Table 10.** Recall and ranking loss recovered from published top- $k$  figures. Rows are comparable only within a correctness criterion; a stereochemistry-strict figure and a connectivity figure differ by 21 points on NMR-Solver’s own identical predictions. Candidate budgets differ and are given, because a recall measured over three candidates is mechanically smaller than one measured over ten.

| system                                                | data ( $n$ )       | top-1 | top- $k$ ( $k$ ) | recall        | prec.         |
|-------------------------------------------------------|--------------------|-------|------------------|---------------|---------------|
| <i><b>scored on connectivity</b></i>                  |                    |       |                  |               |               |
| this work, solver alone                               | literature (194)   | 28.4% | 33.5% (3)        | 33.5%         | 84.6%         |
| this work, + forward-verification                     | literature (194)   | 29.9% | 33.5% (3)        | 33.5%         | <b>89.2%</b>  |
| NMR-Solver <sup>27</sup>                              | literature (450)   | 52.9% | 67.3% (10)       | $\geq 67.3\%$ | $\leq 78.6\%$ |
| NMRAgent <sup>28</sup>                                | literature (450)   | 61.6% | 70.0% (10)       | $\geq 70.0\%$ | $\leq 88.0\%$ |
| <i><b>scored with stereochemistry</b></i>             |                    |       |                  |               |               |
| this work                                             | literature (194)   | 21.1% | 25.8% (3)        | $\geq 25.8\%$ | $\leq 81.8\%$ |
| NMR-Solver <sup>27</sup>                              | simulated (1,000)  | 66.9% | 89.9% (10)       | $\geq 89.9\%$ | $\leq 74.4\%$ |
| NMR-Solver <sup>27</sup>                              | literature (450)   | 31.6% | 53.8% (10)       | $\geq 53.8\%$ | $\leq 58.7\%$ |
| Espejo Morales <sup>30</sup>                          | education (236)    | 80.9% | 90.0% (5)        | $\geq 90.0\%$ | $\leq 89.9\%$ |
| Espejo Morales <sup>30</sup>                          | industrial (34)    | 20.6% | 29.1% (5)        | $\geq 29.1\%$ | $\leq 70.9\%$ |
| <i><b>exact match, stereo handling not stated</b></i> |                    |       |                  |               |               |
| Alberts, 6–13 atoms <sup>18</sup>                     | NIST IR (3,455)    | 63.8% | 84.0% (10)       | 84.0%         | 75.9%         |
| Alberts, 5–35 atoms <sup>18</sup>                     | NIST IR (5,024)    | 59.9% | 78.5% (10)       | 78.5%         | 76.4%         |
| SpecX, random split <sup>20</sup>                     | simulated (99,439) | 59.0% | 81.8% (10)       | 81.8%         | 72.2%         |
| SpecX, scaffold split <sup>20</sup>                   | simulated (99,439) | 29.7% | 50.6% (10)       | 50.6%         | 58.7%         |
| IR-Agent <sup>15</sup>                                | NIST IR (905)      | 10.3% | 21.6% (10)       | 21.6%         | 47.7%         |
| <i><b>criterion not stated</b></i>                    |                    |       |                  |               |               |
| Priessner <sup>16</sup>                               | experimental (34)  | 20.6% | —                | 26.5%         | 77.8%         |

**Three groups have already run the control this paper needed.** Hold the system, the criterion and the candidate budget fixed and change only the data. Since top-1 = recall  $\times$  precision, each term’s share of the resulting collapse is its share of the log ratio; reading the absolute gap  $b - a$  instead would mislead, because that gap is bounded by  $b$  and so shrinks whenever recall does. NMR-Solver moves from simulated spectra to real literature ones and loses 35 points of top-1, of which **recall carries 68%**. SpecX’s transformer moves from a random split to a scaffold split: **70%**.

Espejo Morales et al.’s agent moves from a chemistry-education set to industrial samples: **83%**. Ranking degrades in all three — relative loss 25.6%  $\rightarrow$  41.3%, 27.8%  $\rightarrow$  41.3%, 10.1%  $\rightarrow$  29.1% — so this is not a claim that verification is unaffected. It is a claim about which term dominates, and in every pair it is the same one. None of the three papers reports this, and none was testing it.

**Published gains have come almost entirely from the supply side.** NMRAgent’s own ablation runs 36.89  $\rightarrow$  60.47  $\rightarrow$  64.27  $\rightarrow$  64.42: retrieval and de-novo generation account for 27.38 of the 27.53-point span, and the verifier-driven optimisation loop that earns the paper the word *evidential* is worth 0.15 points<sup>28</sup>. NMRGym’s formula ablation moves recall at  $k=10$  from 17.37% to 36.48% while conditional precision moves only from 36.1% to 49.1% — the constraint buys candidates, not judgment<sup>19</sup>. Even IR-Agent’s agentic layer adds 4.0 points of recall against 0.5 points of top-1<sup>15</sup>.

**The regime has a boundary.** Since  $-\log(\text{top-1}) = -\log(\text{recall}) - \log(\text{precision})$ , recall’s share of a system’s total loss is  $\log(\text{recall})/\log(\text{top-1})$  — scale-free, and unlike a miss-to-ranking ratio not driven by where the reported  $k$  happens to fall. Among the rows where recall is exact that share is 38% on SpecX’s simulated random split and 39% on Alberts’ NIST gas-phase library, where ranking is therefore the larger term; 56% on SpecX’s scaffold split; 67% for IR-Agent on experimental NIST spectra; and 87–91% on ours. What separates the low end from the high is not method sophistication — the two extremes are both trained transformers — but how far the evaluation data are from one curated distribution.

Candidate budget confounds this, and pushes it our way: recall measured over three candidates is mechanically smaller than recall over ten, which raises our share. Our own generate-wide arm supplies the test, since it widens the budget on the same compounds — at recall 42% and top-1 30% the share is 72%, down from 91% but still far above the 38–39% of the simulated and single-library rows (§5.3). The ordering survives the confound; the magnitude should be read with it in mind. Ranking is co-equal or dominant where spectra are simulated or drawn from one curated instrument library with the formula supplied; recall binds where they are real and heterogeneous. The clearest counterexample is Priessner et al., whose reasoning-LLM re-ranker moves top-1 from 41.2% to 67.6% on a fixed candidate pool<sup>16</sup> — 26 points bought purely by re-ordering. That is their condition in which the true structure seeds the pool, which is not the row tabulated above and is why its recall is high by construction; the tabulated row is the harder, analogue-seeded one.

**One number the table explains with our own measurement.** NMR-Solver’s conditional precision is bounded at 78.6% — below our pooled 89.2%, though above the 71.5% we get under corpus weights, and we say so because §3 calls the reweighted figure the better estimate. Either way the ordering is not what a sharper predictor predicts: theirs is roughly twice as sharp as ours. The explanation is pool width rather than predictor sharpness, and we measured it directly: candidates for one target are predicted a median 1.21 ppm apart, inside the predictor’s own error for 82.1% of isomeric pairs (`scripts/isomer_separability.py`), so every candidate added to a pool adds a near-degenerate competitor. Our own ladder shows the same tension from the inside — verification precision falls from 84% to 72% as the pool widens (§5.3) — and NMR-Solver ranks a documented 1,000-molecule pool where we rank three. Conditional precision and recall trade against each other along the candidate budget, which is why neither number means much without the other, and why we report the budget with both.

**What the field does not measure.** Across every paper we read, the frequency with which the true structure was proposed at all appears once, in a supplementary figure caption<sup>16</sup>. Espejo Morales et al. declare a ten-candidate output and report  $k = 1, 2$  and 5<sup>30</sup>; NMRAgent’s verifier emits a `need_bigger_pool` verdict and the paper never reports how often it fires<sup>28</sup>. One collision is

worth naming: NMR-Solver defines *recall@K* as the proportion of cases where the correct structure appears among the top  $K$  candidates<sup>27</sup>, so its “52.89% recall” is a top-1 accuracy and is not the quantity we call recall. Setting our 33.5% beside it would compare incommensurable numbers.

## 7. Discussion

The Claude-family models we tested are reliable scaffold-level elucidators and good *verifiers* (verification precision 89%); exact top-1 is throttled by generation recall and by the regiochemical underdetermination intrinsic to 1D NMR. The contribution is a diagnosis with a bounded, training-free improvement attached, not a method that solves the task.

The diagnosis held under every perturbation we were able to apply — four Claude models, a second chemical domain, four verifiers — and reproduces inside a single application domain (§4.5): the binding constraint is *generation recall*, consistent with the bottleneck being structural rather than incidental.

It also has a boundary, and reading the published literature through the same lens locates it (§6). Recall binds where the spectra are real and heterogeneous; ranking is co-equal or dominant where they are simulated, or drawn from one curated instrument library. Among the systems whose recall can be read exactly from their own figures, recall’s share of the total loss runs from 38–39% on a simulated random split and a NIST gas-phase library to 87–91% on ours. So “recall, not verification” is a statement about a regime rather than about the task, and the sharpest evidence for it is not ours: three groups changed nothing but their data, and in each case recall carried the larger share of the collapse — 68%, 70% and 83%, with ranking taking the rest. We measured the improvement’s ceiling rather than projecting past it: forward-verification moves top-1 from 28% to 30% across the whole benchmark (§5.2), while recall plateaus at 42% (§5.3).

This reframes the engineering problem, and not in the direction of “do not train a model”. Purpose-trained systems win on this task today and we report it: NMR-Solver at 52.9% and a trained IR transformer at 63.8% against a training-free 28–30%, and our own 16M-parameter generator lifts pooled recall to 54.1% and top-1 to 35.1% (§5.6). The claim is narrower and survives that. What ages out each model generation is a *particular* trained artefact; what compounds is the open, experimental data any such artefact needs, and a benchmark honest enough to tell whether it worked. The probe is the evidence: on its own 248-structure held-out split the same architecture recovers 0/248 without IRexp and 25% with it, so the released data, rather than the architecture, is the active ingredient. Inference-time scaffolding is the second durable lever for the same reason — it needs no training, so it rides each new model rather than being replaced by it.

Concurrent work sharpens both halves. Kliuev *et al.* independently place the best of six frontier LLMs at 24% on 105 experimental  $^1\text{H}/^{13}\text{C}$  peak lists<sup>29</sup>, within a few points of our 28%, so the headline is not an artefact of our corpus or our model. Espejo Morales *et al.* then push the lever our decomposition names as binding harder than we do: environment design alone — processing tools, tabulated shifts, hard constraint checks, and no shift-matching ranker — carries a frozen LLM to 71.1% top-1 on the 15-molecule set where graduate students reach 66.7%, against 24.4–40.0% for frontier LLMs given the same inputs in a bare coding harness<sup>30</sup>. That movement is entirely on the proposal side, where we locate the headroom; what it does not report is the split that locates it. And their own 80.9% on curated education spectra against 20.6% on 34 industrial samples is our central empirical claim, measured by someone else.

Protocol is a lever of the same order as capability: a number quoted without its protocol is uninterpretable. We do *not* claim protocol dominates capability — the model axis spans 0% to 46% top-1 on the fixed 24-compound subset (Table 3), wider than the 5%  $\rightarrow$  15% protocol effect in §4.3, and the two are measured on different sets. In molecular property prediction, a systematic benchmark reports that learned, pretrained molecular embeddings rarely outperform classical ECFP fingerprints once evaluation is controlled<sup>54</sup>; we read our forward-verification recipe (§5) in that light, and draw the parallel narrowly.

Two findings transfer to practice: solve each problem in bounded, frequently-reset contexts with tool access (5%  $\rightarrow$  15%, 1/20  $\rightarrow$  3/20, directional at n=20, McNemar  $p \geq 0.25$ , §4.3); and use forward-predicted-vs-observed <sup>13</sup>C agreement to *re-rank* candidates rather than to decide whether to trust the winner. Match distance failed as an absolute confidence gauge in our selective-prediction test (§5.5, a reported null result), so it does not support abstention thresholds.

## 8. Limitations

Several limitations of earlier drafts are now resolved; we state what remains plainly.

**Resolved. Extraction noise:** an RDKit self-consistency audit finds 57/60 ground truths spectrally clean, with metrics unchanged on that subset. **Verifier sample size:** forward-verification now runs over all 194 compounds, so the recall-conditional claim rests on n=65 rather than n=19 (§5.2). **Verifier precision:** precision degrades as near-degenerate regioisomers accumulate (§5.3), so forward-match distance is a strong *re-ranker* but a soft *confidence* gauge. **Non-LLM verifiers** (§5.4): a HOSE-code *lookup* does not beat the LLM verifier (verification precision 85% against 89% at n=65) and a *learned* GNN reaches 91% — directional but not statistically resolved at n=65, so it is suggestive that the deterministic failure was method rather than coverage alone, not a demonstration of it. Beyond all of them lies the near-degenerate-regiochemistry precision ceiling, where DFT-level accuracy or 2D-NMR constraints are the genuine fix. **Projection:** §5.3 now measures what an earlier draft estimated: top-1 30%, recall 42%, below the estimate it replaces, reported as found rather than adjusted.

**Independence checks.** Scoring is mechanical RDKit rather than LLM-judged, and all solver/verifier runs were transcript-audited *at generation time* for zero web and zero ground-truth access. That audit is the one part of the pipeline a reader cannot re-verify from the release: the committed artefacts are the parsed per-compound predictions, and the transcripts are not deposited (they are available on request). The *outcome* of closed-book solving is separately testable: the formula-only control (§4.6) removes the spectra and accuracy collapses, and a permutation negative control (§5.5) collapses verification precision from 89.2% to a chance mean of 73.8% (two-sided  $p=0.002$  on the full benchmark), as a correctly isolated blind evaluation requires. A second Claude model (Sonnet) was comparably recall-bound on a 12-compound subset (recall 33% vs Opus 42% on those compounds), consistent with the recall bound being a property of the task rather than of one model — but within the Claude family, so it speaks to model-instance robustness, not cross-vendor generality.

**Remaining. (i) Pretraining contamination is bounded; we cannot exclude it.** Every benchmark compound was mined from open-access literature, so a frontier model may have encountered it in training. The formula-only control (§4.6) bounds how much of the headline number pure recall can explain: masking the spectra drops top-1 from 23% to 5%, perfectly nested, McNemar exact  $p=0.001$ . It examines the three formula-only successes individually rather than collectively — only

one (an unlisted silyl aryl sulfonate) has a genuinely determining composition, while *N*-tosyl-leucine and [6]-paradol are catalogued compounds for which recall is a live explanation — which leaves the bound unchanged. A second control finds accuracy flat in source publication year across all 194 compounds ( $r=-0.007$ ), the size-adjusted older-versus-newer difference bounded at  $-5.1$  points, 95% CI  $[-17.2$  to  $+7.0]$ . Neither control is randomised: publication year is observational, and the formula-only arm cannot separate memorisation from inference on a near-determining formula. Nor is the recency test anchored to a *known* training cutoff, since the harness does not disclose one (§9); we test recency as a continuous proxy. A replication restricted to compounds published after a disclosed cutoff remains open.

The sharper form of this objection is one neither control addresses, and we state it plainly because a reader is entitled to weigh it. The prompt carries the spectral strings *exactly as printed in the source article*, down to its own typography, because the sampler keeps a compound only when the raw  $^1\text{H}$  payload can be recovered verbatim from the open-access full text (§9). Every benchmark compound therefore has its spectrum present verbatim in a document the model may have read, and that string is a high-entropy fingerprint of the document. Masking the spectra therefore removes the chemistry and a retrieval key together: a collapse from 23% to 5% is what reading predicts, and also what retrieval predicts. Recency cannot separate them either, for the reason just given. The experiment that would be a spectrum meaning the same thing and reading differently — shifts perturbed within reporting precision and typography normalised, so the chemistry is untouched and no verbatim string survives. `scripts/jitter_control.py` builds that set from the released benchmark; running it is the obvious next control and we have not run it. It would bound verbatim retrieval only, not recognition from approximate shift patterns.

**(ii) Human audit — prepared but not yet run.** Two things are unvalidated by hand: the elucidation and forward-prediction outputs, and the *recall* side of dataset extraction (whether the parser found every IR string in every source paper). Transcription fidelity of the records we hold is measured and high (§2.3); record-level recall is not, and only reading papers can settle it. Solver and verifier are both LLMs, so the one validation we cannot perform ourselves is expert-chemist review. We have therefore built and frozen a blinded, pre-registered audit package at `data/audit/`, protocol at `docs/EXPERT_AUDIT_PROTOCOL.md`: a difficulty-stratified 30-compound elucidation panel plus seven ranking-only compounds, with a withheld answer key, testing what a miss actually is (§4) and whether forward verification is a trustworthy re-ranker (§5). We have not yet run the panel; until expert results are in, those claims should be read as machine-validated (RDKit InChIKey) but not yet human-validated. §4 has since narrowed its first question by reporting all 137 analysable misses mechanically (76.6% constitutional isomers, 22.6% scaffold-preserving positional errors), correcting the impressionistic claim the audit was built to check; what remains for a chemist is whether a formula-correct, scaffold-wrong candidate is a *chemically reasonable* reading of the spectra.

**(iii) Single vendor and an underpowered cross-model comparison.** The headline (28.4% top-1,  $n=194$ ) is a single frontier model (Claude Opus) on the full benchmark, and the cross-model evidence (§4.4) is four Claude-family models on a shared 24-compound subset. That comparison is underpowered: the ranking is robust (strictly nested, weakest floored at 0%) but no adjacent gap is established at  $n=24$ , so quantitative claims should be read as holding for the Claude family. The cross-vendor question is narrowed rather than closed: §4.7 finds verification precision exceeding generation recall for Grok 4.6, Gemini 3.7 Flash and GPT-5.6 Sol on the 60-compound arm, as for Claude. What remains is weaker than the original gap but real. Bootstrapping the paired difference within each arm resolves three of the four — Claude, GPT-5.6 Sol and Gemini — and leaves Grok’s directional at  $n=60$  (§4.7). Those models were reached through a coding-assistant harness

exposing reasoning-effort tiers, a decoding control our own arm never had — and the tier serving each arm was not recorded, so the arms are not demonstrably matched to each other either, which bounds the recall *ranking* while leaving each arm’s inequality internally valid (§4.7). The clean-clone contamination control was run for Grok alone, and the cross-vendor arms cover 60 compounds.

**(iv) Domain-subset scope:** the battery-electrolyte case study (§4.5) uses *literature compounds bearing electrolyte-relevant functional chemistry*, not operando or in-cell decomposition spectra, so it demonstrates functional-class transfer of the elucidation bottleneck, not direct assignment of authentic interphase/degradation products, which would require a dedicated operando set.

**(v) Constitution-only scoring:** correctness is judged on InChIKey connectivity, so a correct-constitution / wrong-stereochemistry prediction counts as correct. Scoring the full InChIKey gives 21.1% top-1 (§3), 7.3 points lower, so the headline should be read as an upper bound on full-stereochemistry accuracy. We keep constitution because 1D NMR/IR rarely determines absolute configuration — the 10.3% (20/194) of targets with a defined stereocentre would be penalised for information the prompt never carried — but a stereochemistry-sensitive benchmark would need 2D/chiroptical data.

**(vi) Single-sample scoring:** each headline compound is scored from one solver prediction set, so reported top-1/recall carry no run-to-run (LLM-sampling) variance estimate; the bootstrap CIs reflect compound sampling only. §5.3 pools ten independent generation passes (recall 32% → 42%), indicating generator stochasticity that single-pass scoring understates.

**(vii) Chemical-space coverage:** the benchmark is drawn from open-access organic methodology and total-synthesis literature; it is not representative of organometallic/coordination compounds, large biomolecules (peptides, oligonucleotides, oligosaccharides), or stereochemistry-heavy targets, and our results should not be extrapolated to them.

**(viii) Input regime:** IRSpectra-Bench supplies peak lists as transcribed by the source paper’s authors, not raw instrument files. Concurrent work excludes author-transcribed reports for exactly that reason — they are human-generated and vary by practitioner<sup>30</sup> — and reports 20.6% top-1 on 34 industrial samples processed from raw spectra, scored with stereochemistry retained. Our comparable figure is the full-InChIKey 21.1%, not the 28.4% constitution-only headline (§3). The two settings measure different things: reading the published record across thousands of laboratories, and inference from instrument output in one. Neither bounds the other.

---

## 9. Methods

**Mining and resolution.** PMC-OA full text was fetched from `s3://pmc-oa-opendata`, parsed deterministically, and resolved with OPSIN, RDKit and SELFIES, with an optional cached PubChem fallback.

**Benchmark and agents.** Problems were sampled from `irexp_resolved`, stratified by RDKit ring analysis to half the round per stratum, and de-duplicated across rounds by InChIKey. A compound was admitted only if its raw <sup>1</sup>H payload — multiplicities and *J* values as printed — could be re-extracted verbatim from the PMC-OA full text of its source article (`benchmark_v2.raw_1h_for`), which is what puts genuine coupling information in the prompt and, unavoidably, puts the source article’s exact string there too (§8). The battery-electrolyte subset (§4.5) was drawn from the same corpus by SMARTS filters for six electrolyte functional classes (carbonate, sulfonyl/sulfonate,

nitrile,  $\text{sp}^3\text{-C-F}$ , phosphoryl, glyme/oligoether), balanced to eight compounds per class (48 curated; 46 scored after two yielded no parseable candidate) and excluding every compound used elsewhere. It was  $J$ -enriched — the raw  $^1\text{H}$  string keeps its multiplicities and  $J$  values — and spectrally validated, as in the main rounds. Solver and forward-prediction agents were independent Claude-Opus sub-agents under the same subscription, instructed closed-book and audited by automated transcript search for zero web/answer access. A prediction was scored correct on RDKit InChIKey-connectivity match; similarity used Morgan(2, 2048) Tanimoto, and forward-verification a symmetric chamfer distance over  $^{13}\text{C}$  peak sets. The core protocol trains no model and uses no paid API; the §5.6 generator and the §5.4 learned  $^{13}\text{C}$  verifier are the only two trained components, reported as complements and fenced from the headline results.

*Models and versions.* All experiments used Anthropic Claude models via the consumer subscription (claude.ai). The §4.4 comparison spanned, in capability order, Claude Haiku, Claude Sonnet, Claude Opus and Claude Fable 5; the headline benchmark (§4.1) and forward-verification (§5) used Claude Opus. `docs/MODELS.md` records, per experiment, the model, data directory, collection date, harness and tool access, and scoring code path, across three dated windows:

- **2026-06-09 to 2026-06-11** (UTC) — every candidate structure behind the headline results: §4.1, §4.3–§4.5, and the candidate pools all of §5 re-ranks.
- **2026-07-28** — the formula-only contamination control (§4.6), re-solving the same 60 compounds with spectra masked and generating its own candidates; it touches Table 5 alone.
- **2026-08-07** — three collections forward-predicting  $^{13}\text{C}$  for candidates the June run had already produced (the §5.2 extension to all 194 compounds, the §5.3 coverage-gap closure, and the §5.6 re-run); none introduced a new candidate structure, and no recall number in the paper changes.

One *verified* number does: the §5.6 re-run supersedes a top-1 whose original forward-prediction outputs were never deposited and are lost, disagreeing with the figure it replaces (47% against 41%); §5.6 reports it as a re-run rather than a reproduction.

**No model snapshot can be reported.** The consumer harness exposes no checkpoint identifier and records nothing about which build served a given request. It exposes no decoding parameters either, and none were set or recorded. Claude Opus 4.8 is evidenced only for the 2026-06-09 pilot, so a mid-window build change cannot be excluded, and we do not claim the same build served the main round two days later. Re-running reproduces the protocol distributionally rather than exactly, and exact agreement is not offered as a target. Fixed are the dated collection windows, the frozen per-compound outputs, and mechanical scorers that regenerate every number from them: exact reproducibility of scoring and analysis, not of inference (`docs/MODELS.md` §7).

**Reproducibility.** Every round is frozen: questions, ground-truth answers, per-agent raw outputs, predictions and scorer outputs are released; the sampler, scorer and forward-verification harness are scripted end-to-end.

---

## Supporting Information figures

These are supplied as a separate Electronic Supplementary Information document (`docs/paper_esi.pdf`, built by `scripts/build_pdf.py`), as RSC requires; they are listed here for reference.

- **Fig. S1** (docs/figures/fig0\_overview.png) — study design: open multimodal data (IRexp) → blind, complexity-stratified benchmark → decoupled blind solving → forward-verification re-ranking; training-free core pipeline.
- **Fig. S2** (docs/figures/fig4\_dataset.png) — IRexp composition: IR records → NMR-paired → structure-linked → full IR +  $^1\text{H}$  +  $^{13}\text{C}$  + structure quadruples.
- **Fig. S3** (docs/figures/fig2\_size.png) — accuracy vs molecular size (heavy-atom bucket); the monotonic 60% → 7% top-1 gradient.
- **Fig. S4** (docs/figures/fig6\_electrolyte.png) — top-1 and recovered accuracy on IRSpectra-Bench-Electrolyte by battery-electrolyte chemical class (n=46):  $\text{sp}^3\text{-C-F}$  easiest (50%), sulfonyl and nitrile hardest (12%); overall 26%/28%.
- **Fig. S5** (docs/figures/fig\_generator\_probe.png) — trained-generator probe (§5.6; a complement, not part of the training-free protocol): generation recall and deterministic-HOSE top-1 on the 194-compound benchmark for Claude / + scaffold enumeration / + trained generator. Enumeration’s near-degenerate isomers collapse the verifier (28.4 → 16.0%) while the generator’s formula-correct candidates convert (28.4 → 35.1%).
- **Fig. S6** (docs/figures/fig\_verifier.png) — learned-verifier probe (§5.4; a complement, not part of the training-free protocol). (A) Conditional-on-recall top-1 (n=65, whole benchmark) for the four verifiers: a GNN trained on the same nmrsiftdb2 data as the HOSE lookup reaches the LLM verifier’s level (91% vs 89%) where the lookup (85%) does not move off the solver’s own ranking. (B) Why: held-out  $^{13}\text{C}$  MAE — the learned model is  $\approx 2\times$  sharper (1.70 vs 3.23 ppm).

---

## Author contributions

**I.Y.:** conceptualization, methodology, software, formal analysis, investigation, data curation, visualization, writing — original draft. **R.S.:** methodology, software, formal analysis, investigation, validation (trained-generator and learned-verifier probes, §5.4 and §5.6), writing — review and editing. **R.A.V.-H.:** conceptualization, methodology, supervision, writing — review and editing.

## Conflicts of interest

There are no conflicts to declare.

## Data availability

Table 11 lists every released component and the script that regenerates it, all of them in the project repository. The archival deposit — a complete frozen snapshot of dataset, benchmark, answer keys, predictions, scripts, figure regeneration and the expert-audit package — is deposited on Zenodo, and its DOI is supplied at proof stage; GitHub is the development mirror and the Zenodo record the citable version.

1000 **Table 11. Released artefacts, and the data and code behind each one.**

| component                                                                                                            | data and code                                                                                                                                                                                                                |
|----------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| IRexp and its structure-complete split benchmark rounds and the within-compound control ground-truth integrity audit | data/irexp/, data/irexp_resolved/ — spectro_scraper/ (mining pipeline)<br>data/benchmark*/ — scripts/benchmark_v2.py                                                                                                         |
| headline accuracy (Table 2)<br>battery-electrolyte subset (§4.5)                                                     | data/benchmark*/clean_qids.json — scripts/validate_benchmark.py<br>scripts/score_main.py<br>data/benchmark_electrolyte/ — scripts/build_electrolyte_bench.py, scripts/score_electrolyte.py                                   |
| formula-only contamination control (§4.6)                                                                            | data/modality/ — scripts/modality_ablation.py                                                                                                                                                                                |
| publication-recency control (§4.6)                                                                                   | data/audit/recency_control.json — scripts/contamination_recency.py                                                                                                                                                           |
| forward-verification, original arm (§5.2)                                                                            | data/fverify/ — scripts/forward_verify.py                                                                                                                                                                                    |
| ...extended to all 194 compounds                                                                                     | data/fverify_main/ — scripts/forward_verify_main.py, scripts/forward_verify_all.py                                                                                                                                           |
| generate-wide arm (§5.3)                                                                                             | data/gw/, data/fverify2/ — scripts/score_generate_wide.py, scripts/ladder_significance.py                                                                                                                                    |
| ...its coverage gap, closed non-LLM verifier comparison (Table 9)                                                    | data/fverify_gw/ — scripts/forward_verify_gw.py<br>data/fverify/hose_results.txt, data/fverify/verifier_table_results.txt — scripts/hose_predict.py (incl. coverage), scripts/verifier_table.py, scripts/verifier_leakage.py |
| negative control and selective prediction (§5.5)                                                                     | scripts/verifier_diagnostics.py                                                                                                                                                                                              |
| what a miss is, and isomer separability (§4.1, §5.1)                                                                 | scripts/analyze_misses.py, scripts/isomer_separability.py                                                                                                                                                                    |
| difficulty-threshold sensitivity (§3)                                                                                | scripts/difficulty_sensitivity.py                                                                                                                                                                                            |
| licence-pool split (§2.1)                                                                                            | scripts/split_license_pools.py                                                                                                                                                                                               |
| recall headroom and scaffold enumeration                                                                             | scripts/analyze_recall_headroom.py, scripts/enumerate_isomers.py, scripts/closing_the_gap.py                                                                                                                                 |
| blinded expert-audit package (§8)                                                                                    | data/audit/ — scripts/make_audit_sample.py                                                                                                                                                                                   |
| corpus reweighting (§3)                                                                                              | scripts/corpus_reweight.py                                                                                                                                                                                                   |
| cross-vendor recall at matched budget, and the paired gap (§4.7)                                                     | data/cross_vendor/ — scripts/cross_vendor_budget.py, scripts/cross_vendor_gap.py                                                                                                                                             |
| ring-system names in peak assignments (§3)                                                                           | scripts/prompt_leakage.py                                                                                                                                                                                                    |
| paraphrase-invariant benchmark, for the control not yet run (§8)                                                     | — (regenerated, not deposited) — scripts/jitter_control.py                                                                                                                                                                   |

| component                  | data and code                                                                                                                                                           |
|----------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| manuscript integrity gates | <code>scripts/check_manuscript.py</code> , <code>scripts/verify_statistics.py</code> , <code>scripts/check_compression.py</code> , <code>scripts/check_layout.py</code> |

**Trained complements**, both fenced from the training-free protocol. The §5.6 generator ships as a self-contained bundle at `contrib/generator_probe/` — candidates, the de-leaked split with its InChIKey-14 manifest, the verification scripts (`scripts/closing_the_gap_gen.py`, `scripts/forward_verify_gen.py`, `scripts/verify_leakage_exact40.py`) and the blind forward-prediction outputs behind its verified top-1 (`data/fverify_gen/`), so its scorer runs with no missing predictions. The §5.4 learned verifier ships `scripts/gnn_predict.py` (extract / train / score / control), the trained model `data/nmrshiftdb/gnn_c13.pt`, per-compound results and both leakage checks (`data/fverify/gnn_results.txt`), and the write-up `docs/VERIFIER_PROBE.md`. Model checkpoints are deposited on Zenodo. Both trained predictors of §5.4 need the `nmrshiftdb2` dump, which we cannot redistribute; the README gives the one-line fetch.

Companion technical notes: `docs/BENCHMARK.md`, `docs/FORWARD_VERIFY.md`, `docs/MODALITY_ABLATION.md`, `docs/EXPERT_AUDIT_PROTOCOL.md`, `docs/MODELS.md`.

**Two cautions for re-users.** The public `irexp_release/train` split does not hold out IR Spectra-Bench — it overlaps by 117/200 InChIKey-14 — so de-leak with `contrib/generator_probe/build_exp_manifest.py` before training anything that will be evaluated on the benchmark. And the held-out answer keys `data/audit/key.jsonl` and `data/modality/key.json` are deposited but flagged *withhold from blinded reviewers*.

## Licensing and attribution

IRexp is derived from two open-access source pools and redistributed under terms compatible with each (see `data/NOTICE`). **(a) Redistribution:** we release only *extracted numeric data* — IR band lists,  $^1\text{H}/^{13}\text{C}$  shift lists, and resolved structures (SMILES/SELFIES/InChIKey) — plus each record’s source DOI/accession, not source full text, figures, or PDFs. **(b) Two separable pools:** 119,345 records derive from the PMC Open-Access Subset (CC-BY-4.0) and 1,888 from the Chemotion RADAR4Chem FT-IR deposit (CC-BY-SA-4.0). The two are separable losslessly by `source_doi` — Chemotion records carry the 10.22000 prefix — and `scripts/split_license_pools.py` materialises them as two files with each record stamped license, so users may take the CC-BY pool alone; any combined or Chemotion-derived release carries CC-BY-SA-4.0 to honour the ShareAlike term. Code is released under the MIT License. **(c) Attribution:** re-users must cite this dataset (Zenodo DOI above) and attribute the original publications via each record’s `source_doi`.

## Acknowledgements

### Use of AI tools

This work studies a large language model, and LLMs were also used as instruments and as writing aids; we state both roles explicitly. **As an object of study:** all reported elucidation, forward-prediction and verification results were produced by Claude models invoked under the protocol of §3 and §9 — these are the measurements the paper reports, not assistance in producing it. **As**

**a writing aid:** the authors used an LLM-based coding assistant for figure generation, analysis scripting, manuscript copy-editing, and for the internal review pass that produced several of the corrections recorded in the repository history. No text, number, figure or citation in this manuscript was accepted without author verification against the released data and code; every quantitative claim regenerates from the scripts in `scripts/`. The authors take full responsibility for the content.

## References

- 1 E. Chacko, R. Sondhi, A. Praveen, K. L. Luska and R. A. Vargas-Hernández, *ChemRxiv*, 2024, preprint, DOI: 10.26434/chemrxiv-2024-37v2j.
- 2 D. Kamber, *Making Claude a chemist: How Claude performs on NMR prediction and structure elucidation*, Anthropic, 2026.
- 3 T. Han, S. Wu, J. Wang, Y. Zhou, R. Lv, B. Zhao and W. Hu, *arXiv*, 2026, preprint, DOI: 10.48550/arXiv.2603.25253.
- 4 F. Hu, M. S. Chen, G. M. Rotskoff, M. W. Kanan and T. E. Markland, *ACS Cent. Sci.*, 2024, **10**, 2162–2170.
- 5 L. Yang, Z. Yang, J. Xie, Y. Wang, B. Gao, X. Wei, J. Sun, J. Wu, C. He, Y. Li and Q. Gu, in *Proc. 32nd ACM SIGKDD Conf. Knowledge Discovery and Data Mining (KDD '26)*, Vol. 2, 2026, pp. 12621–12632.
- 6 F. Ottomano, Y. Li and A. M. Ganose, *arXiv*, 2025, preprint, DOI: 10.48550/arXiv.2512.19733.
- 7 U. Garcilazo-Cruz, Q. Nie, R. Sondhi and R. A. Vargas-Hernández, *ChemRxiv*, 2026, preprint, DOI: 10.26434/chemrxiv.15007479/v1.
- 8 R. Sondhi, E. Chacko and R. A. Vargas-Hernández, *ChemRxiv*, 2025, preprint, DOI: 10.26434/chemrxiv-2025-d0j2v.
- 9 A. M. Bran, S. Cox, O. Schilter, C. Baldassari, A. D. White and P. Schwaller, *Nat. Mach. Intell.*, 2024, **6**, 525–535.
- 10 D. A. Boiko, R. MacKnight, B. Kline and G. Gomes, *Nature*, 2023, **624**, 570–578.
- 11 Y. Su, J. Chen, Z. Jiang, Z. Zhong, L. Wang, Q. Liu and Z. Zhang, in *International Conference on Learning Representations (ICLR)*, 2026.
- 12 S. Shen, J. Xie, Z. Yang, A. Zhang, S. Sun, B. Gao, T. Fu, B. Qi and Y. Li, *arXiv*, 2025, preprint, DOI: 10.48550/arXiv.2509.21861.
- 13 X. Zhuang, B. Wu, J. Cui, K. Feng, X. Li, H. Xing, K. Ding, Q. Zhang and H. Chen, in *Proc. 63rd Annu. Meet. Assoc. Comput. Linguist. (Vol. 1: Long Papers)*, 2025, pp. 22561–22576.
- 14 K. Guo, B. Nan, Y. Zhou, T. Guo, Z. Guo, M. Surve, Z. Liang, N. V. Chawla, O. Wiest and X. Zhang, in *Advances in Neural Information Processing Systems 37 (NeurIPS 2024), Datasets and Benchmarks Track*, 2024, pp. 134721–134746.
- 15 H. Noh, N. Lee, G. S. Na, K. Kim and C. Park, *arXiv*, 2025, preprint, DOI: 10.48550/arXiv.2508.16112.
- 16 M. Priessner, R. J. Lewis, M. J. Johansson, J. M. Goodman, J. P. Janet and A. Tomberg, *Digital Discovery*, 2026, **5**, 1237–1251.
- 17 M. Alberts, T. Laino and A. C. Vaucher, *Commun. Chem.*, 2024, **7**, 268.


- 18 M. Alberts, F. Zipoli and T. Laino, *Digital Discovery*, 2025, **4**, 1936–1943.
- 19 Z. Fang, C. Yang, H. Yu, H. Luo, H. He, J. Xie, Z. Yang and J. Xia, *arXiv*, 2026, preprint, DOI: 10.48550/arXiv.2601.15763.
- 20 C. Xiang, T. Ma, Y. Chen, T. Wang, H. Chen and X. Zeng, *arXiv*, 2026, preprint, DOI: 10.48550/arXiv.2605.18791.
- 21 S. G. Smith and J. M. Goodman, *J. Am. Chem. Soc.*, 2010, **132**, 12946–12959.
- 22 N. Grimblat, M. M. Zanardi and A. M. Sarotti, *J. Org. Chem.*, 2015, **80**, 12526–12534.
- 23 C. J. Pickard and F. Mauri, *Phys. Rev. B*, 2001, **63**, 245101.
- 24 S. E. Ashbrook and D. McKay, *Chem. Commun.*, 2016, **52**, 7186–7204.
- 25 M. Elyashberg, K. Blinov, S. Molodtsov and A. Williams, *Magn. Reson. Chem.*, 2012, **50**, 22–27.
- 26 M. E. Elyashberg and A. J. Williams, *Computer-based structure elucidation from spectral data*, Springer Berlin Heidelberg, 2015.
- 27 Y. Jin, J.-J. Wang, F. Xu, X. Ji, Z. Gao, L. Zhang, G. Ke, R. Zhu and W. E, *Nat. Commun.*, 2026, **17**, 4740.
- 28 Z. Fang, C. Yang, Y. Tan, Y. Zhao, F. Xu, H. Xiang, H. Sun, H. Gao, X. Wang, W. Du, Y. Li and J. Xia, *arXiv*, 2026, preprint, DOI: 10.48550/arXiv.2606.29776.
- 29 F. Kliuev, I. Smirnov, M. A. Losev, O. I. Afanasyev and D. Chusov, *ChemRxiv*, 2026, preprint, DOI: 10.26434/chemrxiv.15006195/v1.
- 30 I. Espejo Morales, D. Hinz, M. Alberts, G. Krawezik, H. Jeong and S. Ho, *arXiv*, 2026, preprint, DOI: 10.48550/arXiv.2607.19406.
- 31 National Institute of Standards and Technology, NIST Chemistry WebBook, SRD 69, <https://webbook.nist.gov>, (accessed 1 January 2026).
- 32 N. Jung, P. Tremouilhac, D. Punjabi and P.-C. Huang, *Chemotion Repository – data collection: FT-IR spectroscopy data (Chemotion IR)*, Karlsruhe Institute of Technology, 2024.
- 33 National Institute of Advanced Industrial Science and Technology (AIST), Spectral Database for Organic Compounds (SDBS), <https://sdb.sdb.aist.go.jp>, (accessed 1 January 2026).
- 34 M. C. Swain and J. M. Cole, *J. Chem. Inf. Model.*, 2016, **56**, 1894–1904.
- 35 National Library of Medicine, NCBI PMC Open Access Subset, <https://www.ncbi.nlm.nih.gov/pmc/tools/openftlist/>, (accessed 1 January 2026).
- 36 D. M. Lowe, P. T. Corbett, P. Murray-Rust and R. C. Glen, *J. Chem. Inf. Model.*, 2011, **51**, 739–753.
- 37 G. Landrum and RDKit Contributors, RDKit, <https://www.rdkit.org>, (accessed 1 January 2026).
- 38 M. Krenn, F. Häse, A. Nigam, P. Friederich and A. Aspuru-Guzik, *Mach. Learn.: Sci. Technol.*, 2020, **1**, 045024.

- 1119 39 S. Kim, J. Chen, T. Cheng, A. Gindulyte, J. He, S. He, Q. Li, B. A. Shoemaker, P. A.  
1120 Thiessen, B. Yu, L. Zaslavsky, J. Zhang and E. E. Bolton, *Nucleic Acids Res.*, 2023, **51**,  
D1373–D1380.
- 1121 40 D. Rogers and M. Hahn, *J. Chem. Inf. Model.*, 2010, **50**, 742–754.
- 1122
- 1123 41 Q. McNemar, *Psychometrika*, 1947, **12**, 153–157.
- 1124
- 1125 42 S. Holm, *Scand. J. Stat.*, 1979, **6**, 65–70.
- 1126
- 1127 43 C. Xu, S. Guan, D. Greene and M.-T. Kechadi, *arXiv*, 2024, preprint, DOI:  
1128 10.48550/arXiv.2406.04244.
- 1129 44 D. Williamson, S. Ponte, I. Iglesias, N. Tonge, C. Cobas and E. K. Kemsley, *J. Magn.*  
1130 *Reson.*, 2024, **368**, 107795.
- 1131 45 C. Han, D. Zhang, S. Xia and Y. Zhang, *J. Chem. Theory Comput.*, 2024, **20**, 5250–5258.
- 1132
- 1133 46 F. Xu, W. Guo, F. Wang, L. Yao, H. Wang, F. Tang, Z. Gao, L. Zhang, W. E, Z.-Q. Tian  
1134 and J. Cheng, *Nat. Comput. Sci.*, 2025, **5**, 292–300.
- 1135 47 R. D. Cohen, J. S. Wood, Y.-H. Lam, A. V. Buevich, E. C. Sherer, M. Reibarkh, R. T.  
1136 Williamson and G. E. Martin, *Molecules*, 2023, **28**, 2449.
- 1137 48 X. Wang, J. Wei, D. Schuurmans, Q. Le, E. Chi, S. Narang, A. Chowdhery and D. Zhou,  
1138 in *International Conference on Learning Representations (ICLR)*, 2023.
- 1139 49 W. Bremser, *Anal. Chim. Acta*, 1978, **103**, 355–365.
- 1140
- 1141 50 S. Kuhn and N. E. Schlörer, *Magn. Reson. Chem.*, 2015, **53**, 582–589.
- 1142
- 1143 51 K. Adams, C. C. Wagen, J. Wolford, M. H. Sak, J. Saurí, Z. Feng, A. S. Bhadauria, M. A.  
Bailey, J. S. Downs, S. Z. Li, A. I. Liu, T. Smidt, R. S. Paton, R. Y. Liu, C. W. Coley and  
1144 E. E. Kwan, *ChemRxiv*, 2026, preprint, DOI: 10.26434/chemrxiv.15005989/v1.
- 1145 52 I. M. Novitskiy and A. G. Kutateladze, *J. Org. Chem.*, 2022, **87**, 8589–8598.
- 1146
- 1147 53 C. Wagen, *Simulation tools improve agent problem-solving*, Rowan Scientific, 2026.
- 1148
- 1149 54 M. Praski, J. Adamczyk and W. Czech, *arXiv*, 2025, preprint, DOI:  
1150 10.48550/arXiv.2508.06199.


## Table of contents entry

28% top-1 on blind, real IR + NMR spectra

observed  $^{13}\text{C}$ ,  $\text{C}_{10}\text{H}_{14}\text{N}_2\text{O}$



0.42 ppm selected



1.30 ppm rejected

recall (34%), not verification (89%), is the wall

Recall, not verification, is the wall. On real, blind IR +  $^1\text{H}$  +  $^{13}\text{C}$  literature spectra a frontier LLM proposes the true structure for 34% of 194 compounds and, once proposed, selects it 89% of the time — lifting top-1 from 28% to 30%.