

Electronic Supplementary Information

Recall, not verification, is the bottleneck when frontier LLMs elucidate molecular structures from real spectra

Ilkham Yabbarov, Rudra Sondhi, Rodrigo A. Vargas-Hernández

This document contains the Supporting Information figures referenced in the main article. Data, predictions, and the code that regenerates every figure are released with the manuscript; see *Data and code availability* in the main text.

This document records the operational detail behind the article: which model answered which arm and when (§S1), the prompts and agent protocol (§S2), how the 194-compound benchmark was assembled (§S3), what the correctness criterion accepts and rejects (§S4), the construction of the forward-verification arm (§S5) and of the two non-LLM verifiers (§S6), the four-vendor replication (§S7), the contamination controls (§S8) and the frozen expert-audit package (§S9). Where the repository does not pin a value it says so rather than supplying one. Every number is copied from a released artifact or from a companion note issued with the manuscript (docs/MODELS.md, docs/BENCHMARK.md, docs/FORWARD_VERIFY.md, docs/CROSS_VENDOR.md, docs/VERIFIER_PROBE.md, docs/MODALITY_ABLATION.md, docs/EXPERT_AUDIT_PROTOCOL.md).

S1. Models, versions and collection windows

Every LLM result in the article was produced by Anthropic Claude models invoked as independent sub-agents through the Agent tool under a single consumer claude.ai subscription: no paid API, no fine-tuning, no model training in the core protocol. The two trained probes — the learned ^{13}C verifier of §5.4 and the generator of §5.6 — are not Claude models (§S6, contrib/generator_probe/).

Table S1. Which model answered which arm, and when. “Collected” is the UTC author-date of the commit that first added the prediction artifact. The history is linear and each raw agent-output file was committed once, in the session that produced it, so first-add is the best available proxy for collection time; it is a proxy, not a timestamp emitted by the harness.

arm	model	data directory	collected (UTC)
pilot, n=21, single context, no tools	Claude Opus 4.8	data/benchmark/	2026-06-09 04:40
within-compound control, arm (a): n=20, one context, no tools	Claude Opus	data/benchmark_v2/	2026-06-09 06:28
controlled round v3, n=40, decoupled agents	Claude Opus	data/benchmark_v3/	2026-06-09 06:47–07:37
within-compound control, arm (b): same n=20, 4 agents of 5 compounds, RDKit formula check	Claude Opus	data/benchmark_v2_ctrl/	2026-06-09 17:43
forward verification: 8 blind forward-prediction agents, 126 candidates, 60 compounds	Claude Opus	data/fverify/	2026-06-09 19:51
generate-wide: 10 solver agents, 6 compounds each, up to 6 candidates per compound	Claude Opus	data/gw/	2026-06-10 18:15

arm	model	data directory	collected (UTC)
forward verification of the widened pool: 4 agents, 65 new candidates	Claude Opus	data/fverify2/	2026-06-10 18:19
cross-model recall check on V3-R01...R12 (n=12), identical blind 6-candidate protocol	Claude Sonnet	data/gw/raw/sonnet_b1.json sonnet_b2.json	2026-06-10 18:33–18:38
headline main round, 140 problems (134 spectrally validated), decoupled agents	Claude Opus	data/benchmark_main/raw/	2026-06-11 06:48–09:16
four-model comparison, fixed 24-compound subset	Claude Haiku	data/benchmark_main/haiku/	2026-06-11 16:16
four-model comparison, same subset	Claude Sonnet	data/benchmark_main/sonnet/	2026-06-11 16:41–17:10
electrolyte subset, 48 curated / 46 scored, 8 batches	Claude Opus	data/benchmark_electrolyte/	2026-06-11 18:29–18:43
four-model comparison, same subset	Claude Fable 5	data/benchmark_main/fable5/	2026-06-11 23:06–23:32
formula-only contamination control, same 60 compounds, spectra masked	Claude Opus	data/modality/	2026-07-28 18:37
forward verification extended to the whole benchmark: 15 agents, 247 main-round candidates	Claude Opus	data/fverify_main/	2026-08-07 02:45–02:53
generate-wide coverage gap closed: 9 agents, the 152 wide candidates that had no prediction	Claude Opus	data/fverify_gw/	2026-08-07 03:20–03:30
trained-generator arm re-run: 5 agents, 75 outstanding candidates	Claude Opus	data/fverify_gen/	2026-08-07 03:55–04:05
cross-vendor sweep, non-Claude models (§S7)	see Table S4	sweep_out/	2026-08-13 to 2026-08-17

Claude invocations fall into three dated windows, not one. **2026-06-09 to 2026-06-11** produced every candidate structure behind the headline results and behind the pools all of §5 re-ranks; no headline elucidation artifact exists outside it. **2026-07-28** is the formula-only control, which generates its own candidates (3/60 correct) by design, a masked-input control being meaningful only as a fresh run; it touches Table 5 alone. **2026-08-07** carries three forward-prediction collections — the §5.2 extension to all 194 compounds, the §5.3 coverage-gap closure and the §5.6 re-run — which predict ^{13}C for candidates the June solver had already produced; none introduces a new candidate or moves a recall number. All later commits re-score frozen outputs and re-query no model.

S1.1 Version strings, snapshots and decoding parameters

Claude Opus 4.8 is the only Claude version string anywhere in the repository, and it is evidenced only for the 2026-06-09 pilot (docs/BENCHMARK.md). It is a display version, not a snapshot identifier: it pins no checkpoint. Fable 5 appears as a display name; **no version number of any kind appears for Claude Sonnet or Claude Haiku**, and no dated snapshot identifier exists for any of the four, because the consumer harness exposes no checkpoint identifier, announces no build

change, and records nothing about which build served a request. Two numbers drawn from one window are therefore known to share a window and *not* known to share a build: **a mid-window build change cannot be excluded**, and the article does not claim the pilot build served the main round two days later.

No sampling parameters were set for any run, and none are recorded — no temperature, `top_p`, `top_k`, `max_tokens`, generation seed or thinking budget appears in any script, document, config or artifact. The `seed=` values in the repository are for analysis determinism only (bootstrap resampling, benchmark sampling, audit-sample selection). Reproduction is distributional, not exact.

One protocol asymmetry belongs here, because §4.3 measures a large effect of the variable it concerns. The Opus column of the four-model comparison is not a fresh 24-compound run: it re-scores the main-round predictions on that subset (the SRC map in `scripts/score_models.py`), whose 24 items came from one 6-compound context (`raw/b1.json`) and two 12-compound contexts (`raw/redo_b23.json`, `raw/redo_b45.json`), where Sonnet, Haiku and Fable each saw four 6-compound contexts. Prompts and compounds are identical; context packing is not.

S2. Prompts and agent protocol

What the solver was shown: the molecular formula (as from HRMS), the IR band list, and the ^1H and ^{13}C shift lists with multiplicities and J-couplings where the source paper reported them.

What it was not shown: the compound name, any SMILES, any starting material or scaffold hint, the ground truth, and any other benchmark record. Solver agents were closed-book — no web access, no ground-truth access, no tools beyond an RDKit formula/parse check — except the arm (a) baseline of §4.3, which had no tools at all, that being the variable it isolates. Closed-book status was grep-audited over the task transcripts at run time; the transcripts are not committed (§7).

Candidate budget: up to **three ranked candidate SMILES per compound**, best first, in every arm except generate-wide, where ten independent solver agents each proposed up to **six regiochemistry-aware candidates** and the pools were merged. Scoring reads at most the first three for top-1 and recovered (top-3); generation recall reads the whole pool.

Context discipline: one small batch per agent in a bounded context that is reset between batches — 6 compounds per released batch file (`data/benchmark_main/batch_*.txt`, 23 files: 22 of six plus one of two = the 134 spectrally-validated problems), with some batch pairs merged into a single 12-compound context (`data/benchmark_main/raw/redo_*.json`). Range released: **2–12 compounds per context**.

Forward-prediction agents had **zero tools**, pure reasoning, and were blind: anonymised SMILES only, pooled across compounds, canonicalised, de-duplicated, shuffled and re-labelled, in fixed batches of 17. They saw neither the observed spectrum, nor the compound’s identity, nor which candidates belong to one target. Shuffling does not keep a target’s own candidates in separate batches — in the §5.2 arm 7 of 8 batches held two candidates for some one compound — but with no observed spectrum in hand there is nothing for that to leak: at most a predictor notices that two structures are isomers, evident from either alone.

Verbatim prompts. Both stage prompts are committed (`sweep_prompts/solve_01.md ... solve_10.md`, `sweep_prompts/verify/`) and emitted by `scripts/cross_vendor_sweep.py`. The elucidation header:

```
# Blind structure-elucidation task
```

82

83 You are given real experimental spectra (from the published literature) for a set of

84 organic molecules. For EACH compound you are given the molecular formula (from HRMS),

85 the IR band list, and the ¹H and ¹³C NMR shift lists. No name, SMILES, or hint is given.

86

87 For each compound, propose the 3 most likely structures, best first, as SMILES.

88

89 Rules:

90 - Use only the spectra provided. Do not use external lookups or tools.

91 - Candidates must match the given molecular formula exactly.

92 - Order candidates by your own confidence (most likely first).

93

94 Return ONLY a JSON object mapping each id to a list of 3 SMILES strings, e.g.:

95 {"M001": ["CCO", "COC", "..."], "M002": [...], "...", "..."], ...}

96

97 Compounds:

98 The forward-prediction header:

99 # Forward ¹³C prediction task

100

101 For each candidate structure below (given as SMILES), predict its ¹³C NMR chemical

102 shift list (in ppm) from the structure alone. This is the forward direction only:

103 you are NOT given any observed spectrum and must not assume one.

104

105 Return ONLY a JSON object mapping each id to a list of predicted ¹³C shifts (numbers

106 in ppm), e.g.: {"P001": [21.0, 60.5, 171.2], "P002": [...], ...}

107

108 Candidates:

109 Each compound follows as one block in exactly this form, the NMR strings being the source pa-

110 per's own text reproduced unedited (hence the chemical-shift symbol and whatever assignment

111 annotations the authors wrote):

112 ### S2.1 M001

113 Molecular formula: <formula>

114 IR bands (cm⁻¹): <list of wavenumbers>

115 ¹H NMR: <string as reported>

116 ¹³C NMR: <string as reported>

117 The real blocks are released per arm: data/benchmark_main/batch_*.txt and data/gw/batch_*.txt

118 for solving, data/fverify_main/fbatch_*.txt and data/fverify_gw/gbatch_*.txt for for-

119 ward prediction, data/modality/prompt_full.txt and prompt_formulaonly.txt for the

120 contamination control (whose blocks carry the formula line alone).

121 **One gap, stated rather than papered over.** The wording above is the harness prompt used

122 verbatim for the four-vendor replication, and it is the only instruction text committed in the

123 repository. The Claude sub-agents behind the headline arms were dispatched through the Agent

124 tool with their instruction supplied at dispatch, and **that instruction text was not captured.**

125 What is released for those arms is the per-compound data exactly as the agents received it, plus

126 the candidate budget, tool policy and context sizes above.

S3. Benchmark construction

IRSpectra-Bench is drawn from `irexp_resolved`, the 43,060-record structure-complete split of IRexp. A record enters the sampling pool only if it carries an IR band list, a ^1H shift list, a ^{13}C shift list and a resolved structure; the structure has **8–60 heavy atoms**; each NMR string carries at least three parenthesised entries; and the raw ^1H string, re-fetched from the source article, contains genuine J information. Draws are balanced between the difficulty strata, and every compound revealed in a prior round is excluded by InChIKey-14 beforehand, so no compound appears twice (`scripts/benchmark_v2.py`).

Table S2. Rounds behind the n=194 cohort.

round	data directory	problems	in the cohort	role
main	<code>data/benchmark_main/</code>	140	134	headline, spectrally validated
controlled v3	<code>data/benchmark_v3/</code>	40	40	difficulty control, pre-registered
within-compound control	<code>data/benchmark_v2_ctrl/</code>	20	20	context/tooling control, pre-registered
electrolyte case study	<code>data/benchmark_electrolyte/</code>	48	—	domain subset, 46 scored, held apart

What was excluded, and why. Every main-round ground truth passes an automated RDKit self-consistency check — ^{13}C peak count against symmetry-unique carbons, formula match, SELFIES round-trip — which excluded 6 of 140: five report more ^{13}C peaks than the structure has carbons (R03 28>24, R14 11>8, R65 26>18, R67 23>16, R131 34>17: merged or contaminated spectra) and one is too sparse to constrain (R82, 5 peaks for 22 symmetry-unique carbons). The two controlled rounds are used **whole**, being fixed and pre-registered as controls before the audit existed; the same audit is reported on them rather than applied (57/60 pass), and §4.1 gives the strictly-validated 191-compound cohort as a robustness check. The filter runs over all three rounds in `scripts/validate_benchmark.py`, which regenerates every `clean_qids.json` from the released questions and answers. It tests ^{13}C against the carbon count and does not gate on ^1H : in **13 of the 194** retained records the reported ^1H integral exceeds the reference structure’s hydrogen count (residual solvent, water, exchangeable protons, a rotamer mixture). These are printed as a diagnostic, not excluded, because the cohort was fixed in advance; dropping all 13 moves the headline from 28.4% to 29.3% (53/181).

Complexity stratification, and when it was defined. Difficulty is a declared property of the compound, assigned by RDKit ring analysis at sampling time — before anything was solved — and exhaustive by construction, so nothing falls between the classes. A compound is **complex** if it has a spiro atom, a bridgehead atom, three or more rings, any fused ring system, **or** more than 24 heavy atoms; **simple** if it has at most two rings and at most 22 heavy atoms; anything else (the 23–24-heavy-atom band, 13 compounds) is complex on size alone. The cohort splits 98 simple / 96 complex. The 22-atom threshold is not load-bearing: sweeping it from 18 to 26 moves the simple-minus-complex top-1 gap only between 36 and 40 points, against 39.6 at the released value (`scripts/difficulty_sensitivity.py`). This axis is distinct from the heavy-atom bands (≤ 15 / $16\text{--}25$ / >25) of §4.1 and Fig. S3.

The battery-electrolyte subset (§4.5) was drawn from the same corpus by SMARTS filters for six electrolyte functional classes, balanced to eight per class (48 curated; 46 scored after two yielded no parseable candidate), J-enriched and spectrally validated identically to the main rounds, excluding every compound used elsewhere.

S4. Scoring: what the criterion accepts and rejects

A prediction is **correct** if the first 14 characters of the RDKit InChIKey computed from its SMILES equal those of the reference — the InChIKey **connectivity layer**, which hashes the molecular skeleton. The later blocks, carrying stereochemistry, isotopic substitution and protonation state, are not compared. Top-1 asks this of the first-ranked candidate, recovered (top-3) of the up-to-three returned, generation recall of the whole pool. Nothing is model-mediated: RDKit only, in `scripts/score_main.py` and `scripts/specmetrics.py`. Morgan(2, 2048) Tanimoto¹ is reported alongside as a graded signal, 0.45 being the scaffold-level threshold; intervals are bootstrap 95% over compounds (2,000 resamples, analysis seed 0) and model-vs-model tests McNemar exact with Holm correction.

What it accepts. A candidate with the right constitution and the wrong stereochemistry scores correct. This is a deliberate floor whose cost is measured rather than assumed: re-scoring the *full* InChIKey gives **21.1% top-1 and 25.8% recovered (41/194 and 50/194)** against 28.4% / 33.5% at the connectivity layer, so fourteen top-1 answers are accepted here that a stereochemistry-sensitive scorer rejects. The floor exists because 1D ¹H/¹³C/IR rarely fixes absolute configuration and only 10.3% (20/194) of answers carry a defined (assigned R/S) stereocentre, so a model would be penalised for information the prompt never carried. The cross-vendor arm shows the accepted case per compound: on structures those models got constitutionally right whose reference carries assigned stereocentres, Grok 4.6 reproduced 0/3 correct descriptors, Gemini 3.7 Flash 0/2 and GPT-5.6 Sol 0/2.

What it rejects: every constitutional isomer, however close, and every structure of the wrong composition. Three worked cases, all from released artifacts:

- *A regioisomer the verifier separates.* Picolinamide and nicotinamide differ only in which ring position bears the carboxamide, so they have different connectivity layers and proposing one for the other scores zero; forward-predicted ¹³C separates them at a chamfer of 0.42 ppm against 1.30 ppm (Fig. 5).
- *A regioisomer it does not.* On the 2-(nitrophenyl)-2,3-dihydroquinazolin-4(1*H*)-one targets (C₁₄H₁₁N₃O₃) the predictor cannot resolve the *ortho*- and *meta*-nitrophenyl isomers: it ranks the 2-nitrophenyl isomer first at a chamfer of 1.35 ppm against 1.36 ppm for the true 3-nitrophenyl one, a 0.01 ppm margin. The answer is rejected, as it must be.
- *A pilot miss of the same kind.* The n=21 pilot returned the correct *N*-allyl-2-(pyridinecarbonyl)hydrazinecarbonyl skeleton with the **3-pyridyl** isomer where the truth is **2-pyridyl** (docs/BENCHMARK.md, Q21): right formula, right scaffold, rejected.

That is the dominant shape of failure, not an edge case: over all 139 top-1 misses, **76.6% are constitutional isomers** of the true structure and 23.4% have the wrong formula outright; 22.6% share the true Murcko scaffold, 2.9% reach Tanimoto ≥ 0.85 , and the median Tanimoto between an isomeric miss and the truth is 0.39 (`scripts/analyze_misses.py`). Formula adherence is not part of the criterion — a wrong-composition answer is simply a miss — and it is not uniform across rounds: on top-1 answers it is 95.0% (38/40) in v3, 90.0% (18/20) in the within-compound control and 77.6% (104/134) in the main round.

S5. Forward-verification detail

The distance. Candidates are ranked by a symmetric chamfer distance between the forward-predicted and observed ^{13}C peak sets: a mean of two nearest-neighbour means, so no equal-count requirement is imposed and lower is better. The implementation is the single source of truth for every scorer and diagnostic (`scripts/specmetrics.py`):

```
def chamfer(pred, obs):
    """Symmetric chamfer distance between predicted and observed 13C shift lists."""
    if not pred or not obs:
        return 999.0
    a = sum(min(abs(x - y) for y in obs) for x in pred) / len(pred)
    b = sum(min(abs(y - x) for x in pred) for y in obs) / len(obs)
    return (a + b) / 2
```

The re-ranking. Within a compound the candidate of smallest chamfer becomes the top-1 answer and the solver’s own ordering is discarded. The sentinel matters: a candidate with no forward prediction takes 999.0, effectively infinite, and can never be selected. Prediction coverage is therefore a precondition, not a detail — an arm with unpredicted candidates reports a lower bound on verified top-1 rather than a measurement, which is why the generate-wide arm was re-run to completion.

Table S3. Forward-prediction coverage, by arm. Batch size is 17 anonymised SMILES throughout.

arm	data directory	candidates	forward-predicted	agents
original 60-compound arm	<code>data/fverify/</code>	126	126	8
widened candidate pool	<code>data/fverify2/</code>	65 new	65	4
extension to the whole benchmark	<code>data/fverify_main/</code>	247	247	15
generate-wide coverage-gap closure	<code>data/fverify_gw/</code>	152	152	9
trained-generator arm, re-run	<code>data/fverify_gen/</code>	75	75	5
		outstanding		

Across the whole benchmark that is **373 of 373** candidates forward-predicted by 23 independent agents over all 194 targets, so nothing in Table 7 is a lower bound. In the generate-wide arm all **217 of 217** distinct new candidates are now predicted where the original pass had predicted 65, and **not one number moves**: recall 42%, forward-verified top-1 30%, precision 72%, identical to three significant figures. What the added coverage buys is a direct view of the mechanism — on **18 of the 60** compounds the verifier abandons its previous pick for a newly-selectable candidate, and in **not one** of those 18 does the outcome change; every switch is from one wrong structure to another.

The margin the method works on. Measuring the chamfer between the *predicted* spectra of every pair of candidates proposed for the same target (`scripts/isomer_separability.py`), isomeric pairs sit at a median separation of **1.21 ppm** (quartiles 0.84 and 1.78) and **82% of them are predicted closer together than the predictor’s own ~2 ppm error**; non-isomeric pairs

separate better but not dramatically (median 1.90 ppm, 53% inside the error). That the ranking nevertheless recovers the right structure 89% of the time it is present is what the permutation control tests: re-pairing, as a derangement, which observed spectrum each candidate set is scored against, over 1,000 permutations, drops conditional-on-recall precision from **89.2% (58/65)** to a permuted mean of **73.8%** (95% range 66.2–81.5%; one-sided $p=0.001$, two-sided $p=0.002$). On the 60-compound arm alone it gives 84.2% against 66.4% (one-sided $p=0.019$). The chance floor sits high because recall-positive compounds carry few, near-identical candidates, so the genuine margin over chance is ~ 15 points (§5.5).

A negative result on calibration. Ranking the 138 multi-candidate compounds by chamfer margin (best minus second-best) and answering only the most confident fraction leaves top-1 flat and non-monotonic with coverage: 22% at full coverage, 24% at 75%, 28% at 50%, 24% at 25%. Single-candidate compounds have no margin and must be excluded; an earlier analysis that kept them produced a spurious improvement that was entirely an artefact of those trivial cases.

S6. Non-LLM verifiers: a deterministic lookup and a learned model

Both non-LLM ^{13}C predictors are built from the **same nmrshiftdb2 dump**² and dropped into the verifier slot on the **same candidate sets**, so only the predictor changes.

The deterministic lookup realises the HOSE-code idea³ natively in RDKit: the Morgan per-atom identifier at radius r is a canonical hash of an atom’s environment out to r bonds, so a (radius, identifier) bin groups carbons with identical local environments. The mean assigned shift per bin is learned from nmrshiftdb2, and prediction walks the deepest sufficiently-populated sphere, $r=4 \rightarrow 3 \rightarrow 2 \rightarrow 1$, falling back to a hybridisation prior; a bin needs at least 3 samples to be trusted. Training used **31,000 molecules, 332,595 assigned carbons and 263,366 environment bins**, with held-out **MAE 3.23 ppm (median 1.73)** over 17,456 carbons of 1,647 molecules (scripts/hose_predict.py).

The learned model is a message-passing GNN — 4 layers, hidden width 256, GRU node update, bond features, per-carbon ^{13}C regression — trained on the identical dump: 32,647 molecules and 350,313 assigned carbons, split by molecule into 29,383 train / 1,632 validation / 1,632 test at seed 5 (the lookup’s figures above are this set minus its random 5% held-out split). Optimisation is Adam at learning rate $1e-3$, weight decay $1e-5$, batch size 64, smooth-L1 loss, plateau learning-rate halving, early stopping. Held-out **MAE 1.70 ppm (median 1.02)**, roughly twice as sharp as the lookup (scripts/gnn_predict.py, data/nmrshiftdb/gnn_c13.pt). It is deliberately modest rather than state of the art — purpose-built ^{13}C models reach ~ 1 ppm⁴ and 0.94 ppm with DFT shielding tensors⁵, on benchmarks now standardised for the comparison⁶ — because the question is whether the predictor slot is where the leverage sits, which needs a predictor differing from the lookup in *method* while data and evaluation are held fixed.

Held-out error against benchmark behaviour. Table 9 gives the four verifiers conditional on recall; two readings belong here. The lookup’s tie with self-ranking is not agreement: at $n=65$ it **gains seven compounds and loses seven** (McNemar $b=c=7$, $p=1.00$), reshuffling energetically while carrying no net discriminative signal. Its coverage diagnosis is that of the **6,360 candidate carbons only 2.1% match a training environment at the most specific sphere ($r=4$)**, and 71% resolve only at $r \leq 2$ or fall through to the prior, the benchmark’s exotic chemistry being under-represented in nmrshiftdb2. The GNN’s margins are directional in every pairwise comparison and none reaches significance: against the lookup seven gained and three lost ($p=0.34$), against self-ranking five gained and one lost ($p=0.22$), and against the LLM verifier the two land within one

compound of each other while disagreeing on nine (b=5, c=4, p=1.00).

Leakage analysis is the decisive control for a *learned* verifier, since a GNN can memorise a molecule’s spectrum where a bin average cannot. Exact overlap with the whole nmrshiftdb2 database is **2 of 364** distinct candidate structures by InChIKey-14 — both wrong candidates on recall-negative compounds, which never enter the conditional analysis — and **no benchmark answer appears in the database at all** (0/373; the 60-compound arm is 0/126). Analog overlap is likewise absent: median Morgan(2, 2048) Tanimoto to the nearest training molecule is **0.44**, with three candidates above 0.80. The load-bearing figure is the last: over the **65 true structures the verifier must identify**, the nearest training analog has median Tanimoto **0.50** and maximum **0.81**, so not one compound the conditional analysis scores has a near-duplicate in training. A Y-randomisation control (1,000 derangements) places the learned verifier above the 97.5th percentile of the chance distribution (n=19: real 84% against a permuted mean of 58.6%, 95% range 42.1–73.7%, one-sided p<0.05). Both checks regenerate via `scripts/verifier_leakage.py`. Neither predictor is redistributable end-to-end: the nmrshiftdb2 dump cannot be shipped with the manuscript, so these two rows regenerate only once a reader supplies the same dump.

S7. Cross-vendor replication

Non-Claude models were run through `scripts/cross_vendor_sweep.py` on the `fverify60` arm — the same 60 compounds as the first column of Table 7 — between 2026-08-13 and 2026-08-17. Two arms went through OpenRouter (`scripts/openrouter_run.py`); the rest were driven by cloud coding agents against the committed `sweep_prompts/`, one fresh subagent per batch, at no API cost. Raw replies are released in `sweep_out/`.

Table S4. Generation stage, every model run on the 60-compound arm, with the article’s Claude Opus arm for reference.

model	route	answered	recall	top-1	parse	matches the given formula
grok-4.6	cloud	60/60	53%	38%	99%	95%
	agent					
gemini-3.7-flash	cloud	60/60	50%	38%	98%	94%
	agent					
gpt-5.6-sol	cloud	60/60	42%	35%	100%	100%
	agent					
<i>Claude Opus (reference arm)</i>	subscription	60/60	<i>32%</i>	<i>23%</i>	—	<i>78–95%</i>
composer-2.5	cloud	57/60	20%	17%	95%	67%
	agent					
gpt-5.6-luna	cloud	60/60	15%	10%	90%	76%
	agent					
deepseek/deepseek-v4-pro	OpenRouter	18/60	13%	12%	94%	94%
nvidia/nemotron-3.5-lightning	OpenRouter	60/60	0%	0%	61%	2%

Read the formula-adherence column first: nemotron’s zero is not a chemistry result but a model that could not return a structure of the requested composition, and the deepseek arm answered 18

of 60 compounds — a reasoning model that exhausted its token ceiling without emitting an answer on most batches, where an unanswered compound scores exactly like a wrong one.

Matched across arms: the compounds, the two-stage protocol, the three-candidate budget, the context packing (six compounds per fresh context), the closed-book instruction, the anonymised blind forward-prediction stage, and the scorer. **Not matched:** verification was run only for the three models whose output contract justified it, Composer 2.5 and GPT-5.6 Luna being left out at 67% and 76% formula adherence.

Table S5. Decomposition, the three replicated arms. Recall and precision have different denominators, so the criterion is the inequality, not a difference.

model	generation recall	verification precision given recall	multi-candidate only
<i>Claude Opus (reference arm)</i>	19/60 = 32% [21, 44]	16/19 = 84% [62, 94]	10/13 = 77%
grok-4.6	32/60 = 53% [41, 65]	20/32 = 62% [45, 77]	20/32 = 62%
gemini-3.7-flash	30/60 = 50% [38, 62]	22/30 = 73% [56, 86]	22/30 = 73%
gpt-5.6-sol	25/60 = 42% [30, 54]	17/25 = 68% [48, 83]	16/24 = 67%

What these identifiers name. The models are as served by the coding-agent harness, not stock vendor endpoints. That harness’s allowlist reads `gpt-5.6-sol-xhigh`, `gpt-5.6-sol-high`, `gpt-5.6-luna-high`, `gemini-3.7-flash-high`, `cursor-grok-4.6-high-fast`, `composer-2.5` — every identifier but the last carries a **reasoning-effort suffix**, a decoding parameter the Claude reference arm never had (§S1.1) — and the harness supplies its own agent system prompt above the task text. **The effort tier used for each arm is not recorded**, because the run did not report which allowlist entry it selected; a run at `-xhigh` is not the same measurement as one at `-high`, and the article reports these as vendor-family results accordingly. GPT-5.6 Terra and GLM 5.2 were attempted and rejected: neither is in the agent’s allowlist.

Contamination control. The cloud agents cloned the whole repository and the answer files for these 60 compounds are tracked, so blindness rested on an instruction nothing verifies after the fact. The control is a re-solve with those files physically out of the workspace: Grok 4.6 re-ran all ten batches from a clean clone (`sweep_out/grok-4.6-clean/`).

Table S6. Clean-clone contamination control, Grok 4.6.

arm	recall	95% CI	parse	formula match
original (answer files present)	32/60 = 53%	[41, 65]	179/180	171/180
clean clone (answer files absent)	28/60 = 47%	[35, 59]	176/180	174/180

Paired, that is **b=8, c=4, McNemar exact p=0.39**. The decisive detail is the asymmetry rather than the totals: a model reading the key would solve a superset, and instead **four compounds were**

solved only in the arm that had no key, with 24 of the 60 solved by both. Formula adherence held at 97% against 98%, where a model that had lost a crib would be expected to degrade. The control covers Grok; Gemini 3.7 Flash and GPT-5.6 Sol rest on the shared instruction and on the stereochemistry evidence of §S4.

S8. Modality ablation and contamination controls

The leave-one-modality-out ablation was specified and never run, and no leave-one-out result appears in the article. The staged prompts `prompt_noIR.txt`, `prompt_noH.txt` and `prompt_noC.txt` sit under `data/modality/` (2026-06-16) with no corresponding `out_*.json`. Two attempts were discarded, and why is instructive. The first ran one solver agent per condition, so agent quality was a single draw per condition and swamped the modality effect: full modality came out worst at 3/16 while `-IR` came out best at 8/16, and `-IR` solved compounds the full-information run missed. The second ran a fresh `-IR` arm against the *archived* benchmark predictions as its control, giving full 10/30, `-IR` 22/30 and formula-only 2/30 on the simple stratum — `-IR` beating full modality by 40 points, McNemar $p=0.004$, which is impossible and identifies the confound rather than a modality effect. The rule this yielded is that every arm of an ablation must be generated in one campaign with the same agent configuration, batch size and model, the control included (`docs/MODALITY_ABLATION.md`).

The formula-only control shares that structure and is not impugned by it, because the confound runs against the finding. Only the formula-only arm was freshly generated (2026-07-28); the full-modality comparison arm re-uses the archived June predictions, whose 60 top-1 answers are byte-identical to that run. Fresh agents reason harder per compound than the archived batched run, so the bias favours the formula-only arm — and it still collapsed. The solver received the molecular formula and nothing else, was barred from reading any repository file or searching the web, and was allowed RDKit only to check that a proposed SMILES parses and matches the formula. The outcomes are perfectly nested: **eleven** compounds are solved with the spectra and not without and **none** the other way round ($b=11$, $c=0$, McNemar exact $p=0.001$), for 3/60 against 14/60 top-1 and 3/60 against 19/60 recovered (Table 5).

The three formula-only successes are named rather than rounded away, because they do not all support one reading: **C₁₅H₁₅ClF₂O₃SSi**, a 2-(trimethylsilyl)aryl sulfonate whose Si/S/Cl/F₂ composition is a near-unique benzyne-precursor signature and which is absent from PubChem, so the formula is close to determining; **C₁₃H₁₉NO₄S**, *N*-tosyl-leucine (CAS 67368-40-5), where inference and recall are both available; and **C₁₇H₂₆O₃**, [6]-paradol (CAS 27113-22-0), a catalogued ginger natural product that composition alone constrains very little.

The recency control is independent and agrees. Publication years were resolved for all 194 compounds from their accessions (`scripts/contamination_recency.py`) and span 2008–2026. Accuracy is flat: **28.6%** for the older half (≤ 2020 , $n=112$) against **28.0%** for the newer ($n=82$), point-biserial r between year and correctness of **-0.007**, the most recent bucket (≥ 2024 , $n=25$) in fact highest at 40% [23, 59]. The raw split is biased against the newer half, since newer papers skew larger (median 22 heavy atoms against 20) and size dominates accuracy; stratifying by heavy-atom band removes that, and newer compounds lead in the two bands carrying most of the accuracy (≤ 15 : 64% against 58%; 16–25: 34% against 25%) and trail slightly in the largest, where both are near the floor (> 25 : 6% against 8%). The size-adjusted older-minus-newer difference is **-5.1 points**, **95% CI [-17.2, +7.0]** (Cochran–Mantel–Haenszel $\chi^2=0.42$, $p=0.51$, continuity-corrected) — a bound on any recency effect, not a demonstration of a reversed one.

S9. Human-expert audit: protocol and status

Solver and verifier are both LLMs, so the one validation the authors cannot perform themselves is an expert-chemist review of the outputs. The audit package is therefore built, blinded, pre-registered and **frozen** before any review, and regenerates deterministically at seed 0 from `scripts/make_audit_sample.py` into `data/audit/`.

The panel is three PhD-level synthetic or analytical chemists, independent, with no prior exposure to the benchmark answers, at roughly 3–4 person-hours each. The sample is a difficulty-stratified draw (15 simple, 15 complex) from the 60-compound forward-verification set, the only split carrying both the ranked candidates and the observed spectra. Per compound a reviewer sees the exact solver prompt — formula, IR, ^1H , ^{13}C — and the model’s top-ranked candidate rendered as a 2D structure, with model identity and ground truth hidden.

Table S7. Composition of the frozen audit kit.

item	count
compounds in the kit	37
Task 1 panel (unbiased stratified draw)	30
Task 2 ranking items	37
...carrying the true structure and a real choice, scoring precision	13
...carrying the true structure but a single candidate, excluded	3
...carrying no correct candidate, measuring expert calibration	21
model top-1 exact on the Task 1 panel (held in the withheld key only)	7/30

Task 1 scores the model’s top-1 for consistency with all provided spectra (1–5), gives a categorical verdict (correct / wrong-regiochemistry / wrong-scaffold / uninterpretable) and names the single most diagnostic peak; the read-outs are inter-rater agreement and the fraction of mechanically-wrong top-1 answers judged spectrally consistent but a different regioisomer. Task 2 presents the shuffled, unlabelled candidate set for ranking by spectral fit, and is compared against the LLM forward-verifier’s pick and the deterministic lookup’s pick on identical candidate sets. Two design decisions are worth recording because both were mistakes first: Task 2 is shown on **every** compound, so its presence signals nothing — an earlier version showed it only on recall-positive compounds, which leaked Task 1 twice over and moved the apparent rate of correct answers from 12% to 67%, a five-fold prior available without reading a spectrum — and the seven extra compounds carry Task 2 only, because a correct top-1 is recall-positive by definition, so adding recall-positive compounds to the Task 1 draw would raise the very rate Task 1 exists to judge.

Status. The panel has not been run and the article reports no audit result (§7). One reviewer file is deposited at `data/audit/responses/` (submitted 2026-08-18 UTC) and is incomplete against the 37-item kit; a single partial reviewer supports neither read-out, since inter-rater agreement needs at least two and the pre-registered panel is three. Scoring of completed sheets is mechanical (`scripts/score_audit.py`), with no further model involvement. Three Task 2 sets (A19, A21, A30) hold a single candidate, which cannot be ranked; the scorer excludes and flags them. Until expert results are in, the two load-bearing claims the panel targets — what a miss actually is (§4), and whether forward verification is a trustworthy re-ranker (§5) — should be read as machine-validated (RDKit InChIKey) but not yet human-validated.



Fig. S1: Study design: open multimodal data (IReXP) → blind, complexity-stratified benchmark → decoupled blind solving → forward-verification re-ranking; training-free core pipeline.




Fig. S2: IReXP composition: IR records, NMR-paired, structure-linked, and full IR+¹H+¹³C+structure quadruples.




Fig. S3: Accuracy versus molecular size; the monotonic 60%→7% top-1 gradient with heavy-atom count.




Fig. S4: IRSpectra-Bench-Electrolyte by battery-electrolyte class (n=46): sp³-C-F easiest (50%), sulfonyl and nitrile hardest (12%).




Fig. S5: Trained-generator probe (§5.6; a complement, not part of the training-free protocol). Candidate recall and deterministic-HOSE top-1 on the 194-compound benchmark for Claude / + scaffold enumeration / + trained generator: enumeration’s near-degenerate isomers collapse the verifier (28.4→16.0%) while the generator’s formula-correct candidates convert (28.4→35.1%).




Fig. S6: Learned-verifier probe (§5.4; a complement, not part of the training-free protocol). (A) Conditional-on-recall top-1 over the whole benchmark ($n=65$) across four verifiers: a GNN trained on the same nmrshtdb2 data as the HOSE lookup reaches the LLM verifier’s level (91% against 89%) where the lookup (85%) does not move off the solver’s own ranking. (B) Held-out ^{13}C MAE — the learned model is roughly $2\times$ sharper (1.70 vs 3.23 ppm).

- 1 D. Rogers and M. Hahn, *J. Chem. Inf. Model.*, 2010, **50**, 742–754.
- 2 S. Kuhn and N. E. Schlörer, *Magn. Reson. Chem.*, 2015, **53**, 582–589.
- 3 W. Bremser, *Anal. Chim. Acta*, 1978, **103**, 355–365.
- 4 D. Williamson, S. Ponte, I. Iglesias, N. Tonge, C. Cobas and E. K. Kemsley, *J. Magn. Reson.*, 2024, **368**, 107795.
- 5 C. Han, D. Zhang, S. Xia and Y. Zhang, *J. Chem. Theory Comput.*, 2024, **20**, 5250–5258.
- 6 F. Xu, W. Guo, F. Wang, L. Yao, H. Wang, F. Tang, Z. Gao, L. Zhang, W. E, Z.-Q. Tian and J. Cheng, *Nat. Comput. Sci.*, 2025, **5**, 292–300.