

IRexp: A database of experimental infrared band lists from open literature

Ilkham Yabbarov^{1*} and Rodrigo A. Vargas-Hernández^{1,2,3*}

^{1*}Department of Chemistry and Chemical Biology, McMaster University,
Hamilton, Ontario, L8S 4L8, Canada.

²Brockhouse Institute for Materials Research, McMaster University,
Hamilton, Ontario, L8S 4L8, Canada.

³School of Computational Science and Engineering, McMaster University,
Hamilton, Ontario, L8S 4L8, Canada.

*Corresponding author(s). E-mail(s): yabbaroi@mcmaster.ca;
vargashr@mcmaster.ca;

Abstract

IRexp is a redistributable collection of experimental infrared band lists (cm^{-1} peak positions) mined from open chemistry literature, optionally with author-reported $^1\text{H}/^{13}\text{C}$ NMR strings and resolved structures. The release holds 121,233 records (119,345 PMC Open Access Subset; 1,888 Chemotion/RADAR4Chem), including 43,060 structure-linked entries and 33,201 full IR + ^1H + ^{13}C + structure quadruples. IRexp stores numeric band lists, not absorbance traces; each record carries a source DOI and stamped licence / licence-pool fields. Licensing is mixed: Europe PMC joins yield 87,617 commercial (CC-BY/CC0), 20,938 non-commercial (NC*), 10,781 empty/unknown (excluded from the commercial deposit), and 1,897 ShareAlike (Chemotion plus rare PMC SA). The resource is intended for multimodal training, retrieval, and tool-input in spectroscopic agents and autonomous chemistry workflows that need structured, attributable experimental peak lists rather than paywalled PDFs or view-only archives. Technical validation covers automated transcription and consistency checks; it does not match expert-scale manual audits of larger NMR list peers. Complementary elucidation benchmarks are described in a companion research manuscript and are not analysed here.

Keywords: infrared spectroscopy, band lists, peak lists, PMC Open Access, Chemotion, FAIR data, licence pools, data descriptor

1 Background & Summary

Infrared (IR) spectroscopy is routine in organic characterisation, yet open, redistributable collections of *experimental* IR data remain sparse relative to modern machine-learning and agentic tool-use needs. Digitised absorbance libraries such as the NIST Chemistry WebBook [1] and AIST SDBS [2] are valuable but either modest in size or view-only (no bulk redistribution). Computational IR–NMR resources published in *Scientific Data* expand multimodal coverage with simulated spectra [3], while large literature mines for NMR peak lists (notably NMRexp [4]) show that peer-reviewed experimental spectral *lists* with DOI traceability are in scope for this journal. Concurrent literature corpora that include IR among other modalities (for example NMRSpec/NMRTrans [5]) further motivate an IR-focused, redistributable band-list resource rather than another closed spectrum archive.

Relation to NMRexp and other peers.

NMRexp is the natural comparator: ~3.3 million experimental NMR records mined from supporting-information PDFs, with expert-scale manual checks and replicate consistency metrics [4]. IRexp is *not* an NMR database, does not claim size superiority, and is orders of magnitude smaller. Its contribution is complementary: redistributable *IR band lists* (cm^{-1} positions only) from PMC Open Access full text plus a Chemotion ELN deposit, with per-record licence pools suitable for commercial vs non-commercial reuse. Absorbance-curve libraries (NIST, SDBS) remain the right choice when full digitised spectra are required; computational IR–NMR sets remain the right choice when simulated multimodal coverage is required [3].

In 2025–26, spectroscopic AI agents and autonomous chemistry workflows increasingly consume *structured* experimental peak lists as tool inputs and retrieval substrates — not paywalled PDF prose and not view-only web UIs. Agents that plan characterisation, call spectrum tools, or train IR→structure models need redistributable numeric lists with DOI attribution and explicit licence pools (commercial vs non-commercial vs ShareAlike). IRexp is designed as that substrate: training / retrieval / tool-input material for literature-grounded spectroscopic agents, without embedding diagnosis benchmarks or model leaderboards in this Descriptor.

IRexp fills a specific wedge. Experimental sections of chemistry papers conventionally report per-compound IR band lists (wavenumbers in cm^{-1}) together with $^1\text{H}/^{13}\text{C}$ NMR shift lists. That textual convention is the object language models and many elucidation pipelines consume, and it is a different object from a digitised spectrum. IRexp therefore:

1. Harvests open-access full text from the PMC Open Access Subset bulk S3 mirror and peak lists from the Chemotion FT-IR deposit.
2. Extracts deterministic numeric IR band lists (and co-reported NMR strings where present).
3. Resolves compound names to canonical structures with OPSIN [6], RDKit [7], and SELFIES [8] where possible.
4. Releases extracted numbers only — no PDFs, figures, or article full text — with source accessions for attribution.

Among openly redistributable *text-derived IR band lists*, IRexp is large by record count (121,233). It does not claim to replace SDBS or NIST absorbance libraries; SDBS alone holds more structure-linked *spectra* than IRexp holds structure-linked band lists. The scientific contribution of this Descriptor is the curated dataset, harvest provenance, licence segregation, and validation artefacts — not elucidation accuracy claims. The intended reuse is multimodal pretraining, supervised IR→structure modelling on `irexp_resolved`, and as the literature substrate for complementary elucidation benchmarks described elsewhere (forthcoming ICLR manuscript on IRSpectra-Bench; cite that work for protocol and model results — not reproduced here) [9].

What this Data Descriptor does not contain.

No hypothesis tests, no large-language-model accuracy tables, and no stage-decomposition of elucidation performance. Those belong in the complementary research paper.

2 Methods

2.1 Source corpora

PMC Open Access Subset.

Primary harvest uses the NCBI PMC OA bulk distribution on Amazon S3 (`s3://pmc-oa-opendata`, HTTPS endpoint `https://pmc-oa-opendata.s3.amazonaws.com`) [10]. Harvest window (recoverable from git snapshots): the bulk S3 IR crawl and `seen_papers / ir_harvest_snapshot` artefacts were produced on 2026-06-07 (UTC) (`scripts/s3_ir_harvest.py`; incremental auto-snapshots that day through ~134,893 raw IR rows). Chemotion was ingested the same day. A harvest snapshot records 188,016 distinct PMC identifiers scanned (`data/irexp/seen_papers.txt.gz`). The released corpus retains records from 15,416 unique PMC:* accessions. Extraction operates on open-access plain text objects at flat S3 keys `PMC{id}.{v}/PMC{id}.{v}.txt`. Europe PMC full-text XML is used for some validation re-fetches. The released IRexp DOIs are exclusively PMC:* accessions or the Chemotion deposit DOI; no paywalled publisher DOI appears in the 121,233-record file.

Discovery vs oa_comm / oa_noncomm packages.

Identifier discovery used NCBI E-utilities `esearch` over PMC with an IR-characterisation query and `open access[filter]`, then fetched matching IDs from the flat S3 text layout (`scripts/s3_ir_harvest.py`). The harvest therefore did not pre-filter by walking the PMC OA Subset’s `oa_comm` vs `oa_noncomm` vs other package directories. PMC OA is still not a single licence [10]. Licence truth for redistribution is applied post-hoc: every unique PMCID in IRexp (15,416) was joined to Europe PMC license metadata (`scripts/join_pmc_licences.py`); each record carries `license / license_pool`. Commercial-use (CC-BY/CC0) rows form the Zenodo / Sci Data primary pool; NC* are held aside; empty/unknown are excluded from the commercial

deposit — aligning redistributable intent with the OA commercial/non-commercial distinction via article-level licences rather than S3 package paths (`data/NOTICE`; `docs/scientific_data/LICENCE_REMEDIATION.md`).

Reproducibility of discovery.

Re-running live `esearch` will drift as PMC grows. The frozen discovery set for this release is `data/irexp/seen_papers.txt.gz` (188,016 PMC identifiers); the curated release retains 15,416 of those accessions. Third parties should treat `seen_papers` + the released JSONL as the reproducible snapshot, not a fresh API crawl.

Chemotion / RADAR4Chem.

1,888 records come from the Chemotion Repository FT-IR collection deposited at RADAR4Chem (DOI 10.22000/0GoEQG1sZGE1rgst) [11], licensed CC-BY-SA-4.0. Ingest (`scripts/chemotion_to_irexp.py`, 2026-06-07): download the MD5-verified deposit; for each of 2,116 ATR-IR analyses, flatten the Quill-delta `content` field to plain text; parse an author-curated IR band list with the same regex extractor and quality gates as the PMC path; resolve the deposit’s canonical SMILES with RDKit → InChIKey + SELFIES; keep the richest band list per InChIKey; drop the 8 molecules already present in the PMC pool → +1,888 new structure-resolved rows. All 1,888 Chemotion rows are structure-linked (`has_structure=true`; `inchikey` equals the record `id`). These are not algorithmic peak-picks from absorbance curves; they are author-entered experimental band lists from the ELN deposit, denser than typical paper prose (median 39 bands vs 9 for PMC). Rows carry `license=CC-BY-SA-4.0`, `source=Chemotion`, and `source_doi` equal to the deposit DOI.

Excluded.

AIST SDBS (view-only; no bulk export) [2]; NIST WebBook join code exists in the repository but contributes 0 records to the released IReXP (`ir_source` is always `experimental` for released rows).

2.2 Discovery and fetch (released corpus)

1. Discover IR-reporting OA PMCID IDs via NCBI `esearch` (`open_access[filter]` + characterisation-format IR query; year/month sliced).
2. Fetch OA full text from PMC-OA S3 plain-text objects (`PMC{id}.{v}.txt`).
3. Ingest Chemotion deposit author-curated band lists + structures (`scripts/chemotion_to_irexp.py`).
4. Persist harvest provenance in `seen_papers.txt.gz` and optional rawer rows in `ir_harvest_snapshot.jsonl.gz` (134,893 rows including intermediate prose fields; the curated release is `irexp.jsonl.gz`).

Non-OA / Scrapling fence.

Development adapters for ChemRxiv, Beilstein, and generic publisher pages exist for exploration. They are outside the construction path of the released dataset. No released `source_doi` is a paywalled publisher DOI. Methods for IReXP as published here cite only PMC-OA S3 + Chemotion.

2.3 Extraction (band lists, not spectra)

A deterministic regular-expression pipeline (`spectro_scraper/extract.py`) segments experimental text per compound and extracts:

- IR wavenumbers $\rightarrow$ `ir_bands_cm-1` (list of floats, cm^{-1});
- ^1H and ^{13}C payloads $\rightarrow$ `h_nmr` / `c_nmr` (author strings, when present).

Gates reject scan-range artefacts and common prose false positives. No curve digitisation is performed: figures are not traced; JCAMP-DX is not stored.

2.4 Structure resolution

Where an IUPAC or systematic name is available, names are converted with OPSIN [6], canonicalised with RDKit [7], and encoded as InChIKey and SELFIES [8]. An optional PubChem [12] name fallback (`USE_PUBCHEM`) handles trivial names. Name $\rightarrow$ structure failures are expected for trade names, mixtures, and non-IUPAC prose; unresolved rows keep `has_structure=false` and null structure fields. Structure coverage of the full release is 35.5% (43,060 / 121,233). The structure-complete split is shipped as `irexp_resolved`.

2.5 Licence handling

Table 1 Licence pools stamped on the released IRexp corpus.

Pool	Records	Stamp
commercial (CC-BY + CC0)	87,617	<code>license_pool=commercial</code> ; Zenodo / Sci Data primary
non_commercial (CC-BY-NC*)	20,938	held aside
sharealike (Chemotion + rare PMC SA)	1,897	CC-BY-SA / CC-BY-SA-4.0
empty_unknown	10,781	excluded from commercial deposit

`scripts/join_pmc_licences.py` stamps every row;
`scripts/split_license_pools.py` reports provenance and materialises
pool files under `data/irexp/licence_pools/`. Narrative and policy:
`docs/scientific_data/LICENCE_REMEDIATION.md`.

2.6 Quality tooling

Optional physics gates live in `spectro_scraper/quality.py` (^{13}C peak count $\leq$ carbon count; ^1H integration vs formula; IR wavenumber windows). They were not applied as a hard filter at harvest; Technical Validation reports a full-corpus post-hoc quarantine on `irexp_resolved` (`scripts/quarantine_structure_nmr.py`

→ `data/audit/structure_nmr_quarantine.jsonl.gz`). Transcription and recall-proxy scripts: `scripts/audit_extraction.py`, `scripts/audit_extraction_recall.py`.

3 Data Records

3.1 Object definition

Each IRExp record is a JSON object. Required chemistry fields for an IR entry:

Table 2 Core fields of an IRExp JSON record.

Field	Type	Description
<code>id</code>	string	Stable record id (InChIKey when resolved, else internal)
<code>ir_bands_cm-1</code>	float list	Experimental IR peak positions (cm^{-1})
<code>ir_source</code>	string	experimental for all released rows
<code>source_doi</code>	string	PMC:<pmcid> or Chemotion deposit DOI
<code>pmcid</code>	string/null	PMC accession when PMC-sourced
<code>h_nmr / c_nmr</code>	string/null	Author-reported NMR shift strings (not parsed peak tables)
<code>smiles / selfies / inchikey</code>	string/null	Resolved structure encodings
<code>has_structure</code>	bool	True when SMILES present
<code>license / license_pool</code>	string	Per-article stamp (Europe PMC join or Chemotion)
<code>license_raw / license_source</code>	string	Upstream licence token + join provenance
<code>source</code>	string	Chemotion rows only (Chemotion)

Frozen field definitions and validation counts: `docs/scientific_data/qc_structure_nmr.json`. Not included: absorbance traces, intensities, instrument metadata beyond what appears in source text, PDFs, figures, or full article bodies.

3.2 Files and counts

Paths relative to the project repository / Hugging Face mirror.

Composition of `irexp.jsonl.gz`.

Licence-pool counts (confirmed Track 1 join, 2026-08-26).

Lookup over 15,416 unique PMCID; stamped pool files under `data/irexp/licence_pools/`. Sum of pools = 121,233 (unchanged).

The full `irexp.jsonl.gz` remains multi-licence on disk; commercial redistribution must use the commercial pool (or `license_pool == "commercial"`). Hugging Face was re-issued with stamped pools and honest card text (`scripts/publish_hf.py`, 2026-08-26).

Table 3 Released IReXP files and record counts.

File	Records	Description
data/irexp/irexp.jsonl.gz	121,233	Full curated release
data/irexp_resolved/ irexp_resolved.jsonl.gz	43,060	100% structure-linked
data/irexp_release/ train_no_bench.jsonl.gz	42,808	Resolved minus IRSpectra-Bench InChIKeys
data/irexp_release/ train_no_bench_nmr.jsonl.gz	32,949	Same with both ^1H and ^{13}C
data/irexp_release/ pretrain_ir.jsonl.gz	119,345	PMC-only IR pretrain pool
data/irexp/seen_papers.txt.gz	188,016 lines	PMC IDs scanned at harvest
data/irexp/ ir_harvest_snapshot.jsonl.gz	134,893	Intermediate harvest snapshot
licence_pools/ irexp_commercial.jsonl.gz	87,617	CC-BY + CC0 (Zenodo/Sci Data primary)
licence_pools/ irexp_non_commercial.jsonl.gz	20,938	NC* held aside
licence_pools/ irexp_sharealike.jsonl.gz	1,897	Chemotion + rare PMC SA
licence_pools/ irexp_empty_unknown.jsonl.gz	10,781	Excluded from commercial

Table 4 Composition of the full curated release.

Slice	Count
All IR band-list records	121,233
With ^1H and/or ^{13}C NMR	87,075 (72%)
With resolved structure	43,060 (35.5%)
Structure + any NMR	40,702
Full IR + ^1H + ^{13}C + structure	33,201
PMC OA provenance	119,345
Chemotion provenance	1,888
Unique PMC accessions	15,416

Median bands: 9 (PMC), 39 (Chemotion). All 1,360,866 released IR band values fall inside $350\text{--}4000\text{ cm}^{-1}$ (full-corpus range check; Technical Validation).

3.3 Access

- Hugging Face: <https://huggingface.co/datasets/ilkhamfy/IReXP>
- GitHub development mirror: <https://github.com/IlkhamFY/spectro-agent>
- Zenodo archival snapshot: DOI pending mint after honest multi-licence metadata (data-only deposit; commercial pool primary — not the combined IRSpectra-Bench research archive).

Table 5 Licence-pool file counts (Track 1 join).

Pool file	Count	Notes
<code>irexp_commercial.jsonl.gz</code>	87,617	CC-BY / CC0 — Zenodo primary
<code>irexp_non_commercial.jsonl.gz</code>	20,938	CC-BY-NC* held aside
<code>irexp_empty_unknown.jsonl.gz</code>	10,781	empty / unresolved — excluded from commercial deposit
<code>irexp_other.jsonl.gz</code>	0	reserved (e.g. CC-BY-ND)
<code>irexp_sharealike.jsonl.gz</code>	1,897	Chemotion CC-BY-SA-4.0 (1,888) + rare PMC SA

4 Technical Validation

Machine-readable package: `docs/scientific_data/qc_structure_nmr.json` (and artefacts under `data/audit/`). Numbers below are descriptive fidelity / consistency checks (fixed seeds); we do not report confidence intervals or claim NMRexp-parity expert audits.

4.1 Transcription fidelity (PMC)

On a seed-fixed random sample of 60 PMC-sourced records (`scripts/audit_extraction.py --n 60 --seed 0`), each article was re-fetched from Europe PMC and every recorded wavenumber was checked against the source text ($\pm 1 \text{ cm}^{-1}$ integer match). Result: 560/560 bands and 60/60 records confirmed. Enlarged sample $n = 200$ (same script/seed family, `data/audit/extraction_audit_n200.json`): 2,250/2,261 bands (99.51%) and 196/200 records (98.0%) fully confirmed. The four incomplete records (PMCID: 5713685, 4189708, 12202350, 6259152) missed 11 bands total on re-fetch match — consistent with plain-text / formatting drift between harvest and Europe PMC re-fetch, not wholesale hallucination. This bounds automated transcription error (mis-parsed or unit-mangled values); it is not a human PDF mark-up of the kind reported for NMRexp [4].

4.2 Extraction-recall proxy (PMC, harvest path)

A human mark-up of every IR string in every paper remains the gold standard. As an automatic proxy on the actual harvest path (`scripts/audit_extraction_recall.py --n 40 --seed 0`): re-fetch PMC-OA S3 plain text for 40 distinct source PMCID; re-run `extract_records`; compare band-sets to released rows ($\pm 1 \text{ cm}^{-1}$). Result: 2,640/2,640 released bands confirmed in re-extract; 248/248 released IR lists recovered (list-level recall proxy 1.0; 40/40 papers). Of 258 re-extracted IR lists, 244 matched a released list (0.9457); 4/40 papers yielded *extra* re-extracted lists relative to the curated release (parser over-fire vs post-harvest dedup/curation, not missed released compounds). Summary: `data/audit/extraction_recall_proxy.json`. This is not a substitute for expert human recall.

4.3 Structure–NMR consistency and quarantine (resolved)

Sample (prior).

On 500 `irexp_resolved` records with ^{13}C text (seed 0): 17/500 (3.4%) listed more peaks than carbons. On 500 with ^1H text: integrals > formula H+2 in 17/497 (3.4%).

Full resolved corpus.

`scripts/quarantine_structure_nmr.py` applied the same physics gates to all 43,060 structure-linked rows. 1,882 (4.37%) fail ≥ 1 hard check and are listed in `data/audit/structure_nmr_quarantine.jsonl.gz` (diagnostic only — release files unchanged). Among rows with the relevant modality: ^{13}C peaks > carbons 1,194/34,231 (3.49%); ^1H integral > formula+2 1,141/39,672 (2.88%); IR out-of-range 0; unparseable SMILES 0. Sample rates and full-corpus rates agree closely, supporting stability of the gates. Re-users should filter quarantined IDs before supervised training (default recommendation: drop the quarantine list unless a noisier corpus is intentionally desired). These rates are integrity diagnostics, not an expert skeleton audit (NMRexp reported $\sim 98\%$ skeleton correctness on a manual $n \approx 300$ sample [4]; that check is optional future work for IRexp).

4.4 IR physical window

Every band in the full 121,233-record release lies in $[350, 4000] \text{ cm}^{-1}$ (0 / 1,360,866 out of range). This is a necessary range gate only: IRexp does not store intensities, solvents, or ATR vs transmission metadata.

4.5 QC status

Table 6 Technical validation checklist.

Check	Status
Per-PMCID licence join + pool counts	Done — commercial 87,617
Transcription fidelity n=60 and n=200	Done (automated re-fetch)
Extraction-recall automatic proxy (n=40 papers, S3 path)	Done (human recall still optional)
Full-corpus structure–NMR quarantine	Done — 1,882 / 43,060 flagged
Expert structure spot-check ($n \geq 100$)	Optional / deferred
NMRexp-style replicate MAE for IR lists	Not applicable / not claimed

5 Usage Notes

- **Band lists $\neq$ spectra.** Do not evaluate models trained on IRexp as if they had seen full absorbance curves.

- **AI agents / LLM tool-use.** Prefer the JSON band-list fields and `source_doi` as tool returns or retrieval hits; do not scrape paywalled PDFs or view-only SDBS pages when a redistributable pool already exists. Productized agents must filter by `license_pool` (commercial vs NC* vs ShareAlike) rather than treating the full dump as a single licence.
- **Separate pools by density and licence.** PMC (sparse, author-transcribed) vs Chemotion (denser, author-curated ELN band lists). Higher median band count does not imply more complete vibrational assignment — Chemotion lists are author-entered, not peak-picked curves. Combined redistribution of Chemotion-derived rows must honour CC-BY-SA-4.0 (ShareAlike); do not relicence SA rows as CC-BY.
- **Do not assume PMC = CC-BY-4.0.** Filter to `license_pool == "commercial"` (or use `irexp_commercial.jsonl.gz`) for commercial redistribution; attribute via `source_doi / pmcid`.
- **Structure–NMR quarantine.** Before supervised training on `irexp_resolved`, drop IDs in `data/audit/structure_nmr_quarantine.jsonl.gz` by default ($\sim 4.4\%$ of resolved rows) unless a noisier training set is intentional.
- **Training without benchmark leakage.** If using the complementary IRSpectra-Bench problems (companion research paper), fine-tune from `train_no_bench.jsonl.gz` (or rebuild with `scripts/build_train_no_bench.py`) so benchmark InChIKeys are withheld. Protocol and model results live only in that companion manuscript.
- **Structure coverage.** Prefer `irexp_resolved` for supervised structure tasks; 64.5% of records lack SMILES.
- **Attribution.** Cite this Data Descriptor / Zenodo DOI (when minted) and attribute originating articles through each record’s `source_doi`.

5.1 Limitations

- **Object.** IRExp is a band-list corpus, not an absorbance-spectrum library and not an NMR resource comparable to NMRexp in scale or annotation richness.
- **Technical Validation depth.** Automated transcription ($n = 200$) and harvest-path recall proxies ($n = 40$ papers) are weaker evidence than NMRexp’s manual $n \approx 300$ PDF checks and replicate MAE analyses. No expert molecular-skeleton audit has been completed for IRExp.
- **Metadata sparsity.** Intensities, solvents, and instrument modes are generally absent; the IR window check is necessary but narrow.
- **Structure coverage and name resolution.** Only 35.5% of records are structure-linked; OPSIN/PubChem failures leave many IR lists without SMILES.
- **Licence mix.** The full `irexp.jsonl.gz` is multi-licence; commercial Zenodo/Sci Data redistribution is the commercial pool only.
- **Scope of this paper.** Elucidation benchmarks and model results are out of scope; cite the companion research manuscript for those claims.

6 Data Availability

IRexp numeric extracts are available at:

- Hugging Face Datasets: <https://huggingface.co/datasets/ilkhamfy/IRexp>
- GitHub: <https://github.com/IkhamFY/spectro-agent> (data/irexp/, data/irexp_resolved/, data/irexp_release/)
- Zenodo: archival DOI at proof for a *data-only* deposit (primary artifact = commercial pool, cc-by-4.0 metadata + SA companion). Do not reuse the combined IRSpectra-Bench research .zenodo.json as the Sci Data archival record.
Licensing summary (honest):
- Chemotion (1,888): CC-BY-SA-4.0 [11].
- PMC (119,345): mixed Creative Commons — stamped per article; commercial redistributable 87,617 (CC-BY/CC0); NC* 20,938 held aside; empty/unknown 10,781 excluded from commercial Zenodo (LICENCE_REMEDIATION.md).
- Only extracted numeric fields and identifiers are redistributed; source full texts are not.

7 Code Availability

Harvesting, extraction, structure resolution, licence pool splitting, and validation scripts are in <https://github.com/IkhamFY/spectro-agent> under the MIT License (LICENSE). Principal modules: `spectro_scraper/` (`extract.py`, `normalize.py`, `pipeline.py`, `quality.py`); release harvest `scripts/s3_ir_harvest.py`, `scripts/chemotion_to_irexp.py`; licence join `scripts/join_pmc_licences.py`; validation `scripts/audit_extraction.py`, `scripts/audit_extraction_recall.py`, `scripts/quarantine_structure_nmr.py`; pool split `scripts/split_license_pools.py`. The version corresponding to this Descriptor will be tagged at Zenodo deposit time.

Author contributions

I.Y.: conceptualization, methodology, software, data curation, validation, writing — original draft.

R.A.V.-H.: conceptualization, supervision, writing — review and editing.

Competing interests

The authors declare no competing interests.

Acknowledgements

Funding and institutional support will be confirmed by the authors at submission / proof (human blocker — not invented here).

References

- [1] National Institute of Standards and Technology. NIST Chemistry WebBook, SRD 69. <https://webbook.nist.gov> (2026). Doi:10.18434/T4D303; accessed 2026-01-01.
- [2] National Institute of Advanced Industrial Science and Technology. Spectral Database for Organic Compounds (SDBS). <https://sdb.sdb.aist.go.jp> (2026). Accessed 2026-01-01.
- [3] Zipoli, F., Alberts, M. & Laino, T. IR-NMR multimodal computational spectra dataset for 177K patent-extracted organic molecules. *Sci. Data* **12**, 1380 (2025).
- [4] Wang, J.-J. *et al.* NMReXP: a database of 3.3 million experimental NMR spectra. *Sci. Data* **12**, 1956 (2025).
- [5] Yang, L. *et al.* *NMRTrans: structure elucidation from experimental NMR spectra via set transformers*, 12621–12632 (2026). ArXiv:2602.10158; NMRSpec corpus context.
- [6] Lowe, D. M., Corbett, P. T., Murray-Rust, P. & Glen, R. C. Chemical name to structure: OPSIN, an open source solution. *J. Chem. Inf. Model.* **51**, 739–753 (2011).
- [7] Landrum, G. RDKit: Open-source cheminformatics. <https://www.rdkit.org> (2026). Accessed 2026-01-01.
- [8] Krenn, M., Häse, F., Nigam, A., Friederich, P. & Aspuru-Guzik, A. Self-referencing embedded strings (SELFIES): a 100% robust molecular string representation. *Mach. Learn.: Sci. Technol.* **1**, 045024 (2020).
- [9] Yabbarov, I. & Vargas-Hernández, R. A. Generation recall, not verification, binds LLM structure elucidation from literature spectra. *ICLR* (2026). Conference manuscript, in preparation.
- [10] National Library of Medicine. NCBI PMC Open Access Subset. <https://www.ncbi.nlm.nih.gov/pmc/tools/openftlist/> (2026). Accessed 2026-01-01.
- [11] Jung, N., Tremouilhac, P., Punjabi, D. & Huang, P.-C. Chemotion Repository – data collection: FT-IR spectroscopy data (Chemotion IR). Tech. Rep., Karlsruhe Institute of Technology (2024).
- [12] Kim, S. *et al.* PubChem 2023 update. *Nucleic Acids Res.* **51**, D1373–D1380 (2023).