

IRexp experimental infrared band lists from the PMC Open Access Subset and Chemotion

Ilkham Yabbarov¹ (*corresponding author: yabbaroi@mcmaster.ca*), Rudra Sondhi¹, Rodrigo A. Vargas-Hernández^{2,1,3}

¹ Department of Chemistry and Chemical Biology, McMaster University, Hamilton, Ontario L8S 4L8, Canada.

² Brockhouse Institute for Materials Research, McMaster University, Hamilton, Ontario L8S 4L8, Canada.

³ School of Computational Science and Engineering, McMaster University, Hamilton, Ontario L8S 4L8, Canada.

Abstract

IRexp is a redistributable collection of **experimental infrared band lists** — author-reported peak positions in cm^{-1} — mined from open chemistry literature, optionally paired with $^1\text{H}/^{13}\text{C}$ shift lists and resolved structures. The release holds **121,233** records (119,345 from the PubMed Central Open Access Subset; 1,888 from Chemotion/RADAR4Chem), of which **43,060** are structure-linked and **33,201** are full IR + ^1H + ^{13}C + structure quadruples. IRexp stores **numeric band lists**, not absorbance traces or JCAMP-DX. Records carry `source_doi` for attribution. Licensing is **mixed**: Chemotion rows are CC-BY-SA-4.0; PMC Open Access uses diverse Creative Commons terms, with per-article commercial/non-commercial segregation pending ([PENDING: commercial count after licence join]). Intended reuse includes multimodal spectral model training, literature-scale cheminformatics, and complementary elucidation benchmarks that cite this Descriptor. Data are on Hugging Face (`ilkhamfy/IRexp`); a Zenodo DOI will follow licence remediation ([TODO: 10.5281/zenodo.XXXXXXX]). Harvest, extraction, and validation code is MIT-licensed.

Background & Summary

Infrared (IR) spectroscopy is routine in organic characterisation, yet **open, redistributable** collections of *experimental* IR data remain sparse relative to modern machine-learning needs. Digitised absorbance libraries such as the NIST Chemistry WebBook¹ and AIST SDBS² are valuable but either modest in size or **view-only** (no bulk redistribution). Computational IR–NMR resources published in *Scientific Data* expand multimodal coverage with simulated spectra³, while large literature mines for **NMR** peak lists (notably NMRexp⁴) demonstrate that peer-reviewed experimental spectral *lists* with DOI traceability are in scope for this journal. Concurrent literature corpora that include IR among other modalities (for example NMRSpec/NMRTrans⁵) further motivate an **IR-focused, redistributable** band-list resource rather than another closed spectrum archive.

IRexp fills a specific wedge. Experimental sections of chemistry papers conventionally report per-compound **IR band lists** (wavenumbers in cm^{-1}) together with $^1\text{H}/^{13}\text{C}$ NMR shift lists. That textual convention is the object language models and many elucidation pipelines consume, and it is a **different object** from a digitised spectrum. IRexp therefore:

1. Harvests open-access full text from the **PMC Open Access Subset** bulk S3 mirror and peak lists from the **Chemotion** FT-IR deposit.
2. Extracts deterministic numeric IR band lists (and co-reported NMR strings where present).
3. Resolves compound names to canonical structures with OPSIN⁶, RDKit⁷, and SELFIES⁸ where possible.

4. Releases **extracted numbers only** — no PDFs, figures, or article full text — with source accessions for attribution.

Among openly redistributable *text-derived IR band lists*, IRexp is large by record count (121,233). It does not claim to replace SDBS or NIST absorbance libraries; SDBS alone holds more structure-linked *spectra* than IRexp holds structure-linked band lists. The intended reuse is multimodal pre-training, supervised IR→structure modelling on `irexp_resolved`, and as the literature substrate for complementary elucidation benchmarks described elsewhere (forthcoming ICLR manuscript on IRSpectra-Bench; cite that work for protocol and model results — **not** reproduced here).

What this Data Descriptor does not contain. No hypothesis tests, no large-language-model accuracy tables, and no stage-decomposition of elucidation performance. Those belong in the complementary research paper.

Methods

Source corpora

PMC Open Access Subset. Primary harvest uses the NCBI PMC OA bulk distribution on Amazon S3 (`s3://pmc-oa-opendata`, HTTPS endpoint `https://pmc-oa-opendata.s3.amazonaws.com`)⁹. A harvest snapshot records **188,016** distinct PMC identifiers scanned (`data/irexp/seen_papers.txt.gz`). The released corpus retains records from **15,416** unique PMC:* accessions. Extraction operates on open-access plain text / JATS available through that OA path (and Europe PMC full-text XML for validation re-fetches). The **released** IRexp DOIs are exclusively PMC:* accessions or the Chemotion deposit DOI; no paywalled publisher DOI appears in the 121,233-record file.

PMC OA is **not** a single licence. The subset is partitioned into commercial-use, non-commercial, and other packages⁹. A Europe PMC licence sample of 200 unique PMCIDs drawn from IRexp found approximately 73% cc by, approximately 19% NC or empty licence fields, and the remainder other/empty/no-result (`data/NOTICE`). **Per-article licence join and segregation of commercial vs non-commercial pools is Track 1 work and is not yet complete** — see Data Records placeholders.

Chemotion / RADAR4Chem. **1,888** records come from the Chemotion Repository FT-IR collection deposited at RADAR4Chem (DOI 10.22000/0GoEQGlsZGEIrgst)¹⁰, licensed **CC-BY-SA-4.0**. These rows are **peak lists derived from deposited experimental FT-IR spectra** (not author-transcribed experimental-section prose). They carry a stamped `license=CC-BY-SA-4.0` and `source_doi` equal to the deposit DOI. Median band count is **39** versus **9** for PMC-transcribed lists; users modelling band density should treat the pools separately.

Excluded. AIST SDBS (view-only; no bulk export)²; NIST WebBook join code exists in the repository but contributes **0** records to the released IRexp (`ir_source` is always `experimental` for released rows).

Discovery and fetch (released corpus)

1. Enumerate PMC OA identifiers / packages from the S3 OA mirror.
2. Fetch OA full text for experimental-section mining.
3. Ingest Chemotion deposit peak lists with structure metadata from the RADAR4Chem release.
4. Persist harvest provenance in `seen_papers.txt.gz` and optional rawer rows in `ir_harvest_snapshot.jsonl`.

(134,893 rows including intermediate prose fields used to diagnose materials-text false positives; the curated release is `irexp.jsonl.gz`).

Development adapters for additional publishers exist in the codebase for exploration; they are **outside the construction path of the released dataset**. Methods for IRexp as published here are restricted to **PMC-OA S3 + Chemotion**.

Extraction (band lists, not spectra)

A deterministic regular-expression pipeline (`spectro_scraper/extract.py`) segments experimental text per compound and extracts:

- IR wavenumbers $\rightarrow$ `ir_bands_cm-1` (list of floats, cm^{-1});
- ^1H and ^{13}C payloads $\rightarrow$ `h_nmr` / `c_nmr` (author strings, when present).

Gates reject scan-range artefacts and common prose false positives (for example hydrogel / materials narrative mistaken for IR lists). **No curve digitisation** is performed: figures are not traced; JCAMP-DX is not stored.

Structure resolution

Where an IUPAC or systematic name is available, names are converted with OPSIN (py2opsin)⁶, canonicalised with RDKit⁷, and encoded as InChIKey and SELFIES⁸. An optional PubChem¹¹ name fallback (`USE_PUBCHEM`) handles trivial names. Structure coverage of the full release is **35.5%** (43,060 / 121,233). The structure-complete split is shipped as `irexp_resolved`.

Licence handling (current vs planned)

Pool	Records	Stamp today	Planned (Track 1)
Chemotion	1,888	CC-BY-SA-4.0	unchanged
PMC	119,345	license unset / null on disk	join each PMCID $\rightarrow$ Europe PMC / PMC OA package licence; segregate commercial / NC / other; stamp every row

`scripts/split_license_pools.py` separates Chemotion vs PMC by **DOI prefix** for physical pool files. Until Track 1 lands, treat its PMC output as “PMC-sourced”, **not** as a verified CC-BY-4.0 commercial pool (`data/NOTICE`).

Quality tooling

Optional physics gates live in `spectro_scraper/quality.py` (^{13}C peak count – carbon count; ^1H integration vs formula; IR wavenumber windows). They were **not** applied as a hard filter to every released row at harvest; Technical Validation reports post-hoc checks below. Transcription fidelity uses `scripts/audit_extraction.py`.

Data Records

Object definition

Each IReXP record is a JSON object. Required chemistry fields for an IR entry:

Field	Type	Description
id	string	Stable record id (InChIKey when resolved, else internal)
ir_bands_cm-1	float list	Experimental IR peak positions (cm ⁻¹)
ir_source	string	experimental for all released rows
source_doi	string	PMC:<pmcid> or Chemotion deposit DOI
h_nmr / c_nmr	string or null	Author-reported shift lists
smiles / selfies / inchikey	string or null	Resolved structure encodings
has_structure	bool	Convenience flag
license	string or null	Stamped for Chemotion; PMC pending Track 1

Not included: absorbance traces, intensities, instrument metadata beyond what appears in source text, PDFs, figures, or full article bodies.

Files and counts

Paths relative to the project repository / Hugging Face mirror.

File	Records	Description
data/irexp/irexp.jsonl.gz	121,233	Full curated release
data/irexp_resolved/irexp_resolved.jsonl.gz	43,060	100% structure-linked
data/irexp_release/train_no_bench.jsonl.gz	42,808	Resolved minus IRSpectra-Bench InChIKeys
data/irexp_release/train_no_bench_nmr.jsonl.gz	32,949	Same with both ¹ H and ¹³ C
data/irexp_release/pretrain_ir.jsonl.gz	119,345	PMC-only IR pretrain pool
data/irexp/seen_papers.txt.gz	188,016 lines	PMC IDs scanned at harvest
data/irexp/ir_harvest_snapshot.jsonl.gz	134,893	Intermediate harvest snapshot

Composition of irexp.jsonl.gz:

Slice	Count
All IR band-list records	121,233
With ¹ H and/or ¹³ C NMR	87,075 (72%)

Slice	Count
With resolved structure	43,060 (35.5%)
Structure + any NMR	40,702
Full IR + ^1H + ^{13}C + structure	33,201
PMC OA provenance	119,345
Chemotion provenance	1,888
Unique PMC accessions	15,416

Licence-pool counts after Track 1 (do not invent):

Pool	Count
PMC commercial-use (CC BY / CC0 / ...)	[PENDING: commercial count after licence join]
PMC non-commercial (CC BY-NC*)	[PENDING: NC count after licence join]
PMC other / empty / unresolved	[PENDING: other count after licence join]
Chemotion CC-BY-SA-4.0	1,888

Median bands: **9** (PMC), **39** (Chemotion). All **1,360,866** released IR band values fall inside 350–4000 cm^{-1} (full-corpus range check; Technical Validation).

Access

- **Hugging Face:** <https://huggingface.co/datasets/ilkhamfy/IRexp> (public mirror; re-issue after Track 1).
- **GitHub development mirror:** <https://github.com/IlkhamFY/spectro-agent>
- **Zenodo archival snapshot:** [TODO: 10.5281/zenodo.XXXXXXX] (mint after honest multi-licence metadata).

Technical Validation

Transcription fidelity (PMC)

On a seed-fixed random sample of **60** PMC-sourced records (`scripts/audit_extraction.py --n 60 --seed 0`), each article was re-fetched from Europe PMC and every recorded wavenumber was checked against the source text ($\pm 1 \text{ cm}^{-1}$ integer match). Result: **560/560 bands** and **60/60 records** confirmed (Wilson 95% CI 99.3–100% for bands; 94–100% for records). Frozen machine-readable summary: `data/audit/extraction_audit.json`.

This bounds **transcription** error (hallucinated, mis-parsed, or unit-mangled values). It does **not** measure extraction **recall** (whether every IR string in every paper was found) — that requires human reading of full experimental sections and remains planned (see below).

Structure–NMR consistency (resolved sample)

On a random sample of **500** `irexp_resolved` records with ^{13}C text (seed 0), comma-separated ^{13}C peaks were counted against RDKit carbon counts: **17/500 (3.4%)** listed more peaks than carbons (physically impossible $\rightarrow$ name/structure assignment or parse error). On **500**

records with ^1H text, summed reported integrals exceeded formula hydrogens by more than 2 in **17/497** scored rows (**3.4%**; 3 lacked parseable integrals). Reproducible summary: `docs/scientific_data/qc_structure_nmr.json`.

These rates are reported as integrity diagnostics for re-users (quarantine or filter before supervised training). They are **not** claimed as a full expert structure audit; NMRexp-scale manual skeleton checks (n 300) remain future work.

IR physical window

Every band in the full 121,233-record release lies in **[350, 4000] cm⁻¹** (0 / 1,360,866 out of range).

Planned / deferred QC

Check	Status
Per-PMCID licence join + pool counts	Pending Track 1
Extraction-recall human audit (papers, not bands)	Planned (n 30–50 papers)
Larger transcription sample (n 200)	Planned
Full-corpus <code>quality.py</code> quarantine pass	Planned
Expert structure spot-check (n 100)	Optional / deferred

Usage Notes

- **Band lists spectra.** Do not evaluate models trained on IRexp as if they had seen full absorbance curves.
- **Separate pools by density and licence.** PMC (sparse, author-transcribed) vs Chemotion (denser, deposit peak-picked). Keep Chemotion ShareAlike constraints in mind when combining pools; combined redistribution of Chemotion-derived rows must honour CC-BY-SA-4.0.
- **Do not assume PMC = CC-BY-4.0.** Wait for Track 1 stamps or restrict use to research compatible with mixed upstream terms and per-article attribution via `source_doi`.
- **Training without benchmark leakage.** If using the complementary IRSpectra-Bench problems, fine-tune from `train_no_bench.jsonl.gz` (or rebuild with `scripts/build_train_no_bench.py`) so benchmark InChIKeys are withheld.
- **Structure coverage.** Prefer `irexp_resolved` for supervised structure tasks; 64.5% of records lack SMILES.
- **Attribution.** Cite this Data Descriptor / Zenodo DOI and attribute originating articles through each record’s `source_doi`.

Data Availability

IRexp numeric extracts are available at:

- Hugging Face Datasets: <https://huggingface.co/datasets/ilkhamfy/IRexp>
- GitHub: <https://github.com/IlkhamFY/spectro-agent> (`data/irexp/`, `data/irexp_resolved/`, `data/irexp_release/`)

- Zenodo: [TODO: 10.5281/zenodo.XXXXXXX] (archival DOI at proof, after Track 1 licence metadata)

Licensing summary (honest):

- **Chemotion (1,888):** CC-BY-SA-4.0¹⁰.
- **PMC (119,345):** mixed Creative Commons / PMC OA package terms — **not** uniformly CC-BY; [PENDING: commercial count after licence join].
- Only extracted numeric fields and identifiers are redistributed; source full texts are not.

Code Availability

Harvesting, extraction, structure resolution, licence pool splitting, and validation scripts are in <https://github.com/IlkhamFY/spectro-agent> under the **MIT License** (LICENSE). Principal modules: `spectro_scraper/` (`extract.py`, `normalize.py`, `pipeline.py`, `quality.py`); validation scripts/`audit_extraction.py`; pool split scripts/`split_license_pools.py`. The version corresponding to this Descriptor will be tagged at Zenodo deposit time.

Author contributions

I.Y.: conceptualization, methodology, software, data curation, validation, writing — original draft.

R.S.: methodology, software, validation, writing — review and editing.

R.A.V.-H.: conceptualization, supervision, writing — review and editing.

Competing interests

The authors declare no competing interests.

Acknowledgements

Funding and institutional support will be added at proof.

References

References are maintained in `docs/scientific_data/references.bib` (pandoc/CSL build) and overlap the project bibliography where shared.

- 1 National Institute of Standards and Technology, NIST Chemistry WebBook, SRD 69, <https://webbook.nist.gov>, (accessed 1 January 2026).
- 2 National Institute of Advanced Industrial Science and Technology, Spectral Database for Organic Compounds (SDBS), <https://sdb.sdb.aist.go.jp>, (accessed 1 January 2026).
- 3 F. Zipoli, M. Alberts and T. Laino, *Sci. Data*, 2025, **12**, 1380.
- 4 J.-J. Wang, Y. Jin, C.-Y. Zhi, Y.-J. Liu, X.-H. Huang, F. Xu, X. Ji, X. Fang, H. Tao, W. E, L. Zhang, G. Ke and R. Zhu, *Sci. Data*, 2025, **12**, 1956.
- 5 L. Yang, Z. Yang, J. Xie, Y. Wang, B. Gao, X. Wei, J. Sun, J. Wu, C. He, Y. Li and Q. Gu, in *Proc. 32nd ACM SIGKDD Conf. Knowledge Discovery and Data Mining (KDD '26)*, Vol. 2, 2026, pp. 12621–12632.

- 6 D. M. Lowe, P. T. Corbett, P. Murray-Rust and R. C. Glen, *J. Chem. Inf. Model.*, 2011, **51**, 739–753.
- 7 G. Landrum, RDKit, <https://www.rdkit.org>, (accessed 1 January 2026).
- 8 M. Krenn, F. Häse, A. Nigam, P. Friederich and A. Aspuru-Guzik, *Mach. Learn.: Sci. Technol.*, 2020, **1**, 045024.
- 9 National Library of Medicine, NCBI PMC Open Access Subset, <https://www.ncbi.nlm.nih.gov/pmc/tools/openftlist/>, (accessed 1 January 2026).
- 10 N. Jung, P. Tremouilhac, D. Punjabi and P.-C. Huang, *Chemotion Repository – data collection: FT-IR spectroscopy data (Chemotion IR)*, Karlsruhe Institute of Technology, 2024.
- 11 S. Kim, J. Chen, T. Cheng, A. Gindulyte, J. He, S. He, Q. Li, B. A. Shoemaker, P. A. Thiessen, B. Yu, L. Zaslavsky, J. Zhang and E. E. Bolton, *Nucleic Acids Res.*, 2023, **51**, D1373–D1380.