

# IRexp: A database of experimental infrared band lists from open literature

Ilkham Yabbarov<sup>1\*</sup> and Rodrigo A. Vargas-Hernández<sup>1,2,3\*</sup>

<sup>1\*</sup>Department of Chemistry and Chemical Biology, McMaster University,  
Hamilton, Ontario, L8S 4L8, Canada.

<sup>2</sup>Brockhouse Institute for Materials Research, McMaster University,  
Hamilton, Ontario, L8S 4L8, Canada.

<sup>3</sup>School of Computational Science and Engineering, McMaster University,  
Hamilton, Ontario, L8S 4L8, Canada.

\*Corresponding author(s). E-mail(s): [yabbaroi@mcmaster.ca](mailto:yabbaroi@mcmaster.ca);  
[vargashr@mcmaster.ca](mailto:vargashr@mcmaster.ca);

## Abstract

IRexp is a redistributable collection of experimental infrared band lists ( $\text{cm}^{-1}$  peak positions) mined from open chemistry literature, optionally with author-reported  $^1\text{H}/^{13}\text{C}$  NMR strings and resolved structures. The release holds 121,233 records (119,345 PMC Open Access Subset; 1,888 Chemotion/RADAR4Chem), including 43,060 structure-linked entries and 33,201 full IR +  $^1\text{H}$  +  $^{13}\text{C}$  + structure quadruples. IRexp stores numeric band lists, not absorbance traces; each record carries a source DOI and stamped licence / licence-pool fields. Licensing is mixed: after Europe PMC joins plus conservative Crossref recovery of previously empty licences, 88,545 records are commercial (CC-BY/CC0), 21,823 non-commercial (NC\*), 8,963 empty/unknown (excluded from the commercial deposit), 1,897 ShareAlike (Chemotion plus rare PMC SA), and 5 other (ND). The resource is intended for multimodal training, retrieval, and tool-input in spectroscopic agents and autonomous chemistry workflows that need structured, attributable experimental peak lists rather than paywalled PDFs or view-only archives. Technical validation covers automated transcription, harvest-path recall proxies, stratified chemist-proxy checks, and full-corpus quarantine; it does not claim NMRexp-equivalent human expert audits. Complementary elucidation benchmarks are described in a companion research manuscript and are not analysed here.

**Keywords:** infrared spectroscopy, band lists, peak lists, PMC Open Access, Chemotion, FAIR data, licence pools, data descriptor

# 1 Background & Summary

Infrared (IR) spectroscopy is routine in organic characterisation, yet open, redistributable collections of *experimental* IR data remain sparse relative to modern machine-learning and agentic tool-use needs. Digitised absorbance libraries such as the NIST Chemistry WebBook [1] and AIST SDBS [2] are valuable but either modest in size or view-only (no bulk redistribution). Computational IR–NMR resources published in *Scientific Data* expand multimodal coverage with simulated spectra [3], while large literature mines for NMR peak lists (notably NMRexp [4]) show that peer-reviewed experimental spectral *lists* with DOI traceability are in scope for this journal. Concurrent literature corpora that include IR among other modalities (for example NMRSpec/NMRTrans [5]) further motivate an IR-focused, redistributable band-list resource rather than another closed spectrum archive.

## *Relation to NMRexp and other peers.*

NMRexp is the natural comparator: ~3.3 million experimental NMR records mined from supporting-information PDFs, with expert-scale manual checks and replicate consistency metrics [4]. IRexp is *not* an NMR database, does not claim size superiority, and is orders of magnitude smaller. Its contribution is complementary: redistributable *IR band lists* ( $\text{cm}^{-1}$  positions only) from PMC Open Access full text plus a Chemotion ELN deposit, with per-record licence pools suitable for commercial vs non-commercial reuse. Absorbance-curve libraries (NIST, SDBS) remain the right choice when full digitised spectra are required; computational IR–NMR sets remain the right choice when simulated multimodal coverage is required [3].

In 2025–26, spectroscopic AI agents and autonomous chemistry workflows increasingly consume *structured* experimental peak lists as tool inputs and retrieval substrates — not paywalled PDF prose and not view-only web UIs. Agents that plan characterisation, call spectrum tools, or train IR→structure models need redistributable numeric lists with DOI attribution and explicit licence pools (commercial vs non-commercial vs ShareAlike). IRexp is designed as that substrate: training / retrieval / tool-input material for literature-grounded spectroscopic agents, without embedding diagnosis benchmarks or model leaderboards in this Descriptor.

IRexp fills a specific wedge. Experimental sections of chemistry papers conventionally report per-compound IR band lists (wavenumbers in  $\text{cm}^{-1}$ ) together with  $^1\text{H}/^{13}\text{C}$  NMR shift lists. That textual convention is the object language models and many elucidation pipelines consume, and it is a different object from a digitised spectrum. IRexp therefore:

1. Harvests open-access full text from the PMC Open Access Subset bulk S3 mirror and peak lists from the Chemotion FT-IR deposit.
2. Extracts deterministic numeric IR band lists (and co-reported NMR strings where present).
3. Resolves compound names to canonical structures with OPSIN [6], RDKit [7], and SELFIES [8] where possible.
4. Releases extracted numbers only — no PDFs, figures, or article full text — with source accessions for attribution.

Among openly redistributable *text-derived IR band lists*, IReXP is large by record count (121,233). It does not claim to replace SDBS or NIST absorbance libraries; SDBS alone holds more structure-linked *spectra* than IReXP holds structure-linked band lists. The scientific contribution of this Descriptor is the curated dataset, harvest provenance, licence segregation, and validation artefacts — not elucidation accuracy claims. The intended reuse is multimodal pretraining, supervised IR→structure modelling on `irexp_resolved`, and as the literature substrate for complementary elucidation benchmarks described elsewhere (forthcoming ICLR manuscript on IRSpectra-Bench; cite that work for protocol and model results — not reproduced here) [9].

***What this Data Descriptor does not contain.***

No hypothesis tests, no large-language-model accuracy tables, and no stage-decomposition of elucidation performance. Those belong in the complementary research paper.

## 2 Methods

### 2.1 Source corpora

***PMC Open Access Subset.***

Primary harvest uses the NCBI PMC OA bulk distribution on Amazon S3 (`s3://pmc-oa-opendata`, HTTPS endpoint `https://pmc-oa-opendata.s3.amazonaws.com`) [10]. Harvest window (recoverable from git snapshots): the bulk S3 IR crawl and `seen_papers / ir_harvest_snapshot` artefacts were produced on 2026-06-07 (UTC) (`scripts/s3_ir_harvest.py`; incremental auto-snapshots that day through ~134,893 raw IR rows). Chemotion was ingested the same day. A harvest snapshot records 188,016 distinct PMC identifiers scanned (`data/irexp/seen_papers.txt.gz`). The released corpus retains records from 15,416 unique PMC:\* accessions. Extraction operates on open-access plain text objects at flat S3 keys `PMC{id}.{v}/PMC{id}.{v}.txt`. Europe PMC full-text XML is used for some validation re-fetches. The released IReXP DOIs are exclusively PMC:\* accessions or the Chemotion deposit DOI; no paywalled publisher DOI appears in the 121,233-record file.

***Discovery vs oa\_comm / oa\_noncomm packages.***

Identifier discovery used NCBI E-utilities `esearch` over PMC with an IR-characterisation query and `open access[filter]`, then fetched matching IDs from the flat S3 text layout (`scripts/s3_ir_harvest.py`). The harvest therefore did not pre-filter by walking the PMC OA Subset’s `oa_comm` vs `oa_noncomm` vs other package directories. PMC OA is still not a single licence [10]. Licence truth for redistribution is applied post-hoc: every unique PMCID in IReXP (15,416) was joined to Europe PMC license metadata (`scripts/join_pmc_licences.py`); each record carries `license / license_pool`. A second-pass recovery (`scripts/apply_crossref_licence_recovery.py`) promoted

1,818 previously empty/unknown rows into commercial, non-commercial, or other pools when Crossref exposed an explicit Creative Commons or ACS AuthorChoice CC URL, or when a fresh Europe PMC query returned a licence string; Elsevier/Springer TDM-only Crossref links were ignored. Commercial-use (CC-BY/CC0) rows form the Zenodo / Sci Data primary pool; NC\* are held aside; remaining empty/unknown are excluded from the commercial deposit — aligning redistributable intent with the OA commercial/non-commercial distinction via article-level licences rather than S3 package paths (`data/NOTICE`; `docs/scientific_data/LICENCE_REMEDIATION.md`).

### ***Reproducibility of discovery.***

Re-running live `esearch` will drift as PMC grows. The frozen discovery set for this release is `data/irexp/seen_papers.txt.gz` (188,016 PMC identifiers); the curated release retains 15,416 of those accessions. Third parties should treat `seen_papers` + the released JSONL as the reproducible snapshot, not a fresh API crawl.

### ***Chemotion / RADAR4Chem.***

1,888 records come from the Chemotion Repository FT-IR collection deposited at RADAR4Chem (DOI 10.22000/0GoEQG1sZGE1rgst) [11], licensed CC-BY-SA-4.0. Ingest (`scripts/chemotion_to_irexp.py`, 2026-06-07): download the MD5-verified deposit; for each of 2,116 ATR-IR analyses, flatten the Quill-delta `content` field to plain text; parse an author-curated IR band list with the same regex extractor and quality gates as the PMC path; resolve the deposit’s canonical SMILES with RDKit → InChIKey + SELFIES; keep the richest band list per InChIKey; drop the 8 molecules already present in the PMC pool → +1,888 new structure-resolved rows. All 1,888 Chemotion rows are structure-linked (`has_structure=true`; `inchkey` equals the record `id`). These are not algorithmic peak-picks from absorbance curves; they are author-entered experimental band lists from the ELN deposit, denser than typical paper prose (median 39 bands vs 9 for PMC). Rows carry `license=CC-BY-SA-4.0`, `source=Chemotion`, and `source_doi` equal to the deposit DOI.

### ***Excluded.***

AIST SDBS (view-only; no bulk export) [2]; NIST WebBook join code exists in the repository but contributes 0 records to the released IReXP (`ir_source` is always `experimental` for released rows).

## **2.2 Discovery and fetch (released corpus)**

1. Discover IR-reporting OA PMCID IDs via NCBI `esearch` (`open_access[filter]` + characterisation-format IR query; year/month sliced).
2. Fetch OA full text from PMC-OA S3 plain-text objects (`PMC{id}.{v}.txt`).
3. Ingest Chemotion deposit author-curated band lists + structures (`scripts/chemotion_to_irexp.py`).
4. Persist harvest provenance in `seen_papers.txt.gz` and optional rawer rows in `ir_harvest_snapshot.jsonl.gz` (134,893 rows including intermediate prose fields; the curated release is `irexp.jsonl.gz`).

### *Non-OA / Scrapling fence.*

Development adapters for ChemRxiv, Beilstein, and generic publisher pages exist for exploration. They are outside the construction path of the released dataset. No released `source_doi` is a paywalled publisher DOI. Methods for IReXP as published here cite only PMC-OA S3 + Chemotion.

## 2.3 Extraction (band lists, not spectra)

A deterministic regular-expression pipeline (`spectro_scraper/extract.py`) segments experimental text per compound and extracts:

- IR wavenumbers  $\rightarrow$  `ir_bands_cm-1` (list of floats,  $\text{cm}^{-1}$ );
- $^1\text{H}$  and  $^{13}\text{C}$  payloads  $\rightarrow$  `h_nmr` / `c_nmr` (author strings, when present).

Gates reject scan-range artefacts and common prose false positives. No curve digitisation is performed: figures are not traced; JCAMP-DX is not stored.

## 2.4 Structure resolution

Where an IUPAC or systematic name is available, names are converted with OPSIN [6], canonicalised with RDKit [7], and encoded as InChIKey and SELFIES [8]. An optional PubChem [12] name fallback (`USE_PUBCHEM`) handles trivial names. Name $\rightarrow$ structure failures are expected for trade names, mixtures, and non-IUPAC prose; unresolved rows keep `has_structure=false` and null structure fields. Structure coverage of the full release is 35.5% (43,060 / 121,233). The structure-complete split is shipped as `irexp_resolved`.

## 2.5 Licence handling

**Table 1** Licence pools stamped on the released IReXP corpus.

| Pool                                 | Records | Stamp                                                            |
|--------------------------------------|---------|------------------------------------------------------------------|
| commercial (CC-BY + CC0)             | 88,545  | <code>license_pool=commercial</code> ; Zenodo / Sci Data primary |
| non_commercial (CC-BY-NC*)           | 21,823  | held aside                                                       |
| sharealike (Chemotion + rare PMC SA) | 1,897   | CC-BY-SA / CC-BY-SA-4.0                                          |
| empty_unknown                        | 8,963   | excluded from commercial deposit                                 |
| other (CC-BY-ND)                     | 5       | held aside                                                       |

`scripts/join_pmc_licences.py` stamps every row;  
`scripts/apply_crossref_licence_recovery.py` recovers empty licences when Crossref/EMPC supply explicit CC terms; `scripts/split_license_pools.py` reports provenance and materialises pool files under `data/irexp/licence_pools/`.  
Narrative and policy: `docs/scientific_data/LICENCE_REMEDIATION.md`.

## 2.6 Quality tooling

Optional physics gates live in `spectro_scraper/quality.py` ( $^{13}\text{C}$  peak count  $\leq$  carbon count;  $^1\text{H}$  integration vs formula; IR wavenumber windows). They were not applied as a hard filter at harvest; Technical Validation reports a full-corpus post-hoc quarantine on `irexp_resolved` (`scripts/quarantine_structure_nmr.py`  $\rightarrow$  `data/audit/structure_nmr_quarantine.jsonl.gz`). Transcription and recall-proxy scripts: `scripts/audit_extraction.py`, `scripts/audit_extraction_recall.py`.

## 3 Data Records

### 3.1 Object definition

Each IReXP record is a JSON object. Required chemistry fields for an IR entry:

**Table 2** Core fields of an IReXP JSON record.

| Field                                     | Type        | Description                                                |
|-------------------------------------------|-------------|------------------------------------------------------------|
| <code>id</code>                           | string      | Stable record id (InChIKey when resolved, else internal)   |
| <code>ir_bands_cm-1</code>                | float list  | Experimental IR peak positions ( $\text{cm}^{-1}$ )        |
| <code>ir_source</code>                    | string      | <b>experimental</b> for all released rows                  |
| <code>source_doi</code>                   | string      | PMC:<pmcid> or Chemotion deposit DOI                       |
| <code>pmcid</code>                        | string/null | PMC accession when PMC-sourced                             |
| <code>h_nmr / c_nmr</code>                | string/null | Author-reported NMR shift strings (not parsed peak tables) |
| <code>smiles / selfies / inchikey</code>  | string/null | Resolved structure encodings                               |
| <code>has_structure</code>                | bool        | True when SMILES present                                   |
| <code>license / license_pool</code>       | string      | Per-article stamp (Europe PMC join or Chemotion)           |
| <code>license_raw / license_source</code> | string      | Upstream licence token + join provenance                   |
| <code>source</code>                       | string      | Chemotion rows only ( <b>Chemotion</b> )                   |

Frozen field definitions and validation counts: `docs/scientific_data/qc_structure_nmr.json`. Not included: absorbance traces, intensities, instrument metadata beyond what appears in source text, PDFs, figures, or full article bodies.

### 3.2 Files and counts

Paths relative to the project repository / Hugging Face mirror.


*Composition of `irexp.jsonl.gz`.*

*Licence-pool counts (confirmed Track 1 join, 2026-08-26).*

Lookup over 15,416 unique PMCIDs; stamped pool files under `data/irexp/licence_pools/`. Sum of pools = 121,233 (unchanged).

**Table 3** Released IReXP files and record counts.

| File                                           | Records       | Description                                     |
|------------------------------------------------|---------------|-------------------------------------------------|
| data/irexp/irexp.jsonl.gz                      | 121,233       | Full curated release                            |
| data/irexp_resolved/irexp_resolved.jsonl.gz    | 43,060        | 100% structure-linked                           |
| data/irexp_release/train_no_bench.jsonl.gz     | 42,808        | Resolved minus IRSpectra-Bench InChIKeys        |
| data/irexp_release/train_no_bench_nmr.jsonl.gz | 32,949        | Same with both $^1\text{H}$ and $^{13}\text{C}$ |
| data/irexp_release/pretrain_ir.jsonl.gz        | 119,345       | PMC-only IR pretrain pool                       |
| data/irexp/seen_papers.txt.gz                  | 188,016 lines | PMC IDs scanned at harvest                      |
| data/irexp/ir_harvest_snapshot.jsonl.gz        | 134,893       | Intermediate harvest snapshot                   |
| licence_pools/irexp_commercial.jsonl.gz        | 88,545        | CC-BY + CC0 (Zenodo/Sci Data primary)           |
| licence_pools/irexp_non_commercial.jsonl.gz    | 21,823        | NC* held aside                                  |
| licence_pools/irexp_sharealike.jsonl.gz        | 1,897         | Chemotion + rare PMC SA                         |
| licence_pools/irexp_empty_unknown.jsonl.gz     | 8,963         | Excluded from commercial                        |

**IReXP release overview (n = 121,233)****Fig. 1** IReXP release overview from frozen counts. **(a)** Provenance (PMC Open Access Subset vs Chemotion/RADAR4Chem). **(b)** Licence pools after Europe PMC joins and Crossref empty-licence recovery (commercial pool is the Sci Data / Zenodo primary artifact). **(c)** Composition cascade from all band-list records to full IR +  $^1\text{H}$  +  $^{13}\text{C}$  + structure quadruples. Numeric values match Tables 1–5.

The full `irexp.jsonl.gz` remains multi-licence on disk; commercial redistribution must use the commercial pool (or `license_pool == "commercial"`). Hugging Face was remirrored with Crossref-recovered pools and honest card text (`scripts/publish_hf.py`, 2026-08-27; commercial 88,545).

Median bands: 9 (PMC), 39 (Chemotion). All 1,360,866 released IR band values fall inside 350–4000  $\text{cm}^{-1}$  (full-corpus range check; Technical Validation).

### 3.3 Access

- Hugging Face: <https://huggingface.co/datasets/ilkhamfy/IReXP>
- GitHub development mirror: <https://github.com/IkhamFY/spectro-agent>

**Table 4** Composition of the full curated release.

| Slice                                                | Count          |
|------------------------------------------------------|----------------|
| All IR band-list records                             | 121,233        |
| With $^1\text{H}$ and/or $^{13}\text{C}$ NMR         | 87,075 (72%)   |
| With resolved structure                              | 43,060 (35.5%) |
| Structure + any NMR                                  | 40,702         |
| Full IR + $^1\text{H}$ + $^{13}\text{C}$ + structure | 33,201         |
| PMC OA provenance                                    | 119,345        |
| Chemotion provenance                                 | 1,888          |
| Unique PMC accessions                                | 15,416         |

**Table 5** Licence-pool file counts (Track 1 join).

| Pool file                                  | Count  | Notes                                                 |
|--------------------------------------------|--------|-------------------------------------------------------|
| <code>irexp_commercial.jsonl.gz</code>     | 88,545 | CC-BY / CC0 — Zenodo primary                          |
| <code>irexp_non_commercial.jsonl.gz</code> | 21,823 | CC-BY-NC* held aside                                  |
| <code>irexp_empty_unknown.jsonl.gz</code>  | 8,963  | empty / unresolved — excluded from commercial deposit |
| <code>irexp_other.jsonl.gz</code>          | 5      | CC-BY-ND (Crossref recovery)                          |
| <code>irexp_sharealike.jsonl.gz</code>     | 1,897  | Chemotion CC-BY-SA-4.0 (1,888) + rare PMC SA          |

- Zenodo archival snapshot: DOI pending mint after honest multi-licence metadata (data-only deposit; commercial pool primary — not the combined IRSpectra-Bench research archive).

## 4 Technical Validation

Machine-readable package: `docs/scientific_data/qc_structure_nmr.json` (and artefacts under `data/audit/`; methods note `docs/scientific_data/TV_AND_SCALE_PASS.md`). Numbers below are seed-fixed descriptive checks; Wilson 95% intervals are reported for the expanded recall proxy and chemist-proxy pass rates. **All checks below are automated** unless marked otherwise — they do **not** claim NMRexp-parity human expert audits.

### 4.1 Transcription fidelity (PMC)

On a seed-fixed random sample of 60 PMC-sourced records (`scripts/audit_extraction.py --n 60 --seed 0`), each article was re-fetched from Europe PMC and every recorded wavenumber was checked against the source text ( $\pm 1\text{cm}^{-1}$  integer match). Result: 560/560 bands and 60/60 records confirmed. Enlarged sample  $n = 200$  (same script/seed family, `data/audit/extraction_audit_n200.json`): 2,250/2,261 bands (99.51%) and 196/200 records (98.0%) fully confirmed. The four incomplete records (PMcIDs 5713685, 4189708, 12202350, 6259152) missed 11 bands total on re-fetch match —



consistent with plain-text / formatting drift between harvest and Europe PMC re-fetch, not wholesale hallucination. This bounds automated transcription error (mis-parsed or unit-mangled values); it is not a human PDF mark-up of the kind reported for NMRexp [4].

## 4.2 Extraction-recall proxy (PMC, harvest path)

A human mark-up of every IR string in every paper remains the gold standard. As an automatic proxy on the actual harvest path (`scripts/audit_extraction_recall.py --n 120 --seed 0`): re-fetch PMC-OA S3 plain text for 120 distinct source PMCID; re-run `extract_records`; compare band-sets to released rows ( $\pm 1\text{ cm}^{-1}$ ). Result (`data/audit/extraction_recall_proxy_n120.json`): 7,981/8,059 released bands confirmed in re-extract (rate 0.9903; Wilson 95% CI [0.9879, 0.9922]); 845/858 released IR lists recovered (list-level recall proxy 0.9848; Wilson 95% CI [0.9743, 0.9911]); 115/120 papers recovered every released list (0.9583; Wilson 95% CI [0.9062, 0.9821]). Of 871 re-extracted IR lists, 811 matched a released list (0.9311); 18/120 papers yielded *extra* re-extracted lists relative to the curated release (parser over-fire vs post-harvest dedup/curation, not missed released compounds). A smaller prior  $n = 40$  run remains archived as `data/audit/extraction_recall_proxy.json`. This is not a substitute for expert human recall.

## 4.3 Stratified chemist-proxy audit (automated)

To approach NMRexp’s  $n \approx 300$  audit *sample size* without fabricating chemist signatures, we ran a seed-fixed stratified multi-check **automated chemist-proxy** (`scripts/audit_chemist_proxy.py --n 280 --seed 0`; `data/audit/chemist_proxy_audit.json`). Strata (quotas): PMC structure-linked commercial (98), PMC structure-linked other licence (42), PMC IR-only (98), Chemotion (42). Checks: IR window and band-list sanity; RDKit formula vs  $^{13}\text{C}$  peak count /  $^1\text{H}$  integral physics on structure-linked rows; PMC Europe PMC + S3 transcription re-confirm ( $\pm 1\text{ cm}^{-1}$ ).

**Table 6** Automated chemist-proxy audit ( $n = 280$ , seed 0). Not a human skeleton audit.

| Stratum / metric                             | $n$   | Pass / rate    | Wilson 95%       |
|----------------------------------------------|-------|----------------|------------------|
| All strata (joint automated checks)          | 280   | 271 / 0.9679   | [0.9401, 0.9830] |
| Chemotion                                    | 42    | 42 / 1.000     | [0.9162, 1.000]  |
| PMC IR-only                                  | 98    | 97 / 0.9898    | [0.9444, 0.9982] |
| PMC structure commercial                     | 98    | 92 / 0.9388    | [0.8728, 0.9716] |
| PMC structure other licence                  | 42    | 40 / 0.9524    | [0.8421, 0.9868] |
| Transcription bands (PMC refetch)            | 2,508 | 2,500 / 0.9968 | —                |
| Transcription records fully confirmed        | 238   | 236 / 0.9916   | —                |
| Structure physics pass (RDKit formula gates) | 182   | 177 / 0.9725   | —                |

Dominant fail modes in-sample:  $^1\text{H}$  integral > formula+2 (4), duplicate integer bands (2), transcription miss (2),  $^{13}\text{C}$  peaks > carbons (1). The 97.25% structure-physics pass rate on 182 structure-linked rows is a **chemist-proxy consistency** metric, not NMRexp’s human molecular-skeleton correctness ( $\sim 98\%$  on  $n \approx 300$  manual checks).

#### 4.4 Structure–NMR consistency and quarantine (resolved)

##### *Sample (prior).*

On 500 `irexp_resolved` records with  $^{13}\text{C}$  text (seed 0): 17/500 (3.4%) listed more peaks than carbons. On 500 with  $^1\text{H}$  text: integrals > formula H+2 in 17/497 (3.4%).

##### *Full resolved corpus.*

`scripts/quarantine_structure_nmr.py` applied the same physics gates to all 43,060 structure-linked rows. 1,882 (4.37%) fail  $\geq 1$  hard check and are listed in `data/audit/structure_nmr_quarantine.jsonl.gz` (diagnostic only — release files unchanged). Among rows with the relevant modality:  $^{13}\text{C}$  peaks > carbons 1,194/34,231 (3.49%);  $^1\text{H}$  integral > formula+2 1,141/39,672 (2.88%); IR out-of-range 0; unparseable SMILES 0. Sample rates and full-corpus rates agree closely, supporting stability of the gates. Re-users should filter quarantined IDs before supervised training (default recommendation: drop the quarantine list unless a noisier corpus is intentionally desired). These rates are integrity diagnostics, not an expert skeleton audit (NMRexp reported  $\sim 98\%$  skeleton correctness on a manual  $n \approx 300$  sample [4]; that human check remains pending for IRexp).

#### 4.5 IR physical window

Every band in the full 121,233-record release lies in  $[350, 4000] \text{ cm}^{-1}$  (0 / 1,360,866 out of range). This is a necessary range gate only: IRexp does not store intensities, solvents, or ATR vs transmission metadata.

#### 4.6 QC status

### 5 Usage Notes

- **Band lists  $\neq$  spectra.** Do not evaluate models trained on IRexp as if they had seen full absorbance curves.
- **AI agents / LLM tool-use.** Prefer the JSON band-list fields and `source_doi` as tool returns or retrieval hits; do not scrape paywalled PDFs or view-only SDBS pages when a redistributable pool already exists. Productized agents must filter by `license_pool` (commercial vs NC\* vs ShareAlike) rather than treating the full dump as a single licence.
- **Separate pools by density and licence.** PMC (sparse, author-transcribed) vs Chemotion (denser, author-curated ELN band lists). Higher median band count does not imply more complete vibrational assignment — Chemotion lists are

**Table 7** Technical validation checklist.

| Check                                               | Status                             |
|-----------------------------------------------------|------------------------------------|
| Per-PMCID licence join + Crossref empty recovery    | Done — commercial 88,545           |
| Transcription fidelity $n=60$ and $n=200$           | Done (automated re-fetch)          |
| Extraction-recall automatic proxy ( $n=120$ papers) | Done (human recall still optional) |
| Stratified chemist-proxy audit ( $n=280$ )          | Done (automated; not human expert) |
| Full-corpus structure-NMR quarantine                | Done — 1,882 / 43,060 flagged      |
| Expert human structure spot-check ( $n \geq 100$ )  | Deferred (human)                   |
| NMRexp-style replicate MAE for IR lists             | Not applicable / not claimed       |

author-entered, not peak-picked curves. Combined redistribution of Chemotion-derived rows must honour CC-BY-SA-4.0 (ShareAlike); do not relicence SA rows as CC-BY.

- **Do not assume PMC = CC-BY-4.0.** Filter to `license_pool == "commercial"` (or use `irexp_commercial.jsonl.gz`) for commercial redistribution; attribute via `source_doi / pmcid`.
- **Structure-NMR quarantine.** Before supervised training on `irexp_resolved`, drop IDs in `data/audit/structure_nmr_quarantine.jsonl.gz` by default ( $\sim 4.4\%$  of resolved rows) unless a noisier training set is intentional.
- **Training without benchmark leakage.** If using the complementary IRSpectra-Bench problems (companion research paper), fine-tune from `train_no_bench.jsonl.gz` (or rebuild with `scripts/build_train_no_bench.py`) so benchmark InChIKeys are withheld. Protocol and model results live only in that companion manuscript.
- **Structure coverage.** Prefer `irexp_resolved` for supervised structure tasks; 64.5% of records lack SMILES.
- **FAIR reuse example.** Load only `license_pool == "commercial"` rows for a redistributable training set; drop quarantine IDs; keep `source_doi / pmcid` for attribution; prefer `irexp_resolved` when SMILES is required. ShareAlike Chemotion rows may be remixed only under CC-BY-SA terms.
- **Attribution.** Cite this Data Descriptor / Zenodo DOI (when minted) and attribute originating articles through each record’s `source_doi`.

## 5.1 Limitations

- **Object.** IRExp is a band-list corpus, not an absorbance-spectrum library and not an NMR resource comparable to NMRexp in scale or annotation richness.
- **Technical Validation depth.** Automated transcription ( $n = 200$ ), harvest-path recall proxies ( $n = 120$  papers with Wilson intervals), and a stratified chemist-proxy ( $n = 280$ ) close the *sample-size* gap versus NMRexp’s  $n \approx 300$  audits, but remain machine checks. They are weaker evidence than NMRexp’s manual PDF

mark-up and replicate MAE analyses. No human molecular-skeleton audit has been completed for IRexp.

- **Metadata sparsity.** Intensities, solvents, and instrument modes are generally absent; the IR window check is necessary but narrow.
- **Structure coverage and name resolution.** Only 35.5% of records are structure-linked; OPSIN/PubChem failures leave many IR lists without SMILES.
- **Licence mix.** The full `irexp.jsonl.gz` is multi-licence; commercial Zenodo/Sci Data redistribution is the commercial pool only (88,545 after Crossref recovery of previously empty rows).
- **Scope of this paper.** Elucidation benchmarks and model results are out of scope for this Data Descriptor; cite the companion research manuscript (IRSpectra-Bench / ICLR track) for those claims. Do not treat this Descriptor and the companion paper as interchangeable publications of the same results.

## 6 Data Availability

IRexp numeric extracts are available at:

- Hugging Face Datasets: <https://huggingface.co/datasets/ilkhamfy/IRexp>
- GitHub: <https://github.com/IlkhamFY/spectro-agent> (data/irexp/, data/irexp\_resolved/, data/irexp\_release/)
- Zenodo: archival DOI at proof for a *data-only* deposit (primary artifact = commercial pool, cc-by-4.0 metadata + SA companion). Do not reuse the combined IRSpectra-Bench research `.zenodo.json` as the Sci Data archival record.  
Licensing summary (honest):
- Chemotion (1,888): CC-BY-SA-4.0 [11].
- PMC (119,345): mixed Creative Commons — stamped per article; commercial redistributable 88,545 (CC-BY/CC0); NC\* 21,823 held aside; empty/unknown 8,963 excluded from commercial Zenodo (LICENCE\_REMEDIATION.md).
- Only extracted numeric fields and identifiers are redistributed; source full texts are not.

## 7 Code Availability

Harvesting, extraction, structure resolution, licence pool splitting, and validation scripts are in <https://github.com/IlkhamFY/spectro-agent> under the MIT License (LICENSE). Principal modules: `spectro_scraper/` (`extract.py`, `normalize.py`, `pipeline.py`, `quality.py`); release harvest `scripts/s3_ir_harvest.py`, `scripts/chemotion_to_irexp.py`; licence join `scripts/join_pmc_licences.py`, `scripts/apply_crossref_licence_recovery.py`; validation `scripts/audit_extraction.py`, `scripts/audit_extraction_recall.py`, `scripts/audit_chemist_proxy.py`, `scripts/quarantine_structure_nmr.py`; pool split `scripts/split_license_pools.py`. The version corresponding to this Descriptor will be tagged at Zenodo deposit time.

## Author contributions

**I.Y.:** conceptualization, methodology, software, data curation, validation, writing — original draft.

**R.A.V.-H.:** conceptualization, supervision, writing — review and editing.

## Competing interests

The authors declare no competing interests.

## Acknowledgements

Funding and institutional support will be confirmed by the authors at submission / proof (human blocker — not invented here).

## References

- [1] National Institute of Standards and Technology. NIST Chemistry WebBook, SRD 69. <https://webbook.nist.gov> (2026). Doi:10.18434/T4D303; accessed 2026-01-01.
- [2] National Institute of Advanced Industrial Science and Technology. Spectral Database for Organic Compounds (SDBS). <https://sdb.sdb.aist.go.jp> (2026). Accessed 2026-01-01.
- [3] Zipoli, F., Alberts, M. & Laino, T. IR-NMR multimodal computational spectra dataset for 177K patent-extracted organic molecules. *Sci. Data* **12**, 1380 (2025).
- [4] Wang, J.-J. *et al.* NMReXP: a database of 3.3 million experimental NMR spectra. *Sci. Data* **12**, 1956 (2025).
- [5] Yang, L. *et al.* *NMRTrans: structure elucidation from experimental NMR spectra via set transformers*, 12621–12632 (2026). ArXiv:2602.10158; NMRSpec corpus context.
- [6] Lowe, D. M., Corbett, P. T., Murray-Rust, P. & Glen, R. C. Chemical name to structure: OPSIN, an open source solution. *J. Chem. Inf. Model.* **51**, 739–753 (2011).
- [7] Landrum, G. RDKit: Open-source cheminformatics. <https://www.rdkit.org> (2026). Accessed 2026-01-01.
- [8] Krenn, M., Häse, F., Nigam, A., Friederich, P. & Aspuru-Guzik, A. Self-referencing embedded strings (SELFIES): a 100% robust molecular string representation. *Mach. Learn.: Sci. Technol.* **1**, 045024 (2020).

- [9] Yabbarov, I. & Vargas-Hernández, R. A. Generation recall, not verification, binds LLM structure elucidation from literature spectra. *ICLR* (2026). Conference manuscript, in preparation.
- [10] National Library of Medicine. NCBI PMC Open Access Subset. <https://www.ncbi.nlm.nih.gov/pmc/tools/openftlist/> (2026). Accessed 2026-01-01.
- [11] Jung, N., Tremouilhac, P., Punjabi, D. & Huang, P.-C. Chemotion Repository – data collection: FT-IR spectroscopy data (Chemotion IR). Tech. Rep., Karlsruhe Institute of Technology (2024).
- [12] Kim, S. *et al.* PubChem 2023 update. *Nucleic Acids Res.* **51**, D1373–D1380 (2023).